

CS-E4600

Algorithmic methods for data mining

Aristides Gionis

Dept of Computer Science

Slide set 4: Similarity search

reading assignment

LRU book : rest of chapter 3

An introductory tutorial on k-d trees
by Andrew Moore

finding similar objects

nearest-neighbor search

again, objects can be

documents

records of users

images


videos

strings

time series

similarity search: applications

in machine learning : nearest-neighbor rule



similarity search: applications

in information retrieval

a user wants to find similar documents or similar images
to a given one

for clustering algorithms

the k-means algorithm assigns points to their
nearest centers

finding similar objects

informal definition

two problems

1. similarity search problem

given a set X of objects (off-line)

given a query object q (query time)

find the object in X that is most similar to q

2. all-pairs similarity problem

given a set X of objects (off-line)

find all pairs of objects in X that are similar

naïve solutions

(assume a distance function $d : X \times X \rightarrow \mathbb{R}$)

I. similarity search problem

given a set X of objects (off-line)

given a query object q (query time)

find the object in X that is most similar to q

naïve solution:

compute $d(q, x)$ for all $x \in X$

return $x^* = \arg \min_{x \in X} d(q, x)$

naïve solutions

(assume a distance function $d : X \times X \rightarrow \mathbb{R}$)

2. all-pairs similarity problem

given a set X of objects (off-line)

find all pairs of objects in X that are similar
(say distance less than t)

naïve solution:

compute $d(x, y)$ for all $x, y \in X$

return all pairs such that $d(x, y) \leq t$

naïve solutions too inefficient

1. similarity search problem

given a set X of objects (off-line)

given a query object q (query time)

find the object in X that is most similar to q

complexity $O(nd)$

applications often require fast answers (milliseconds)

we cannot afford scanning through all objects

goal to beat linear-time algorithm

what does it mean?

$O(\log n)$ $O(\text{poly}(\log n))$ $O(n^{1/2})$ $O(n^{1-\epsilon})$ $O(n+d)$?

naïve solutions too inefficient

2. all-pairs similarity problem

given a set X of objects (**off-line**)

find all pairs of objects in X that are similar

complexity $O(n^2d)$

quadratic time is **prohibitive** for almost any setting
for billions of data points, computation takes years

warm up

let's focus on problem 1

how to solve a problem for 1-d points?

example:

given $X = \{ 5, 9, 1, 11, 14, 3, 21, 7, 2, 17, 26 \}$

given $q=6$, what is the nearest point of q in X ?

answer: sorting and binary search!



any lessons to learn?

1. trade-off preprocessing for query time
2. with one comparison prune away many points

generalization of the idea

space-partition algorithms

many algorithms that follow these principles

k-d trees is a popular variant

k-d trees in 2-d

a data structure to support range queries in \mathbb{R}^2

not the most efficient solution in theory

everyone uses it in practice

preprocessing time : $O(n \log n)$

space complexity : $O(n)$

query time : $O(n^{1/2} + k)$ (for range queries, not similarity search)

k-d trees in 2-d

algorithm :

- choose x or y coordinate (alternate)
- choose the median of the coordinate;
(this defines a horizontal or vertical line)
- recurse on both sides


we get a binary tree

size : $O(n)$



depth : $O(\log n)$

construction time : $O(n \log n)$



construction of k-d trees



the complete k-d tree



how do we search
for a nearest neighbor
in a k-d tree ?



searching in k-d trees

searching for **nearest neighbor** of a query **q**

start from the **root** and visit down the tree

at each point keep the **NN** found so far

before visiting a tree node


estimate a **lower bound** distance

if **lower bound** larger than the current distance to **NN**,

do not visit (**prune**)


(possible to visit both children of a node)

region of a node



region(v) : the subtree rooted at **v** stores the points in **black** dots

lower bound and pruning



green point: query

red point: current NN

purple line: lower bound

searching in k-d trees

range searching in X

given a rectangle R

find all points of X contained in R

consider only **axis-aligned** rectangles R

range searching in k-d trees

start from $v = \text{root}$

$\text{search}(v, R)$

if v is a leaf

then report the point stored in v if it lies in R

otherwise, if $\text{region}(v)$ is contained in R

report all points in the $\text{subtree}(v)$

otherwise:

if $\text{region}(\text{left}(v))$ intersects R

then $\text{search}(\text{left}(v), R)$

if $\text{reg}(\text{right}(v))$ intersects R

then $\text{search}(\text{right}(v), R)$

query time analysis

time required by range searching in k-d trees is $O(n^{1/2}+k)$

where k is the number of points reported

total time to report all points is $O(k)$


just need to bound the number of nodes v such that $\text{region}(v)$

intersects R

but is not contained in R

query time analysis

let $Q(n)$ be the max number of regions in an n -point k-d tree intersecting a line l , boundary of R



if l intersects $\text{region}(v)$

then after two levels it intersects 2 regions

the number of regions intersecting l is $Q(n)=2+2Q(n/4)$

solving the recurrence gives $Q(n)=(n^{1/2})$

k-d trees in d dimensions

supporting range queries in \mathbb{R}^d

preprocessing time : $O(n \log n)$

space complexity : $O(n)$

query time : $O(n^{1-1/d} + k)$

k-d trees in d dimensions

construction is similar as in 2-d

split at the median by alternating coordinates

recursion stops when there is only one point left,
which is stored as a leaf

impact of high dimensionality in similarity search

as dimension grows the similarity search problem becomes harder

for the **range searching** problem this is shown by the $O(n^{1-1/d}+k)$ bound

for the **nearest neighbor** problem, the pruning rule becomes **not effective**

as dimension grows the performance of any index **degrades** to **linear search**

point of **frustration** in the research community

a.k.a. the **curse of the dimensionality**

any catch?

idea **relies** on having **vector-space** objects

what happens with points in a **metric space**?

the **space-partition** idea **generalizes** to **metric spaces**

vantage-point algorithm

consider a metric space (X, d)

partition the objects in X using a binary tree

at each step, when partitioning n objects, choose a point v in X (vantage point)

right subtree $R(v)$:


the set of the $n/2$ points that are closest to v

left subtree $L(v)$:


the rest of the points

recurse on $R(v)$ and $L(v)$


vantage-point algorithm




vantage-point algorithm



vantage-point algorithm




vantage-point algorithm




vantage point
space partition

vantage-point algorithm




vantage-point algorithm




query with distance to
current **NN** : pruning

vantage-point algorithm



query with distance to
current **NN** : pruning

vantage-point algorithm



query with distance to
current **NN** : **NO** pruning

similarity search in metric spaces

1. what are the pruning rules ?
2. can you see how the triangle inequality is used for the vantage-point pruning rules ?

these two questions are left as exercise

problem in metric spaces becomes more difficult than in vector spaces

how to fight against the curse of dimensionality?

idea : approximations!

find approximate nearest neighbors

find approximately similar pairs

why does it make sense?

distance functions are proxies to human notion of similarity

approximate nearest neighbor

given a set X of objects (off-line)


given accuracy parameter ϵ (off-line or query time)

given a query object q (query time)


find an object z in X , such that

$$d(q, z) \leq (1 + \epsilon)d(q, x) \text{ for all } x \text{ in } X$$

k-d trees for approximate similarity search




k-d trees for approximate similarity search



solid circle has radius $d(q, x)$

k-d trees for approximate similarity search



dashed circle has radius $d(q, x)/(1 + e)$

next lecture

locality sensitive hashing