

CS-E4600
Mining data streams
slide set 7

Aristides Gionis
Department of Computer Science
Aalto University

reading assignment

- LRU book: chapter 4
- optional reading
 - paper by Alon, Matias, and Szegedy
[Alon et al., 1999]
 - paper by Charikar, Chen, and Farach-Colton
[Charikar et al., 2002]
 - paper by Cormode and Muthukrishnan
[Cormode and Muthukrishnan, 2005]

data streams

- a data stream is a massive sequence of data
- too large to store (on disk, memory, cache, etc.)
- examples:
 - social media (e.g., twitter feed, foursquare checkins)
 - sensor networks (weather, radars, cameras, etc.)
 - network traffic (trajectories, source/destination pairs)
 - satellite data feed
- how to deal with such data?
- what are the issues?


issues when working with data streams

- space
 - data size is very large
 - often not possible to store the whole dataset
 - inspect each data item, make some computations, do not store it, and never get to inspect it again
 - sometimes data is stored, but making one single pass takes a lot of time, especially when the data is stored on disk
 - can afford a small number of passes over the data
- time
 - data “flies by” at a high speed
 - computation time per data item needs to be small

data streams

- data items can be of **complex types**
 - documents (tweets, news articles)
 - images
 - geo-located time-series
 - ...
- to study basic algorithmic ideas we **abstract away** application-specific details
- consider the data stream as a **sequence of numbers**

data-stream model



data-stream model

- **stream**: m elements from universe of size n , e.g.,

$$\langle x_1, x_2, \dots, x_m \rangle = 6, 1, 7, 4, 9, 1, 5, 1, 5, \dots$$

- **goal**: compute a function over the elements of the stream, e.g., median, number of distinct elements, quantiles, ...
- **constraints**:
 - ① limited working memory, sublinear in n and m
e.g., $\mathcal{O}(\log n + \log m)$,
 - ② access data sequentially
 - ③ limited number of passes, in some cases only one
 - ④ process each element quickly, e.g., $\mathcal{O}(1)$, $\mathcal{O}(\log n)$, etc.

warm up: computing some simple functions

- assume that a number can be stored in $\mathcal{O}(\log n)$ space
- **max**, **min** can be computed with $\mathcal{O}(\log n)$ space
- **sum**, **mean** (average) need $\mathcal{O}(\log n + \log m)$ space

$$\mu_X = \mathbb{E}[X] = \mathbb{E}[x_1, \dots, x_m] = \frac{1}{m} \sum_{i=1}^m x_i$$

- what about **variance**?

$$\begin{aligned}\mathbb{V}ar[X] &= \mathbb{V}ar[x_1, \dots, x_m] = \mathbb{E}[(X - \mathbb{E}[X])^2] \\ &= \frac{1}{m} \sum_{i=1}^m (x_i - \mu_X)^2\end{aligned}$$

- two passes? one pass?

warm up: computing some simple functions

- assume that a number can be stored in $\mathcal{O}(\log n)$ space
- **max**, **min** can be computed with $\mathcal{O}(\log n)$ space
- **sum**, **mean** (average) need $\mathcal{O}(\log n + \log m)$ space

$$\mu_X = \mathbb{E}[X] = \mathbb{E}[x_1, \dots, x_m] = \frac{1}{m} \sum_{i=1}^m x_i$$

- what about **variance**?

$$\begin{aligned}\mathbb{V}ar[X] &= \mathbb{V}ar[x_1, \dots, x_m] = \mathbb{E}[(X - \mathbb{E}[X])^2] \\ &= \frac{1}{m} \sum_{i=1}^m (x_i - \mu_X)^2\end{aligned}$$

- two passes? one pass?

warm up: computing some simple functions

- assume that a number can be stored in $\mathcal{O}(\log n)$ space
- **max**, **min** can be computed with $\mathcal{O}(\log n)$ space
- **sum**, **mean** (average) need $\mathcal{O}(\log n + \log m)$ space

$$\mu_X = \mathbb{E}[X] = \mathbb{E}[x_1, \dots, x_m] = \frac{1}{m} \sum_{i=1}^m x_i$$

- what about **variance**?

$$\begin{aligned}\mathbb{V}ar[X] &= \mathbb{V}ar[x_1, \dots, x_m] = \mathbb{E}[(X - \mathbb{E}[X])^2] \\ &= \frac{1}{m} \sum_{i=1}^m (x_i - \mu_X)^2\end{aligned}$$

- two passes? one pass?

warm up: computing some simple functions

- assume that a number can be stored in $\mathcal{O}(\log n)$ space
- **max**, **min** can be computed with $\mathcal{O}(\log n)$ space
- **sum**, **mean** (average) need $\mathcal{O}(\log n + \log m)$ space

$$\mu_X = \mathbb{E}[X] = \mathbb{E}[x_1, \dots, x_m] = \frac{1}{m} \sum_{i=1}^m x_i$$

- what about **variance**?

$$\begin{aligned}\mathbb{V}ar[X] &= \mathbb{V}ar[x_1, \dots, x_m] = \mathbb{E}[(X - \mathbb{E}[X])^2] \\ &= \frac{1}{m} \sum_{i=1}^m (x_i - \mu_X)^2\end{aligned}$$

- two passes? one pass?

how to tackle massive data streams?

- a general and powerful technique: **sampling**
- idea:
 - ① keep a random sample of the data stream
 - ② perform the computation on the sample
 - ③ extrapolate
- example: compute the median of a data stream
(how to extrapolate in this case?)
- but ... how to keep a random sample of a data stream?

how to tackle massive data streams?

- a general and powerful technique: **sampling**
- idea:
 - ① keep a random sample of the data stream
 - ② perform the computation on the sample
 - ③ extrapolate
- example: compute the median of a data stream
(how to extrapolate in this case?)
- but ... how to keep a random sample of a data stream?

reservoir sampling

- **problem:** take a uniform sample s from a stream of unknown length
- **algorithm:**
 - initially $s \leftarrow x_1$
 - on seeing the t -th element, $s \leftarrow x_t$ with probability $1/t$
- **analysis:**
 - what is the probability that $s = x_i$ at some time $t \geq i$?

$$\begin{aligned}\Pr[s = x_i] &= \frac{1}{i} \cdot \left(1 - \frac{1}{i+1}\right) \cdot \dots \cdot \left(1 - \frac{1}{t-1}\right) \cdot \left(1 - \frac{1}{t}\right) \\ &= \frac{1}{i} \cdot \frac{i}{i+1} \cdot \dots \cdot \frac{t-2}{t-1} \cdot \frac{t-1}{t} = \frac{1}{t}\end{aligned}$$

- how much space? $\mathcal{O}(\log n)$
- to get k samples we need $\mathcal{O}(k \log n)$ bits

reservoir sampling

- **problem:** take a uniform sample s from a stream of unknown length
- **algorithm:**
 - initially $s \leftarrow x_1$
 - on seeing the t -th element, $s \leftarrow x_t$ with probability $1/t$
- **analysis:**
 - what is the probability that $s = x_i$ at some time $t \geq i$?

$$\begin{aligned}\Pr[s = x_i] &= \frac{1}{i} \cdot \left(1 - \frac{1}{i+1}\right) \cdot \dots \cdot \left(1 - \frac{1}{t-1}\right) \cdot \left(1 - \frac{1}{t}\right) \\ &= \frac{1}{i} \cdot \frac{i}{i+1} \cdot \dots \cdot \frac{t-2}{t-1} \cdot \frac{t-1}{t} = \frac{1}{t}\end{aligned}$$

- how much space? $\mathcal{O}(\log n)$
- to get k samples we need $\mathcal{O}(k \log n)$ bits

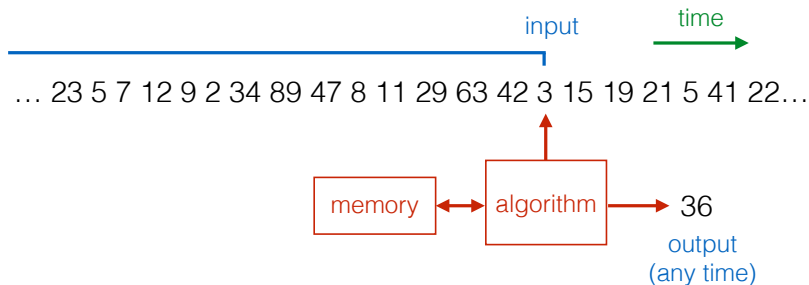
reservoir sampling

- **problem:** take a uniform sample s from a stream of unknown length
- **algorithm:**
 - initially $s \leftarrow x_1$
 - on seeing the t -th element, $s \leftarrow x_t$ with probability $1/t$
- **analysis:**
 - what is the probability that $s = x_i$ at some time $t \geq i$?

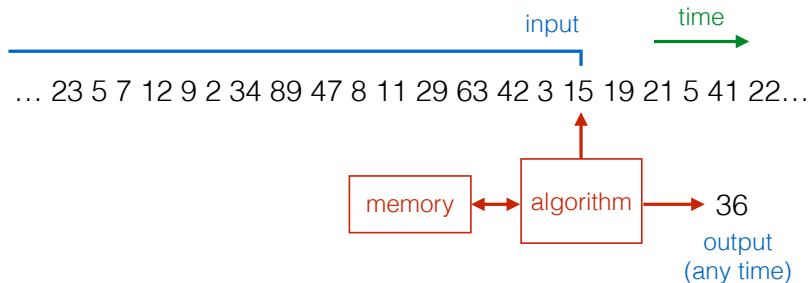
$$\begin{aligned}\Pr[s = x_i] &= \frac{1}{i} \cdot \left(1 - \frac{1}{i+1}\right) \cdot \dots \cdot \left(1 - \frac{1}{t-1}\right) \cdot \left(1 - \frac{1}{t}\right) \\ &= \frac{1}{i} \cdot \frac{i}{i+1} \cdot \dots \cdot \frac{t-2}{t-1} \cdot \frac{t-1}{t} = \frac{1}{t}\end{aligned}$$

- how much space? $\mathcal{O}(\log n)$
- to get k samples we need $\mathcal{O}(k \log n)$ bits

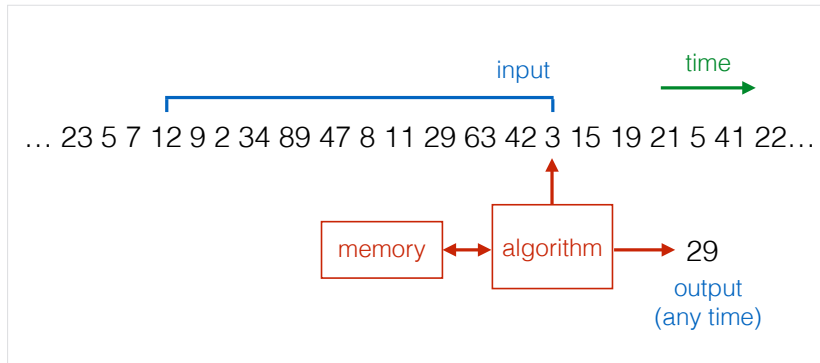
infinite data-stream model



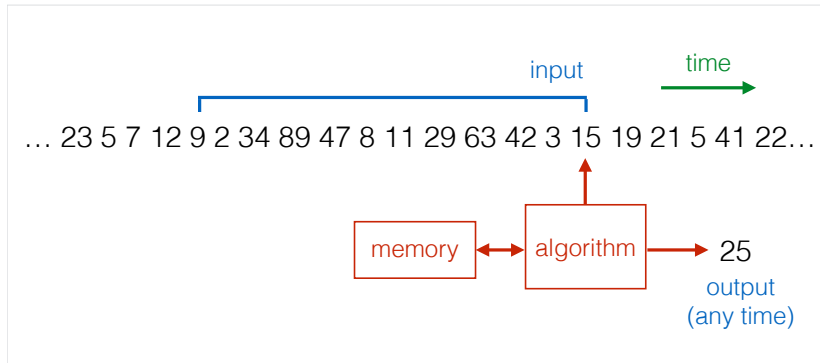
infinite data-stream model



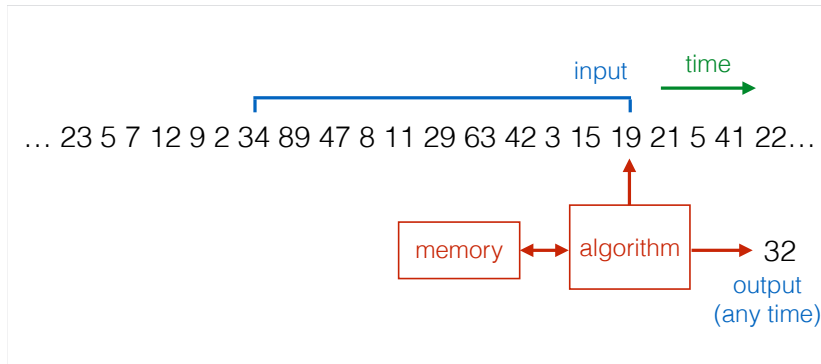
sliding-window data-stream model



sliding-window data-stream model



sliding-window data-stream model



sliding-window data-stream model

- does sliding-window model makes computation **easier** or **harder**?
- how to compute **sum**?
- how to keep a **random sample**?
- all computations can be done with $\mathcal{O}(w)$ space
- can we do better?

sliding-window data-stream model

- does sliding-window model makes computation **easier** or **harder**?
- how to compute **sum**?
- how to keep a **random sample**?
- all computations can be done with $\mathcal{O}(w)$ space
- can we do better?

sliding-window data-stream model

- does sliding-window model makes computation **easier** or **harder**?
- how to compute **sum**?
- how to keep a **random sample**?
- all computations can be done with $\mathcal{O}(w)$ space
- can we do better?

sliding-window data-stream model

- does sliding-window model makes computation **easier** or **harder**?
- how to compute **sum**?
- how to keep a **random sample**?
- all computations can be done with $\mathcal{O}(w)$ space
- can we do better?


priority sampling for sliding window

- maintain a uniform sample from the last w items
- reservoir sampling does not work in this model
- algorithm:
 - 1 for each x_i we pick a random value $v_i \in (0, 1)$
 - 2 for window $\langle x_{j-w+1}, \dots, x_j \rangle$ return x_i with smallest v_i
- to do this, maintain set of all elements in sliding window whose v value is minimal among all subsequent values


priority sampling for sliding window

- maintain a uniform sample from the last w items
- reservoir sampling does not work in this model
- algorithm:
 - 1 for each x_i we pick a random value $v_i \in (0, 1)$
 - 2 for window $\langle x_{j-w+1}, \dots, x_j \rangle$ return x_i with smallest v_i
- to do this, maintain set of all elements in sliding window whose v value is minimal among all subsequent values


priority sampling for sliding window




priority sampling for sliding window




priority sampling for sliding window




priority sampling for sliding window




priority sampling for sliding window




priority sampling for sliding window




priority sampling for sliding window




priority sampling for sliding window



priority sampling for sliding window



priority sampling for sliding window



priority sampling for sliding window

- **correctness 1**: in any given window each item has equal chance to be selected as a random sample
- **correctness 2**: each minimal element x removed from memory has a smaller element y that comes after; x will expire before y , so x will never be needed again
- **correctness 3**: memory has always at least one element

priority sampling for sliding window

- **correctness 1**: in any given window each item has equal chance to be selected as a random sample
- **correctness 2**: each minimal element x removed from memory has a smaller element y that comes after; x will expire before y , so x will never be needed again
- **correctness 3**: memory has always at least one element

priority sampling for sliding window

- **correctness 1**: in any given window each item has equal chance to be selected as a random sample
- **correctness 2**: each minimal element x removed from memory has a smaller element y that comes after; x will expire before y , so x will never be needed again
- **correctness 3**: memory has always at least one element

priority sampling for sliding window

- **correctness 1**: in any given window each item has equal chance to be selected as a random sample
- **correctness 2**: each minimal element x removed from memory has a smaller element y that comes after; x will expire before y , so x will never be needed again
- **correctness 3**: memory has always at least one element

priority sampling for sliding window

- **space efficiency**: how many minimal elements do we expect at any given point?
 - expected number of minimal elements is $\mathcal{O}(\log w)$
 - so, expected space requirement is $\mathcal{O}(\log w \log n)$
- **time efficiency**: maintaining list of minimal elements requires $\mathcal{O}(\log w)$ time

priority sampling for sliding window

- **space efficiency**: how many minimal elements do we expect at any given point?
- expected number of minimal elements is $\mathcal{O}(\log w)$
- so, expected space requirement is $\mathcal{O}(\log w \log n)$
- **time efficiency**: maintaining list of minimal elements requires $\mathcal{O}(\log w)$ time

priority sampling for sliding window

- **space efficiency**: how many minimal elements do we expect at any given point?
- expected number of minimal elements is $\mathcal{O}(\log w)$
- so, expected space requirement is $\mathcal{O}(\log w \log n)$
- **time efficiency**: maintaining list of minimal elements requires $\mathcal{O}(\log w)$ time

mining data streams

- what are **real-world applications**?
- imagine monitoring a **social feed stream**
 - a stream of hashtags in twitter
 - what are interesting questions to ask?
 - do data stream considerations (space/time) really matter?

mining data streams

- what are **real-world applications**?
- imagine monitoring a **social feed stream**
 - a stream of hashtags in twitter
 - what are interesting questions to ask?
 - do data stream considerations (space/time) really matter?

how to tackle massive data streams?

- a general and powerful technique: **sketching**
- general idea:
- apply a linear projection that takes high-dimensional data to a smaller dimensional space
- post-process lower dimensional image to estimate the quantities of interest

computing statistics on data streams

- $X = (x_1, x_2, \dots, x_m)$ a sequence of elements
- each x_i is a member of the set $N = \{1, \dots, n\}$
- $m_i = |\{j : x_j = i\}|$ the number of occurrences of i
- define the k -th frequency moment

$$F_k = \sum_{i=1}^n m_i^k$$

- F_0 is the number of distinct elements
- F_1 is the length of the sequence
- F_2 is the second moment: index of homogeneity, size of self-join, and other applications
- F_∞ frequency of most frequent element

computing statistics on data streams

- $X = (x_1, x_2, \dots, x_m)$ a sequence of elements
- each x_i is a member of the set $N = \{1, \dots, n\}$
- $m_i = |\{j : x_j = i\}|$ the number of occurrences of i
- define the k -th frequency moment

$$F_k = \sum_{i=1}^n m_i^k$$

- F_0 is the number of distinct elements
- F_1 is the length of the sequence
- F_2 is the second moment: index of homogeneity, size of self-join, and other applications
- F_∞^* frequency of most frequent element

computing statistics on data streams

- $X = (x_1, x_2, \dots, x_m)$ a sequence of elements
- each x_i is a member of the set $N = \{1, \dots, n\}$
- $m_i = |\{j : x_j = i\}|$ the number of occurrences of i
- define the k -th frequency moment

$$F_k = \sum_{i=1}^n m_i^k$$

- F_0 is the number of distinct elements
- F_1 is the length of the sequence
- F_2 is the second moment: index of homogeneity, size of self-join, and other applications
- F_∞^* frequency of most frequent element

computing statistics on data streams

- $X = (x_1, x_2, \dots, x_m)$ a sequence of elements
- each x_i is a member of the set $N = \{1, \dots, n\}$
- $m_i = |\{j : x_j = i\}|$ the number of occurrences of i
- define the k -th frequency moment

$$F_k = \sum_{i=1}^n m_i^k$$

- F_0 is the number of distinct elements
- F_1 is the length of the sequence
- F_2 is the second moment: index of homogeneity, size of self-join, and other applications
- F_∞^* frequency of most frequent element

computing statistics on data streams

- $X = (x_1, x_2, \dots, x_m)$ a sequence of elements
- each x_i is a member of the set $N = \{1, \dots, n\}$
- $m_i = |\{j : x_j = i\}|$ the number of occurrences of i
- define the k -th frequency moment

$$F_k = \sum_{i=1}^n m_i^k$$

- F_0 is the number of distinct elements
- F_1 is the length of the sequence
- F_2 is the second moment: index of homogeneity, size of self-join, and other applications
- F_∞^* frequency of most frequent element

computing statistics on data streams

- how much space I need to compute the frequency moments in a straightforward manner?
- how to compute the frequency moments using less than $O(n \log m)$ space?
- problem studied by Alon, Matias, Szegedy [Alon et al., 1999]
- **sketching**: create a sketch that takes much less space and gives an estimation of F_k

computing statistics on data streams

- how much space I need to compute the frequency moments in a straightforward manner?
- how to compute the frequency moments using less than $O(n \log m)$ space?
- problem studied by Alon, Matias, Szegedy
[Alon et al., 1999]
- sketching: create a sketch that takes much less space and gives an estimation of F_k

computing statistics on data streams

- how much space I need to compute the frequency moments in a straightforward manner?
- how to compute the frequency moments using less than $O(n \log m)$ space?
- problem studied by Alon, Matias, Szegedy
[Alon et al., 1999]
- **sketching**: create a sketch that takes much less space and gives an estimation of F_k

estimating the number of distinct values (F_0)

[Flajolet and Martin, 1985]

- consider a bit vector \mathbf{b} with $O(\log n)$ bits
- initialize \mathbf{b} to $[0, \dots, 0]$
- consider a hash function f that maps each item x to the j -th bit of the bit-vector \mathbf{b} with probability $1/2^j$
- for each item x_i in the data stream
 - set the bit $j = f(x_i)$ of \mathbf{b} equal to 1
 - (important: bits are set deterministically for each x_i)
- let R be the index of the largest bit set
- return $Y = 2^R$

estimating the number of distinct values (F_0)

[Flajolet and Martin, 1985]

intuition:

- the j -th bit of \mathbf{b} is set with probability $1/2^j$
- e.g., after seeing 32 distinct elements
 - the bits 1, 2, 3, 4, 5 are most likely set
 - the bits 6, 7, ... are most likely not set
- i.e., we expect the bit vector to be 00000011111, and thus the estimate is 32

estimating number of distinct values (F_0)

Theorem. For every $c > 2$, the algorithm computes a number Y using $\mathcal{O}(\log n)$ memory bits, such that the probability that the ratio between Y and F_0 is not between $1/c$ and c is at most $2/c$.

Theorem proven in [Alon et al., 1999]

estimating F_2

- $X = (x_1, x_2, \dots, x_m)$ a sequence of elements
- each x_i is a member of the set $N = \{1, \dots, n\}$
- $m_i = |\{j : x_j = i\}|$ the number of occurrences of i
- $F_k = \sum_{i=1}^n m_i^k$
- algorithm:
- hash each $i \in \{1, \dots, n\}$ to a random $\epsilon_i \in \{-1, +1\}$
- maintain sketch $Z = \sum_i \epsilon_i m_i$
just need space $\mathcal{O}(\log n + \log m)$
- take $X = Z^2$
- return the average Y of k such estimates X_1, \dots, X_k
- $Y = \frac{1}{k} \sum_{j=1}^k X_j$ where $k = \frac{16}{\lambda^2}$

estimating F_2

- $X = (x_1, x_2, \dots, x_m)$ a sequence of elements
- each x_i is a member of the set $N = \{1, \dots, n\}$
- $m_i = |\{j : x_j = i\}|$ the number of occurrences of i
- $F_k = \sum_{i=1}^n m_i^k$
- algorithm:
- hash each $i \in \{1, \dots, n\}$ to a random $\epsilon_i \in \{-1, +1\}$
- maintain sketch $Z = \sum_i \epsilon_i m_i$
just need space $\mathcal{O}(\log n + \log m)$
- take $X = Z^2$
- return the average Y of k such estimates X_1, \dots, X_k
- $Y = \frac{1}{k} \sum_{j=1}^k X_j$ where $k = \frac{16}{\lambda^2}$

expectation of the estimate is correct

$$\begin{aligned}\mathbb{E}[X] &= \mathbb{E}[Z^2] \\ &= \mathbb{E}\left[\left(\sum_{i=1}^n \epsilon_i m_i\right)^2\right] \\ &= \sum_{i=1}^n m_i^2 \mathbb{E}[\epsilon_i^2] + 2 \sum_{i < j} m_i m_j \mathbb{E}[\epsilon_i] \mathbb{E}[\epsilon_j] \\ &= \sum_{i=1}^n m_i^2 = F_2\end{aligned}$$

accuracy of the estimate

easy to show

$$\mathbb{E}[X^2] = \sum_{i=1}^n m_i^4 + 6 \sum_{i < j} m_i^2 m_j^2$$

which gives

$$\mathbb{V}ar[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2 = 4 \sum_{i < j} m_i^2 m_j^2 \leq 2F_2^2$$

and by Chebyshev's inequality

$$\Pr[|Y - F_2| \geq \lambda F_2] \leq \frac{\mathbb{V}ar[Y]}{\lambda^2 F_2^2} = \frac{\mathbb{V}ar[X]/k}{\lambda^2 F_2^2} \leq \frac{2F_2^2/k}{\lambda^2 F_2^2} = \frac{2}{k\lambda^2} = \frac{1}{8}$$

accuracy of the estimate

easy to show

$$\mathbb{E}[X^2] = \sum_{i=1}^n m_i^4 + 6 \sum_{i < j} m_i^2 m_j^2$$

which gives

$$\mathbb{V}ar[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2 = 4 \sum_{i < j} m_i^2 m_j^2 \leq 2F_2^2$$

and by Chebyshev's inequality

$$\Pr[|Y - F_2| \geq \lambda F_2] \leq \frac{\mathbb{V}ar[Y]}{\lambda^2 F_2^2} = \frac{\mathbb{V}ar[X]/k}{\lambda^2 F_2^2} \leq \frac{2F_2^2/k}{\lambda^2 F_2^2} = \frac{2}{k\lambda^2} = \frac{1}{8}$$

accuracy of the estimate

easy to show

$$\mathbb{E}[X^2] = \sum_{i=1}^n m_i^4 + 6 \sum_{i < j} m_i^2 m_j^2$$

which gives

$$\mathbb{V}ar[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2 = 4 \sum_{i < j} m_i^2 m_j^2 \leq 2F_2^2$$

and by Chebyshev's inequality

$$\Pr[|Y - F_2| \geq \lambda F_2] \leq \frac{\mathbb{V}ar[Y]}{\lambda^2 F_2^2} = \frac{\mathbb{V}ar[X]/k}{\lambda^2 F_2^2} \leq \frac{2F_2^2/k}{\lambda^2 F_2^2} = \frac{2}{k\lambda^2} = \frac{1}{8}$$

estimate of F_2 : summing up

Theorem. Let X_1, \dots, X_k be AMS sketches, with $k = \frac{16}{\lambda^2}$, and Y be their average $Y = \frac{1}{k} \sum_{j=1}^k X_j$.

Then, Y is an unbiased estimator of F_2 , and the quality of approximation is given by

$$\Pr[|Y - F_2| \geq \lambda F_2] \leq \frac{1}{8}$$

references I



Alon, N., Matias, Y., and Szegedy, M. (1999).

The space complexity of approximating the frequency moments.

J. Comput. Syst. Sci., 58(1):137–147.



Charikar, M., Chen, K., and Farach-Colton, M. (2002).

Finding frequent items in data streams.

In *International Colloquium on Automata, Languages, and Programming*, pages 693–703.



Cormode, G. and Muthukrishnan, S. (2005).

An improved data stream summary: the count-min sketch and its applications.

Journal of Algorithms, 55(1):58–75.



Flajolet, P. and Martin, N. G. (1985).

Probabilistic counting algorithms for data base applications.

Journal of Computer and System Sciences, 31(2):182–209.