

CS-E4600

Mining data streams II

slide set 8

Aristides Gionis

Department of Computer Science

Aalto University

reading assignment

- LRU book: chapter 4
- recent Communications of the ACM paper by Cormode
[Cormode, 2017]
- optional reading
 - paper by Charikar, Chen, and Farach-Colton
[Charikar et al., 2002]
 - paper by Cormode and Muthukrishnan
[Cormode and Muthukrishnan, 2005]

finding frequent items in a data stream

- consider again a data stream
- $X = (x_1, x_2, \dots, x_m)$ a data stream
- each x_i is a member of the set $N = \{1, \dots, n\}$
- $m_i = |\{j : x_j = i\}|$ the number of occurrences of i
- $f_i = m_i/m$ the frequency of item i
- problem : estimate most frequent items in data stream

finding frequent items in a data stream

- problem formalization
- rename items $\{o_1, \dots, o_n\}$ so that $m_1 \geq \dots \geq m_n$
- given $k < n$ want to return top- k items o_1, \dots, o_k

finding frequent items in a data stream

- problem formalization — first attempt
- problem $\text{FINDCANDIDATETOP}(X, k, \ell)$
 - given stream X and integers k and ℓ
 - return list of ℓ items, so that top most frequent k items of X occur in the list
- should return all most frequent items

finding frequent items in a data stream

- $\text{FINDCANDIDATETOP}(X, k, \ell)$ can be too hard to solve
- consider the case $m_k = m_{\ell+1} + 1$
 - i.e., number of occurrences of k -th frequent item exceeds only by 1 the number of occurrences of the $(\ell + 1)$ -th frequent item
- almost impossible to find a list that contains the k most frequent items

finding frequent items in a data stream

- problem formalization — second attempt
- problem $\text{FINDAPPROXTOP}(X, k, \epsilon)$
 - given stream X , integer k , and real $\epsilon < 1$
 - return list of k items, so that for each item i in the list it is $m_i \geq (1 - \epsilon)m_k$
- no guarantee to return all most frequent items, but if return an item it should be frequent enough

finding frequent items in a data stream

- problem : $\text{FINDAPPROXTOP}(X, k, \epsilon)$
- algorithm : COUNTSKETCH
 - based on sketching techniques
- intuition
 - use a hash function s and a counter c
 - function s hashes objects to $\{-1, +1\}$
 - for each item o_i seen in the stream, set $c \leftarrow c + s[o_i]$
 - then $\mathbb{E}[c \cdot s[o_i]] = m_i$ (prove it!)
 - so, estimate m_i by $c \cdot s[o_i]$

the COUNTSKETCH algorithm

- problem with using one hash function and one counter
 - very high variance
- remedy 1
 - use t hash functions s_1, \dots, s_t and t counters c_1, \dots, c_t
 - for each item o_i seen in the stream,
 - set $c_j \leftarrow c_j + s_j[o_i]$, for all $j = 1, \dots, t$
 - to estimate m_i take **median** of $\{c_1 \cdot s_1[o_i], \dots, c_t \cdot s_t[o_i]\}$
(as before $\mathbb{E}[c_j \cdot s_j[o_i]] = m_i$ for all $j = 1, \dots, t$)

the COUNTSKETCH algorithm

- problem with using one hash function and one counter
 - very high variance
- remedy 1
 - use t hash functions s_1, \dots, s_t and t counters c_1, \dots, c_t
 - for each item o_i seen in the stream,
 - set $c_j \leftarrow c_j + s_j[o_i]$, for all $j = 1, \dots, t$
 - to estimate m_i take median of $\{c_1 \cdot s_1[o_i], \dots, c_t \cdot s_t[o_i]\}$
(as before $\mathbb{E}[c_j \cdot s_j[o_i]] = m_i$ for all $j = 1, \dots, t$)

the COUNTSKETCH algorithm

- problem with previous idea
 - high-frequency items (e.g., o_1) may spoil estimates of lower-frequency items (e.g., o_k)
- remedy 2
 - do not update all counters with all items
 - replace each counter with a hash table of b counters
 - items update different subsets of counters, one per hash table
 - each item gets enough high-confidence estimates (those avoiding collisions with high-frequency elements)


the COUNTSKETCH algorithm

- problem with previous idea
 - high-frequency items (e.g., o_1) may spoil estimates of lower-frequency items (e.g., o_k)
- remedy 2
 - do not update all counters with all items
 - replace each counter with a hash table of b counters
 - items update different subsets of counters, one per hash table
 - each item gets enough high-confidence estimates (those avoiding collisions with high-frequency elements)

the COUNTSKETCH algorithm

- use parameters t and b
- let h_1, \dots, h_t be hash functions from items to $1, \dots, b$
- let s_1, \dots, s_t be hash functions from items to $\{-1, +1\}$
- consider $t \times b$ table of counters
- for each item o_i seen in the stream,
set $h_j[o_i] \leftarrow h_j[o_i] + s_j[o_i]$, for all $j = 1, \dots, t$
- to estimate m_i take median of
 $\{h_1[o_i] \cdot s_1[o_i], \dots, h_t[o_i] \cdot s_t[o_i]\}$

the COUNTSKETCH data structure



an improved data stream summary

- the COUNTMINSKETCH data stream summary
- see [Cormode, 2017]
- optional reading
[Cormode and Muthukrishnan, 2005]

the COUNTMINSKETCH data stream summary

- **limitations** of existing sketches
 - model limitations (a sequence of items / numbers)
 - space required is $\mathcal{O}(\frac{1}{\epsilon^2})$
recall that guarantees are quantified by ϵ , δ parameters
 ϵ : accuracy
 δ : probability of failure
 - update time proportional to the whole sketch
 - different sketch for each summary
- **COUNTMINSKETCH** addresses all those limitations

incremental data-stream model

- consider a vector $\mathbf{x}(t) = \{x_1(t), \dots, x_n(t)\}$
- number of coordinates n potentially very large
- $\mathbf{x}(t)$ the values of vector at time t
- at each time t a vector coordinate is updated
- data stream : updates (i_t, c_t) for $t = 1, \dots$
- then

$$x_{i_t}(t) \leftarrow x_{i_t}(t-1) + c_t$$

and

$$x_j(t) \leftarrow x_j(t-1), \text{ for } j \neq i_t$$

incremental data-stream model

- generalization of previous model
previous model was $c_t = 1$
- special cases
 - cash register model : $c_t \geq 0$
 - turnstile model : c_t can be negative
 - non-negative turnstile model : $x_i(t) \geq 0$
 - general turnstile model : $x_i(t)$ can be negative

the COUNTMINSKETCH data stream summary

- interesting queries that we would like to handle
 - point query $Q(i)$: approximate x_i
 - range query $Q(\ell, r)$: approximate $\sum_{i=\ell}^r x_i$
 - inner product $Q(\mathbf{x}, \mathbf{y})$: approximate $\mathbf{x} \cdot \mathbf{y} = \sum_{i=1}^n x_i y_i$
 - ϕ -quantiles
 - heavy-hitters : most frequent items
given frequency threshold ϕ , find items i for which
 $x_i \geq (\phi - \epsilon) \|\mathbf{x}\|_1$ for some $\epsilon < \phi$

the COUNTMINSKETCH data structure

- similar to COUNTSKETCH
- a table of counters C of dimension $d \times w$
- d hash functions h_1, \dots, h_d from $\{1, \dots, n\}$ to $\{1, \dots, w\}$ chosen from a pairwise-independent family

$$C = \begin{pmatrix} C[1, 1] & \cdots & C[1, w] \\ \vdots & \ddots & \vdots \\ C[d, 1] & \cdots & C[d, w] \end{pmatrix}$$

- parameters d and w specify the space requirements depend on error bounds we want to achieve

COUNTMINSKETCH : update summary

- given (i_t, c_t) update one counter in each row of C ,
in particular

$$C[j, h_j(i_t)] \leftarrow C[j, h_j(i_t)] + c_t$$

for all $j = 1, \dots, d$

the COUNTMINSKETCH data structure

Figure 2. Count-min sketch data structure with four rows, nine columns.




Figure from "Data Sketching", Cormode, CACM, 2017

COUNTMINSKETCH : point query

- the answer to $Q(i)$ is $\hat{x}_i = \min_j C[j, h_j(i)]$
- theorem : the estimate \hat{x}_i satisfies
 - (i) $x_i \leq \hat{x}_i$
 - (ii) $\hat{x}_i \leq x_i + \epsilon \|\mathbf{x}\|_1$ with prob at least $1 - \delta$

COUNTMINSKETCH

- similar type of estimates for other queries
 - range, inner product, etc.
- parameters are set to

$$d = \left\lceil \log \frac{1}{\delta} \right\rceil \quad \text{and} \quad w = \left\lceil \frac{1}{\epsilon} \right\rceil$$

- improved space ; instead of usual $\mathcal{O}(\frac{1}{\epsilon^2})$
- improved update time : access only d counters

references I



Charikar, M., Chen, K., and Farach-Colton, M. (2002).

Finding frequent items in data streams.

In International Colloquium on Automata, Languages, and Programming, pages 693–703.



Cormode, G. (2017).

Data sketching.

Communications of the ACM, 80(9):48–55.



Cormode, G. and Muthukrishnan, S. (2005).

An improved data stream summary: the count-min sketch and its applications.

Journal of Algorithms, 55(1):58–75.