

CS-E4600

# Algorithmic methods for data mining

Aristides Gionis

dept of Computer Science

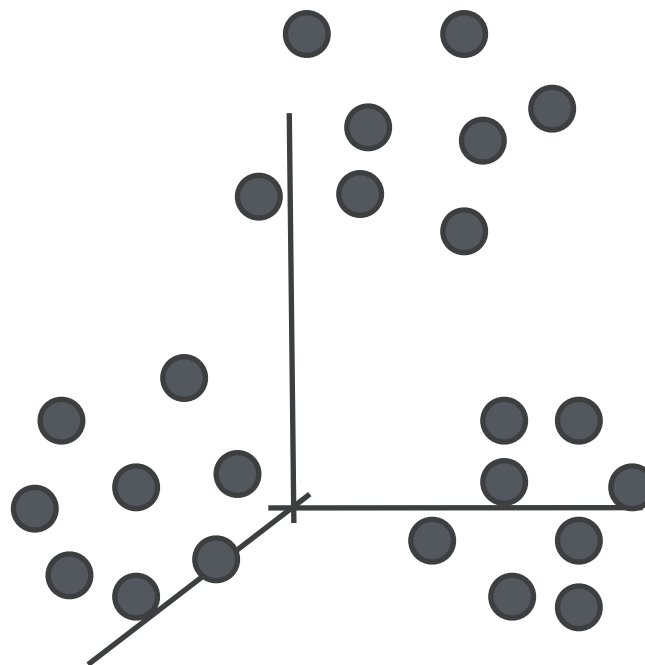
slide set 9: data clustering

# reading assignment

LRU book : chapter 7

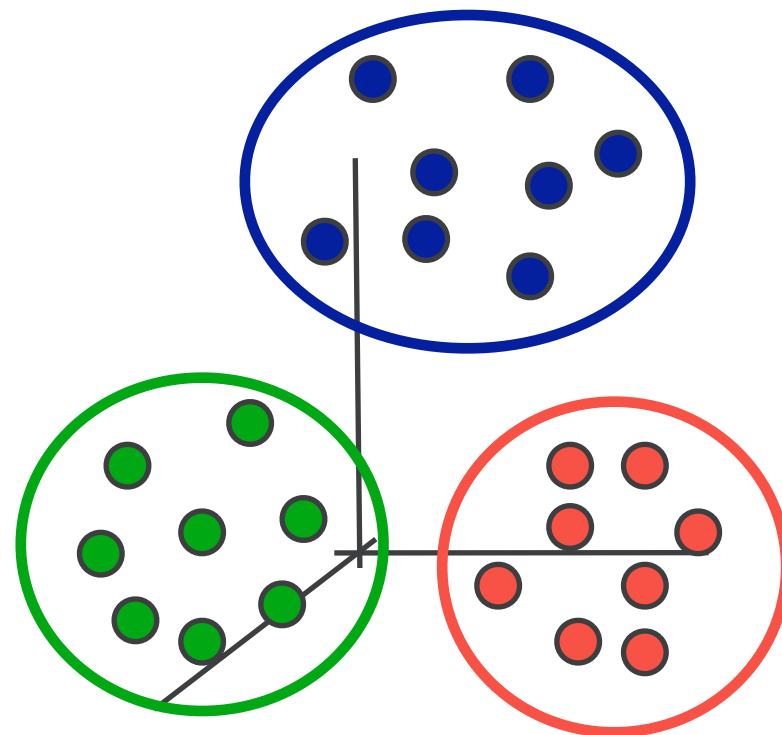
# what is clustering?

a **grouping** of data objects such that the **objects within a group** are **similar** (or **near**) to one another and **dissimilar** (or **far**) from the **objects in other groups**



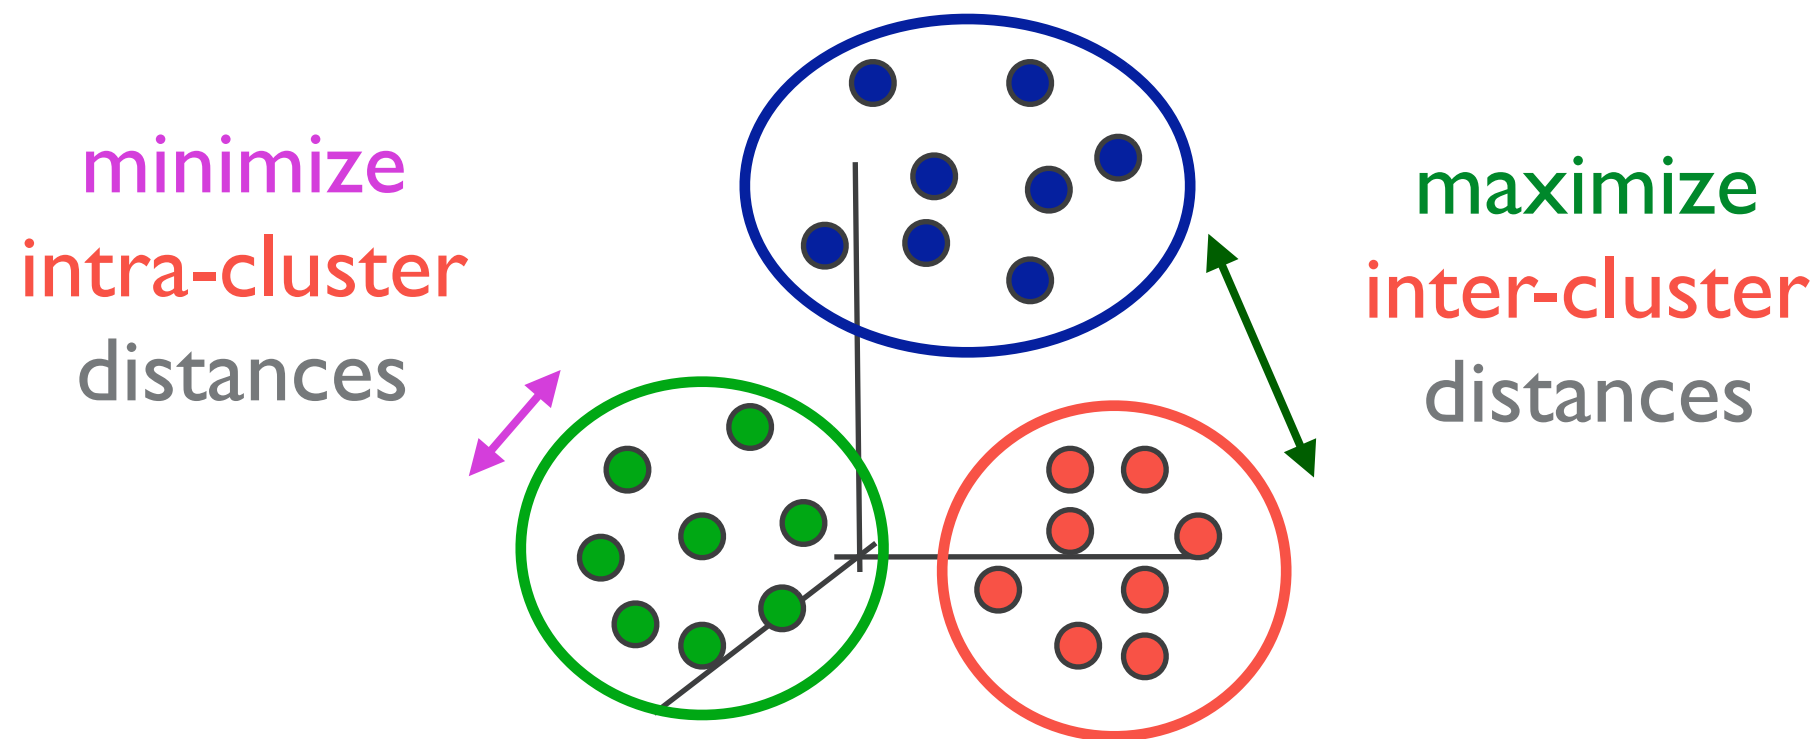
# what is clustering?

a **grouping** of data objects such that the **objects within a group** are **similar** (or **near**) to one another and **dissimilar** (or **far**) from the **objects in other groups**



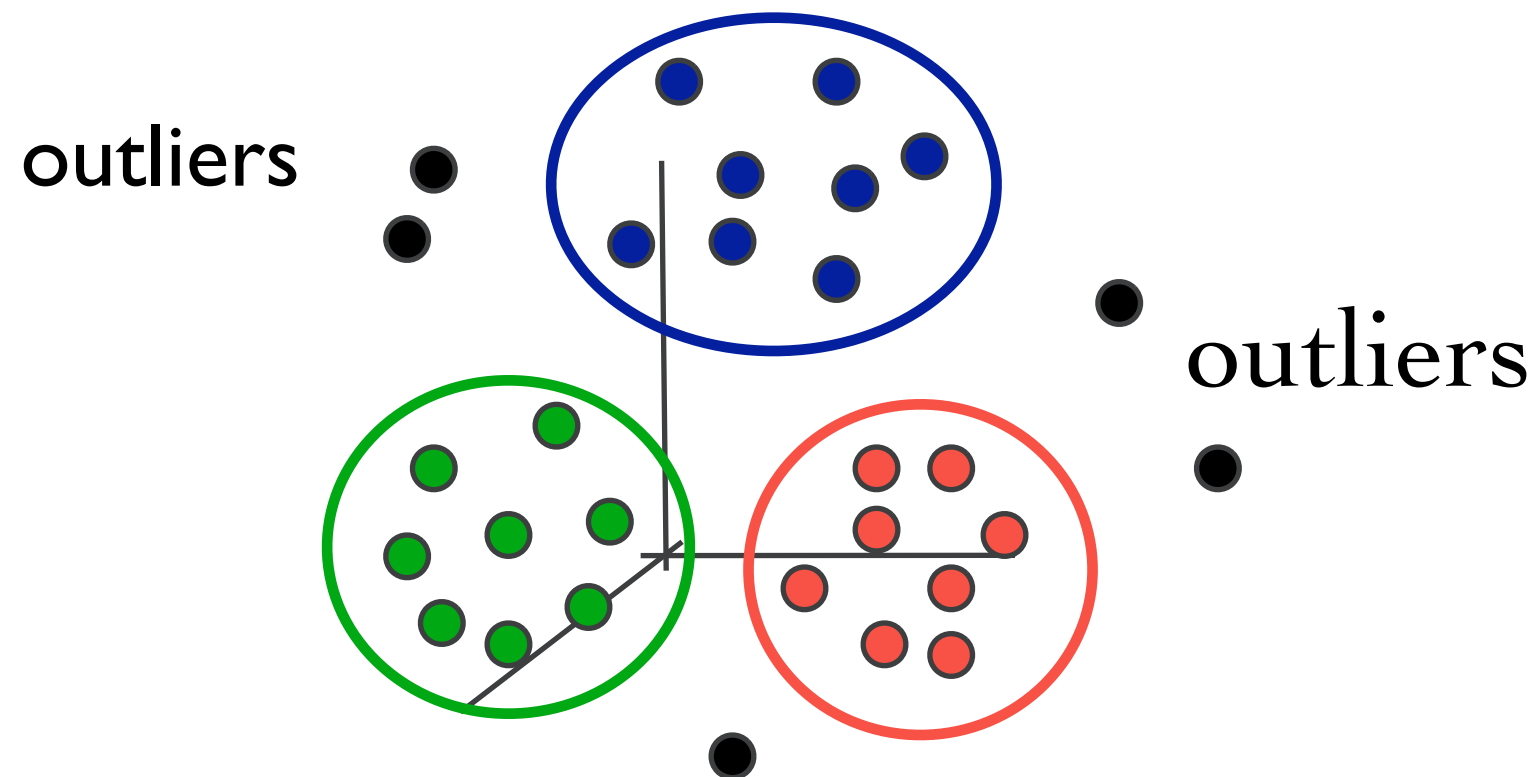
# how to capture this objective?

a **grouping** of data objects such that the **objects within a group** are **similar** (or **near**) to one another and **dissimilar** (or **far**) from the **objects in other groups**



# outliers

outliers are objects that do not belong to any cluster, or form very small clusters



sometimes, we are interested in discovering outliers, not clusters (outlier detection)

# clustering — why care?

stand-alone tool to gain insight into the data

visualization

preprocessing step for other algorithms

indexing or compression often relies on clustering

# applications of clustering

## image processing

cluster images based on their visual content

## market segmentation

cluster customers based on their behavior

## bioinformatics

cluster similar proteins together (similarity wrt chemical structure and/or functionality etc)

many more...



# clustering — high-level definition

given a collection of data objects

find a grouping so that

similar objects are in the same cluster

dissimilar objects are in different clusters

# clustering — basic questions

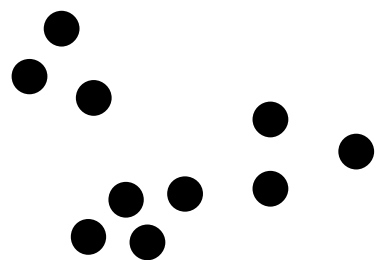
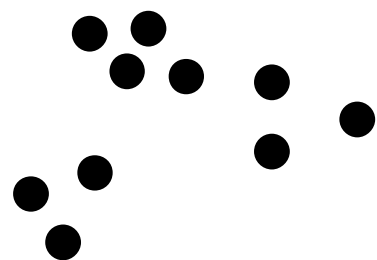
what does **similar** mean?

what is a **good partition** of the objects?

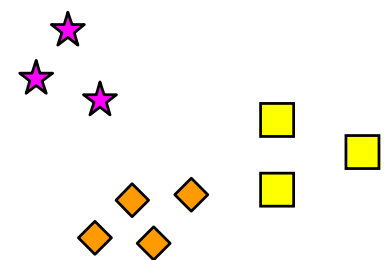
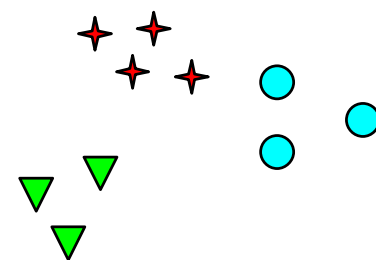
i.e., how is the quality of a solution measured?

**how to find** a good partition?

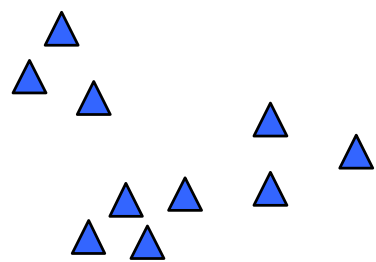
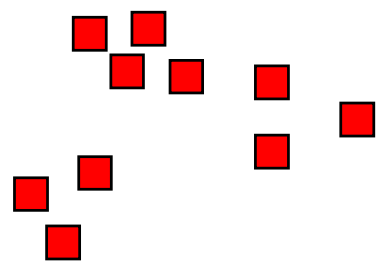
# notion of a cluster can be ambiguous



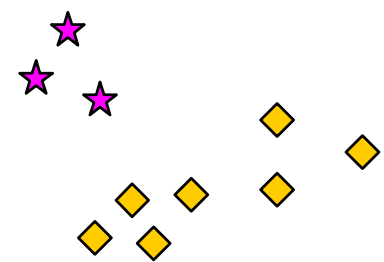
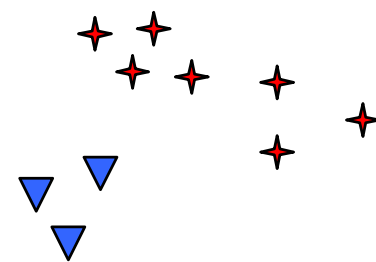
How many clusters?



Six Clusters



Two Clusters



Four Clusters

# types of clusterings

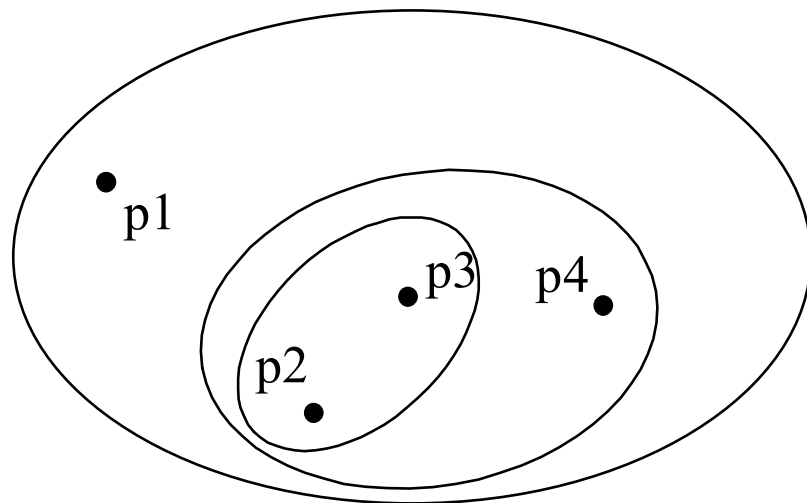
## partitional

each object belongs in exactly one cluster

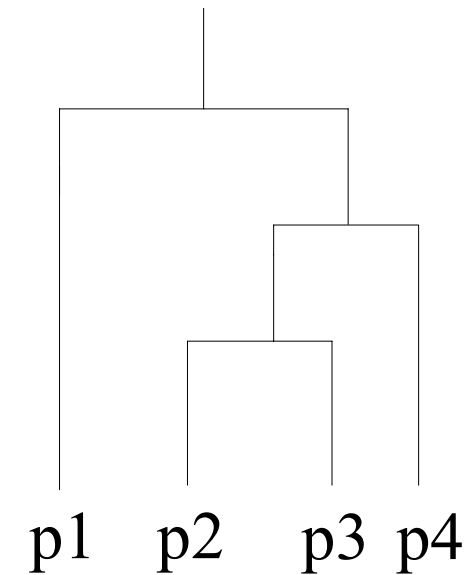
## hierarchical

a set of nested clusters organized in a tree

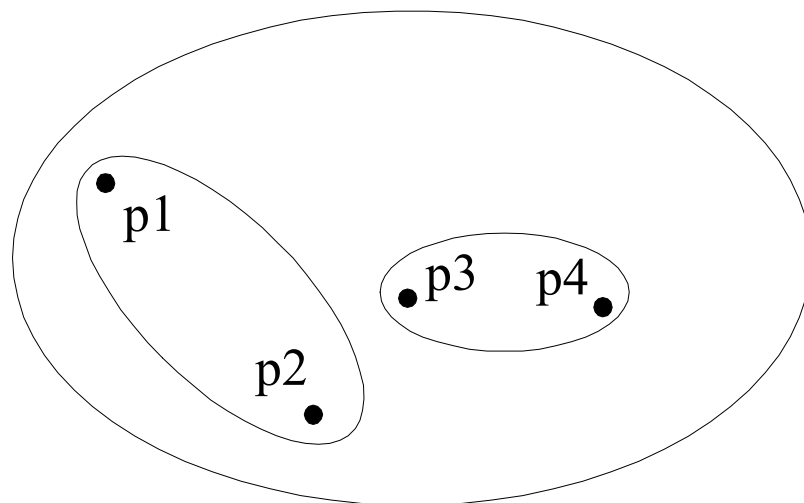
# hierarchical clustering



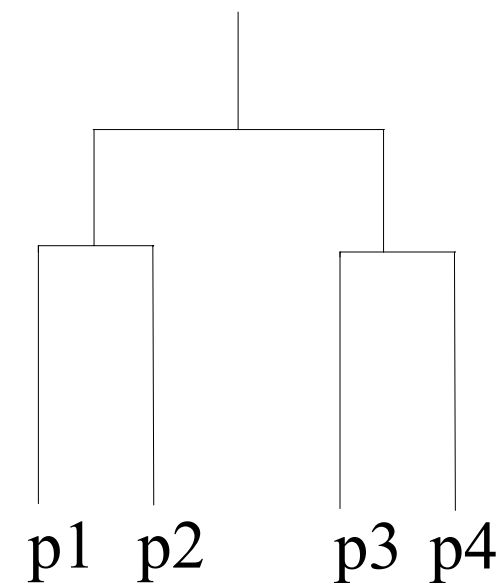
**Hierarchical Clustering**



**Dendrogram**

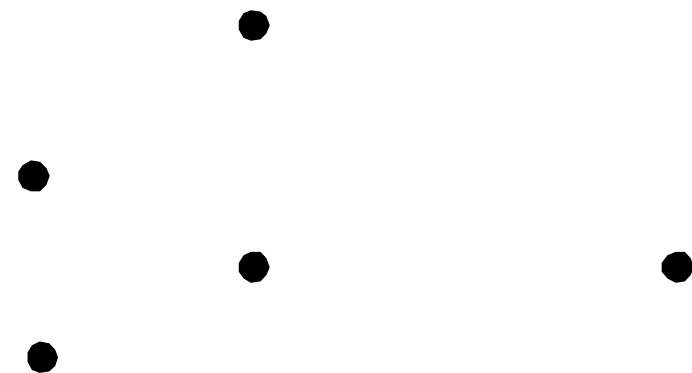
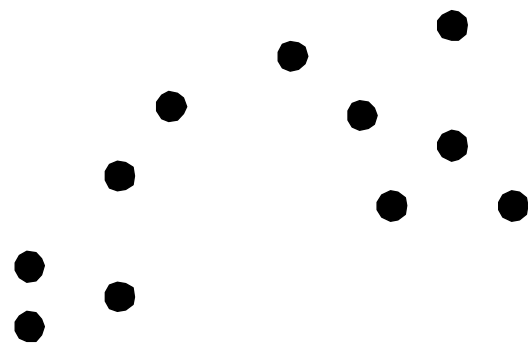


**Hierarchical Clustering**

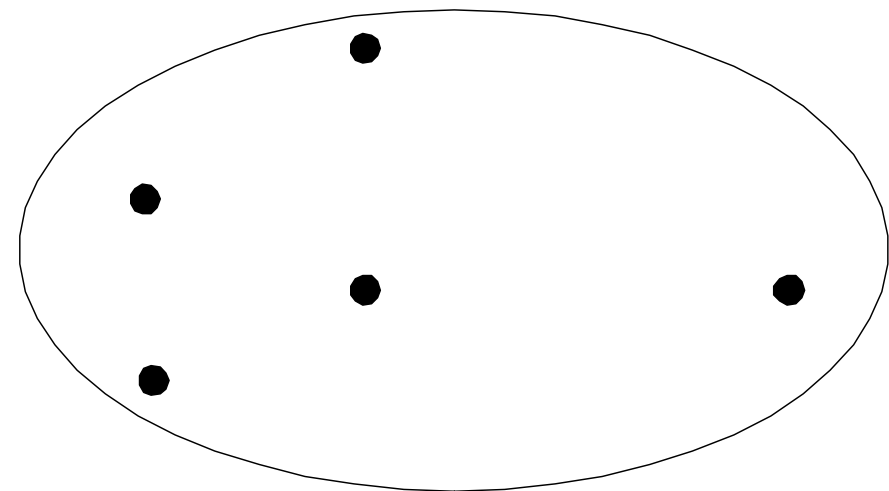
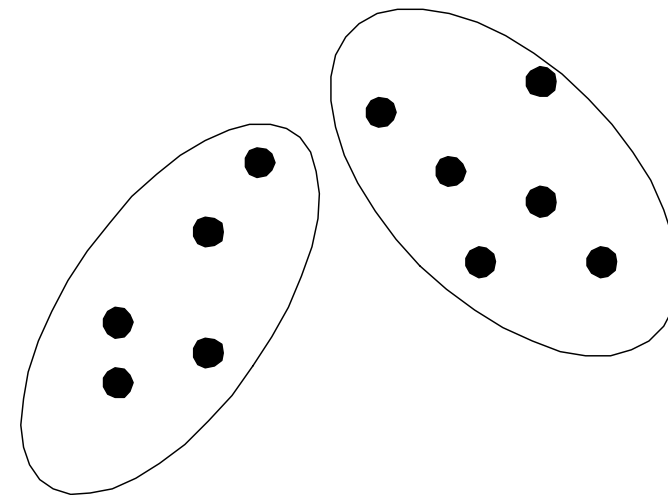


**Dendrogram**

# partitional clustering



**Original Points**



**A Partitional Clustering**

# partitional algorithms

partition the  $n$  objects into  $k$  clusters

each object belongs to exactly one cluster

the number of clusters  $k$  is given in advance

# the k-means problem

consider set  $X = \{x_1, \dots, x_n\}$  of  $n$  points in  $\mathbb{R}^d$

assume that the number  $k$  is given

problem:

find  $k$  points  $c_1, \dots, c_k$  (named centers or means)

so that the cost

$$\sum_{i=1}^n \min_j \{L_2^2(x_i, c_j)\} = \sum_{i=1}^n \min_j \|x_i - c_j\|_2^2$$

is minimized



# the k-means problem

consider set  $X = \{x_1, \dots, x_n\}$  of  $n$  points in  $\mathbb{R}^d$

assume that the number  $k$  is given

problem:

find  $k$  points  $c_1, \dots, c_k$  (named centers or means)

and partition  $X$  into  $\{X_1, \dots, X_k\}$  by assigning each point  $x_i$  in  $X$  to its nearest cluster center,

so that the cost

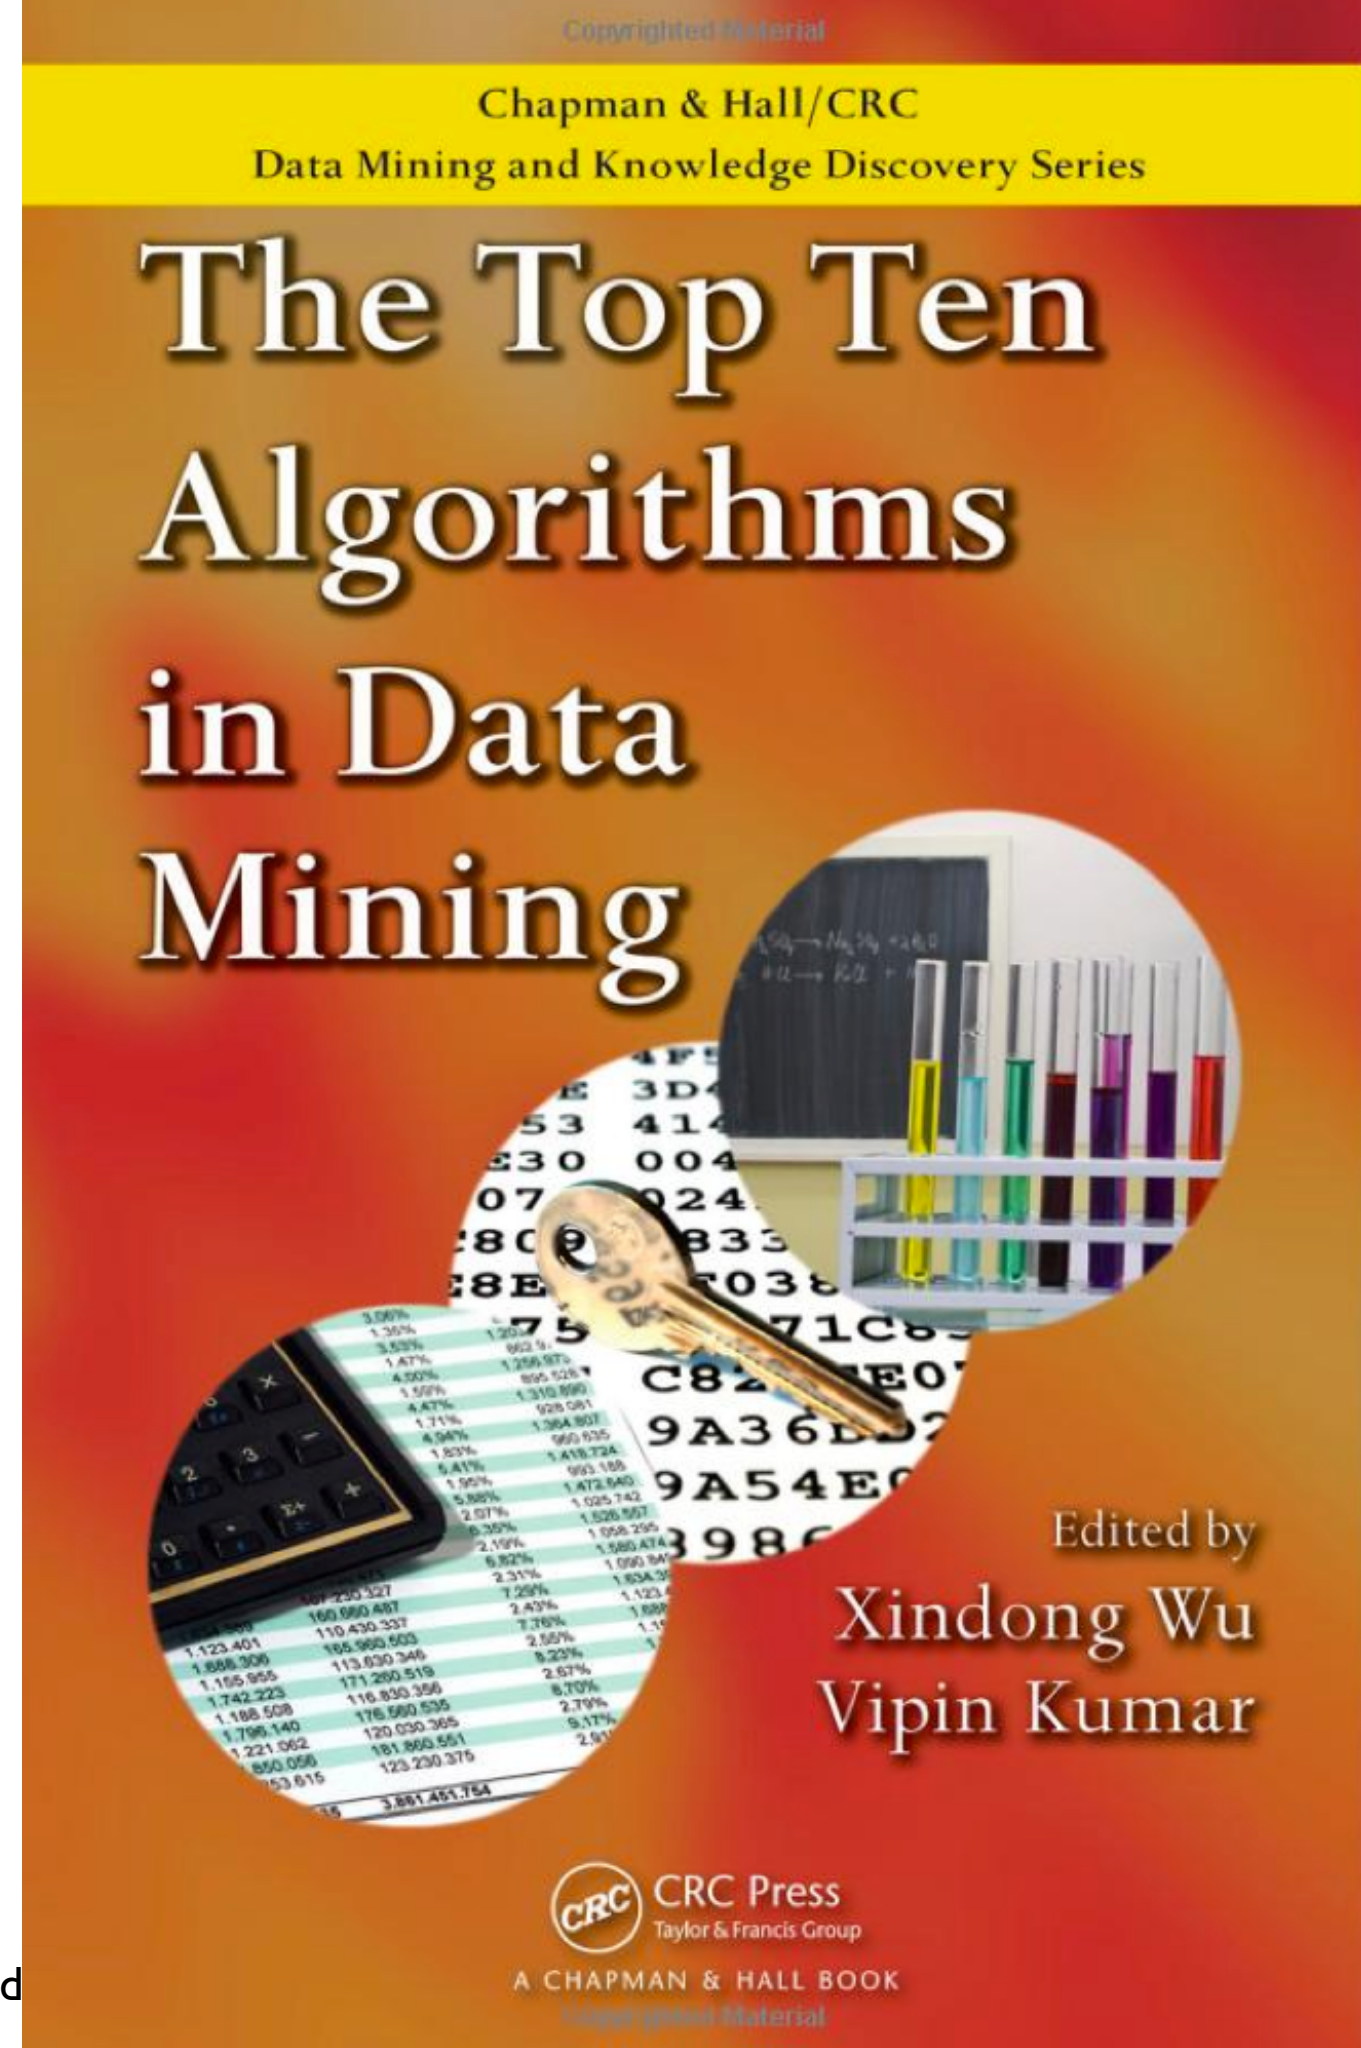
$$\sum_{i=1}^n \min_j \|x_i - c_j\|_2^2 = \sum_{j=1}^k \sum_{x \in X_j} \|x - c_j\|_2^2$$

is minimized

# the k-means algorithm

voted among the **top-10 algorithms** in data mining

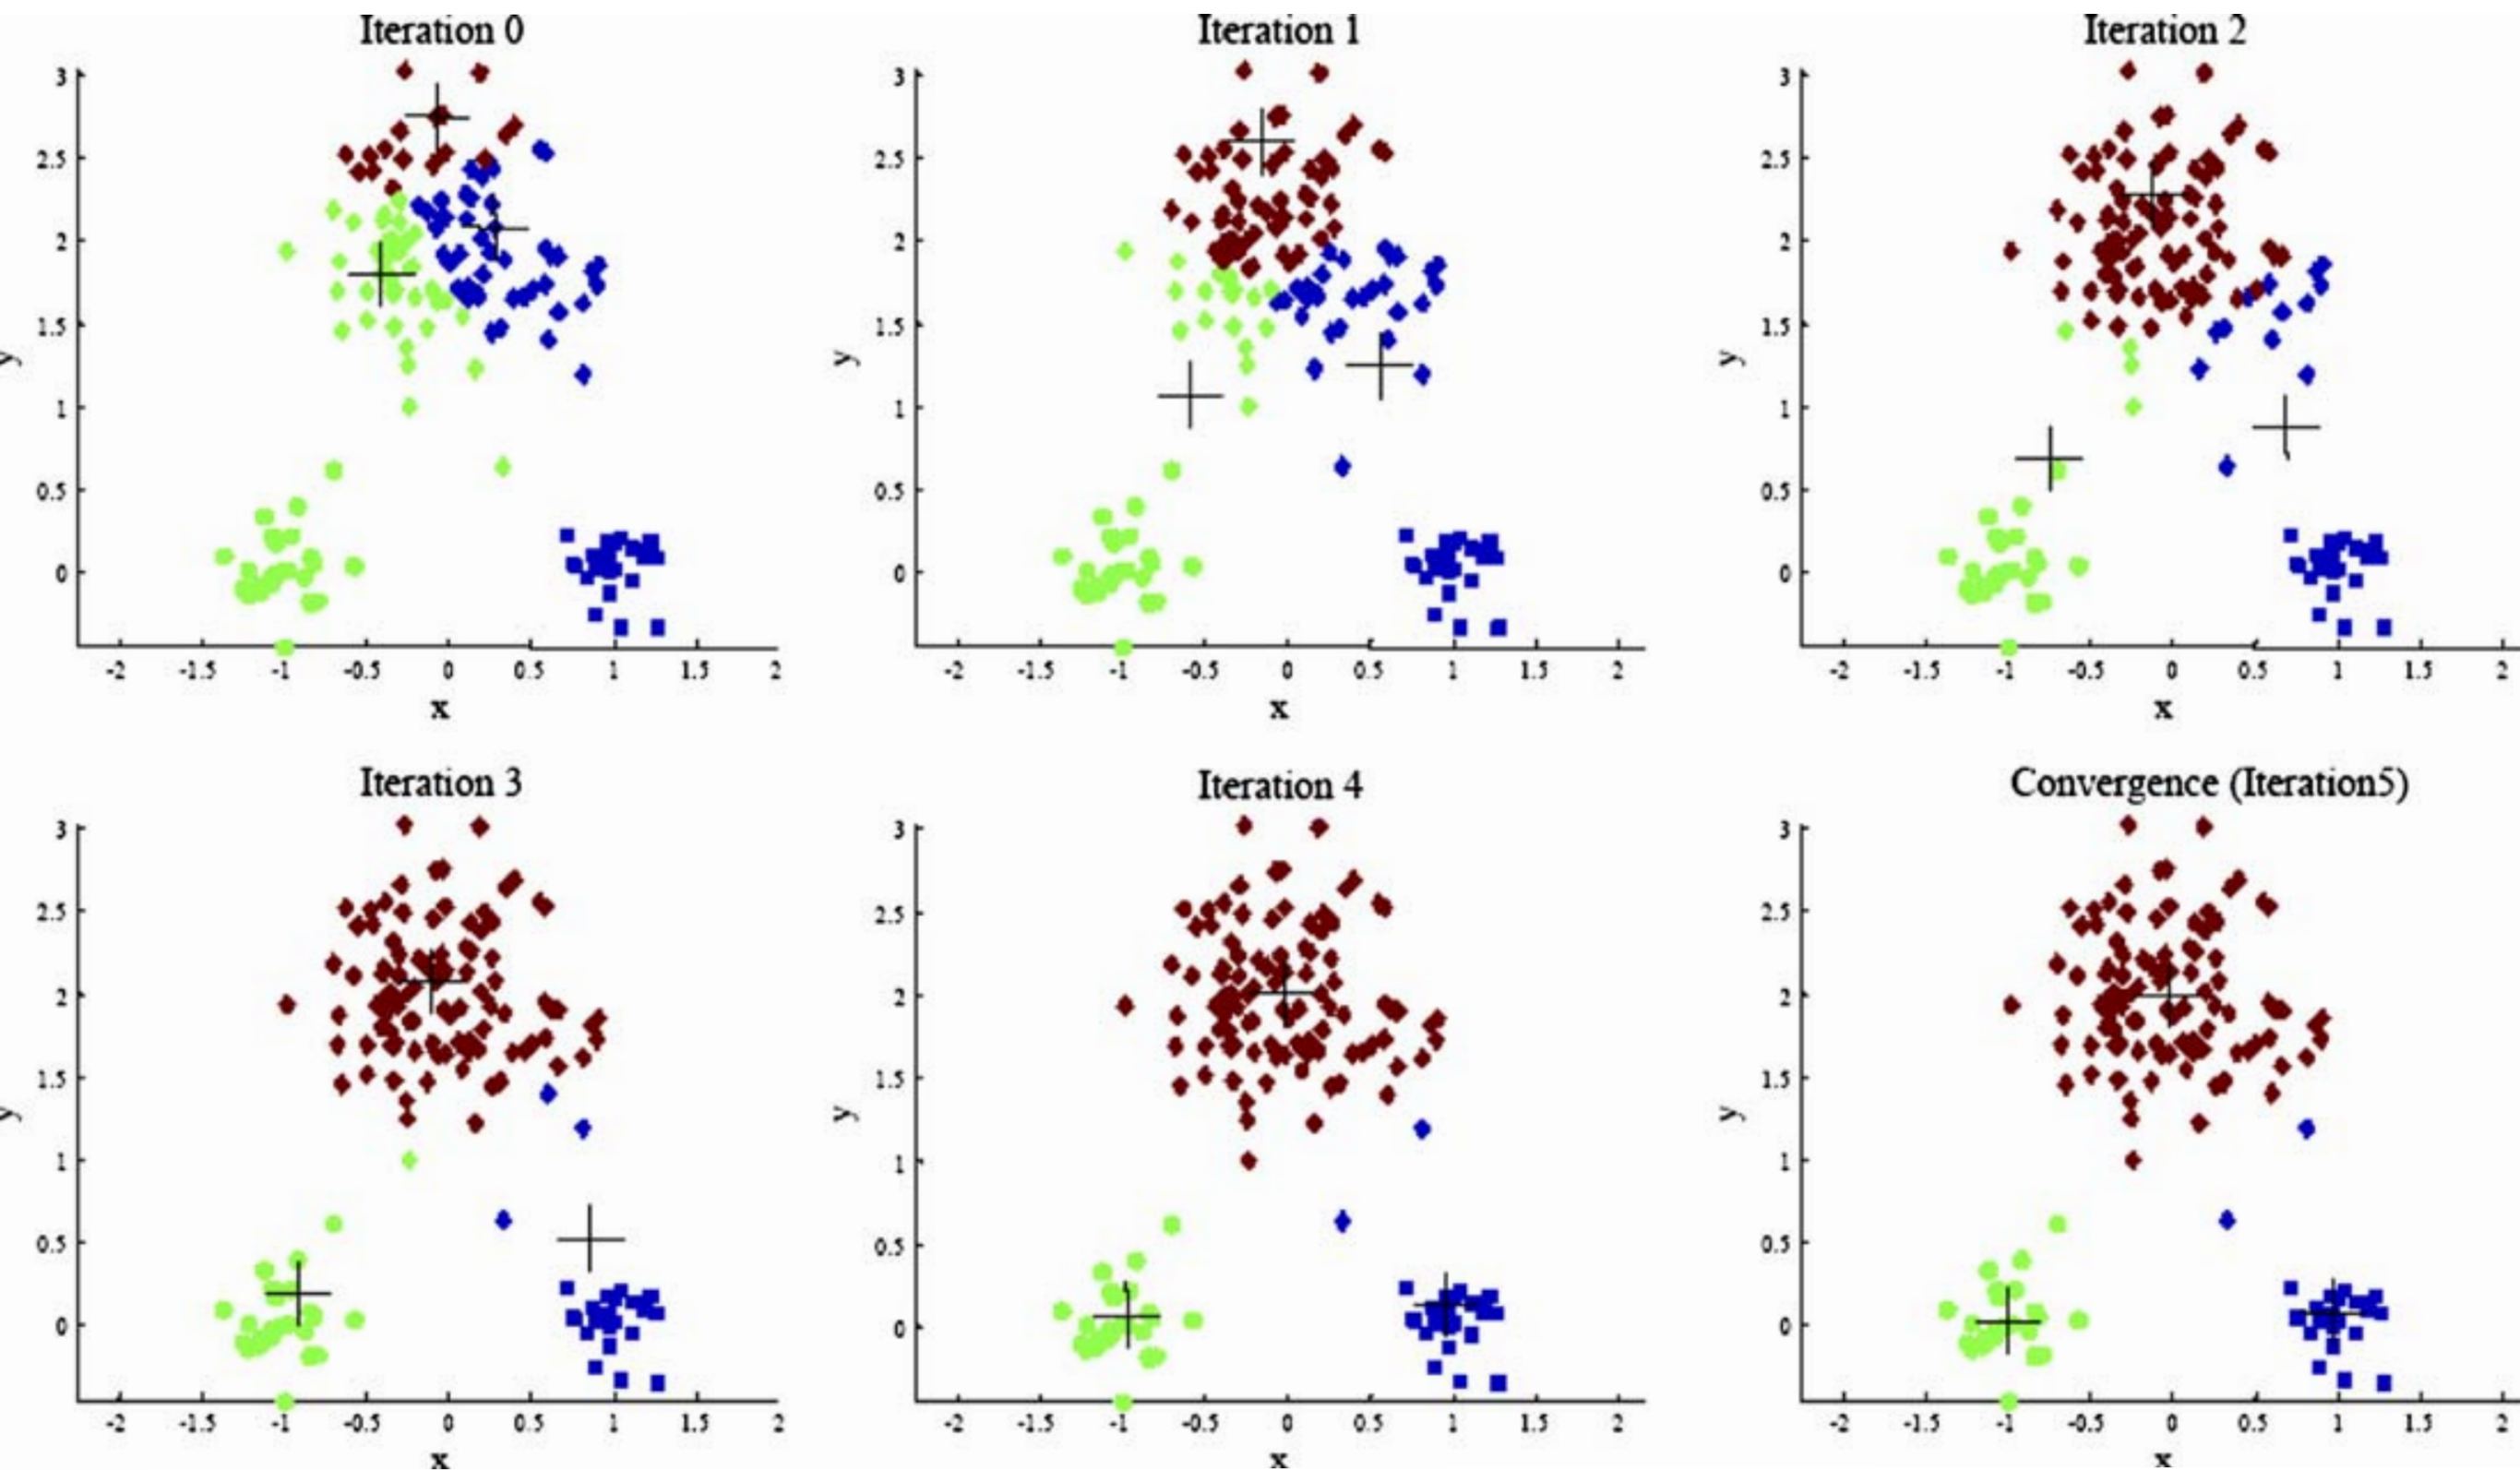
**one way** of solving the **k-means** problem



# the k-means algorithm

1. randomly (or with another method) pick  $k$  cluster centers  $\{c_1, \dots, c_k\}$
2. for each  $j$ , set the cluster  $X_j$  to be the set of points in  $X$  that are the closest to center  $c_j$
3. for each  $j$  let  $c_j$  be the center of cluster  $X_j$  (mean of the vectors in  $X_j$ )
4. repeat (go to step 2) until convergence

# sample execution





# properties of the k-means algorithm

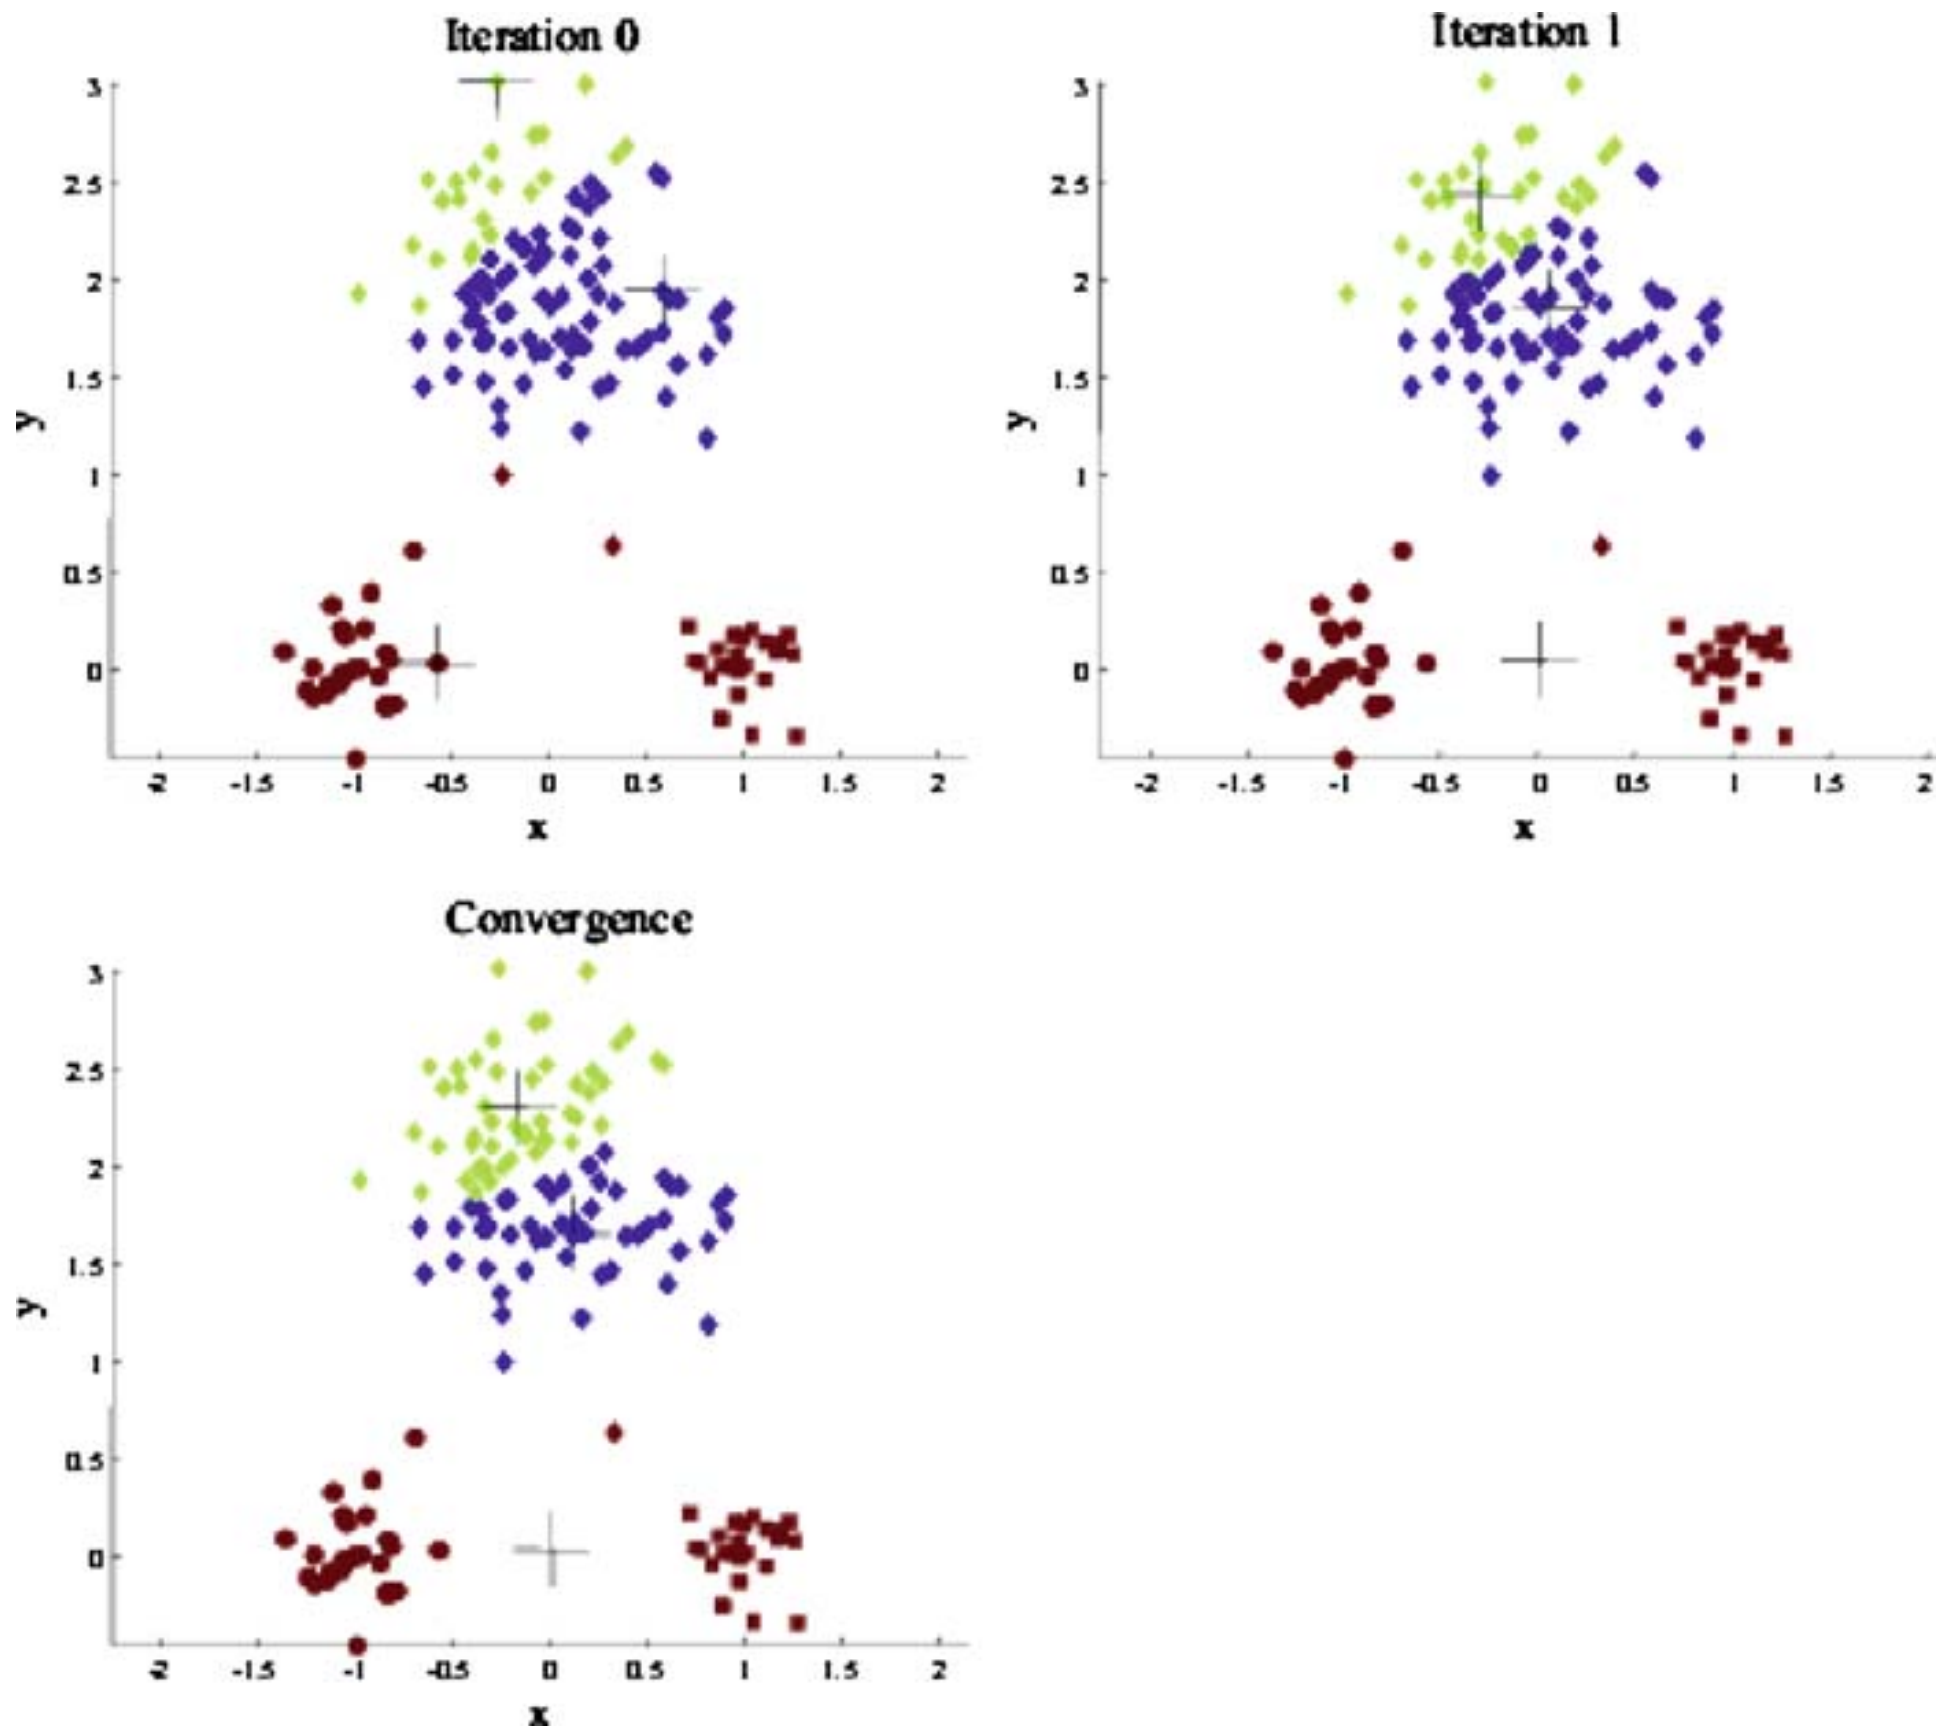
finds a local optimum

often converges quickly

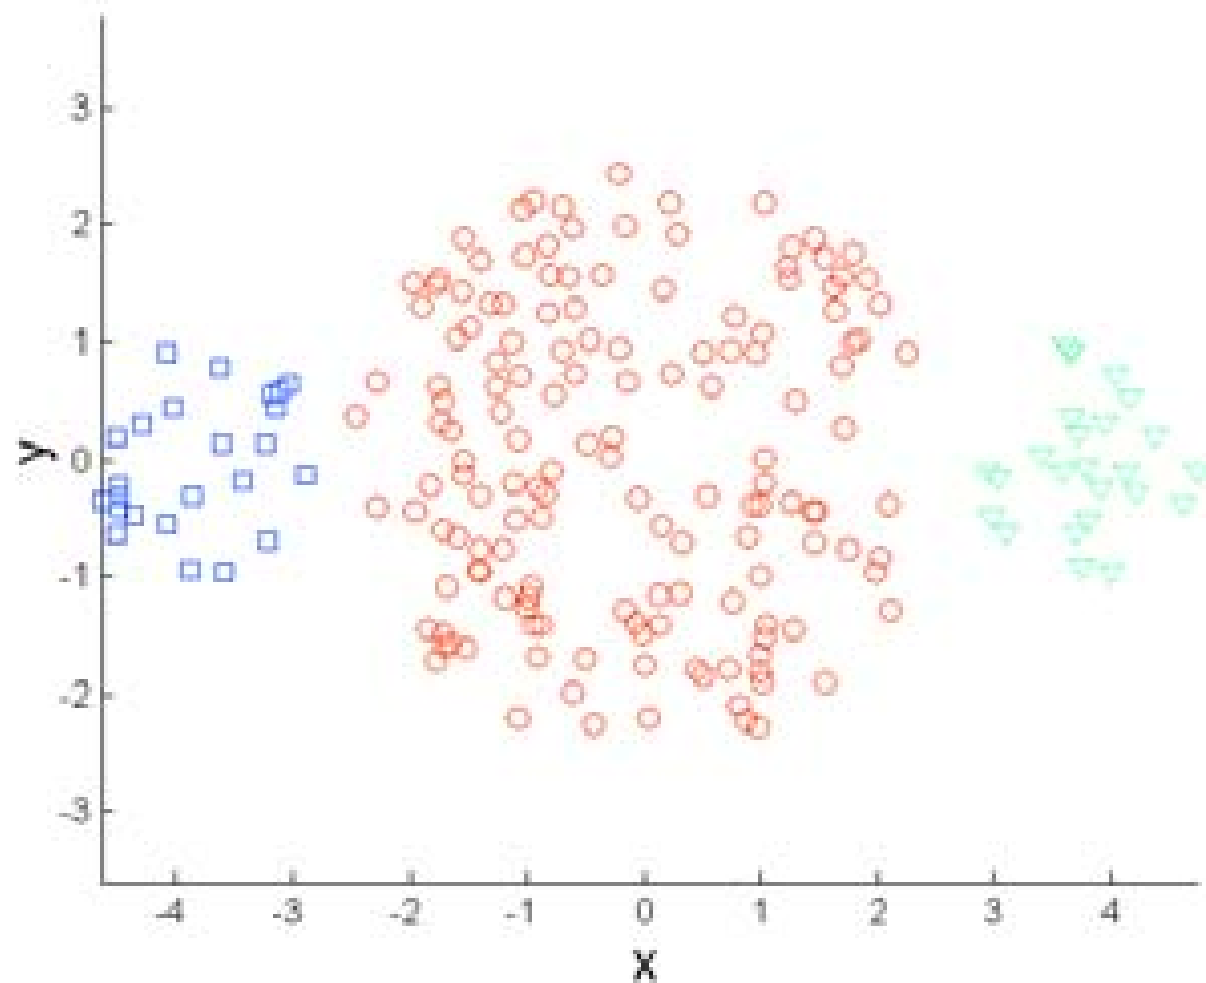
but not always

the choice of initial points can have large influence in the result

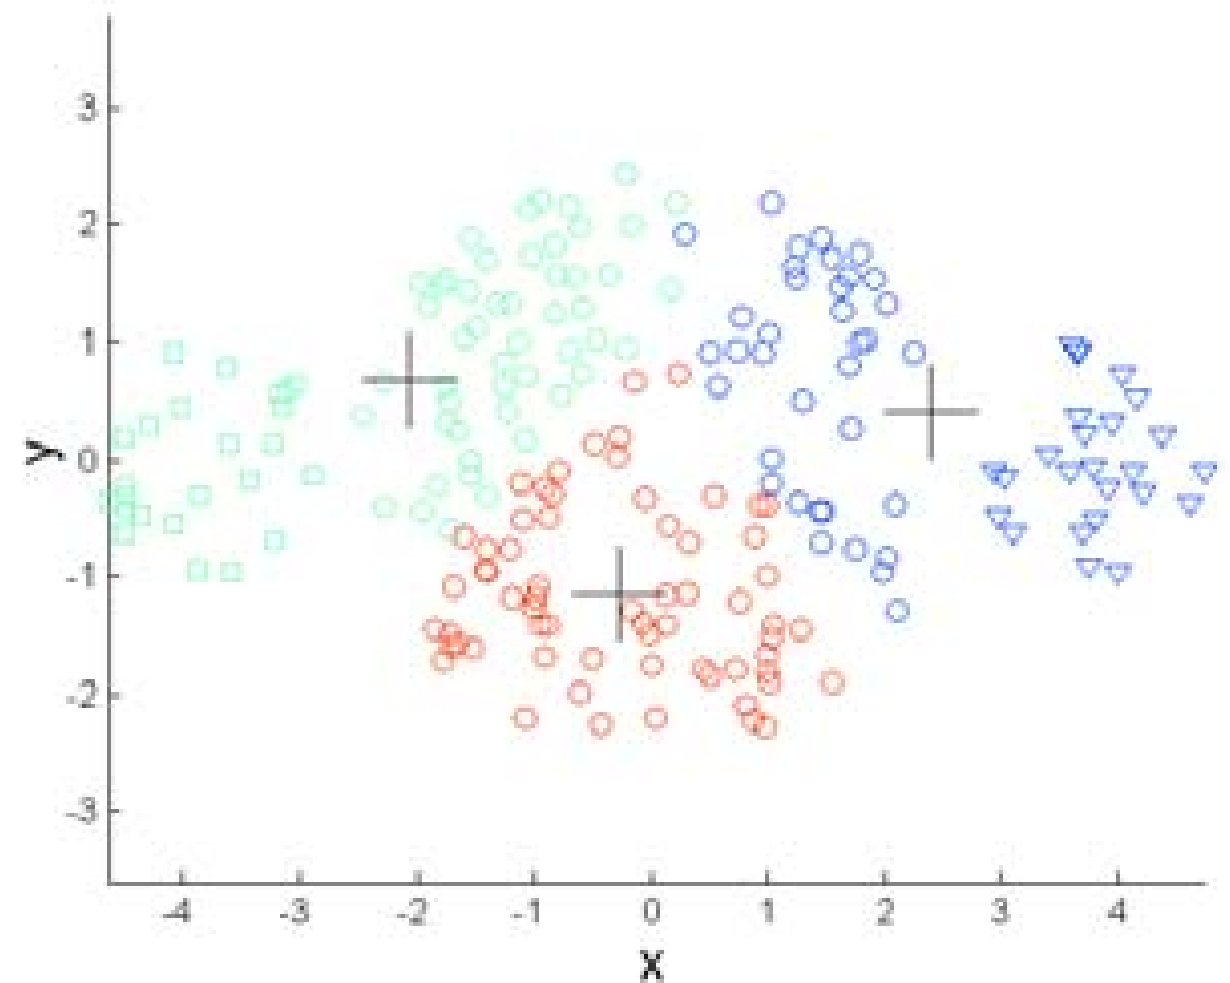
# effects of bad initialization



# limitations of k-means: different sizes

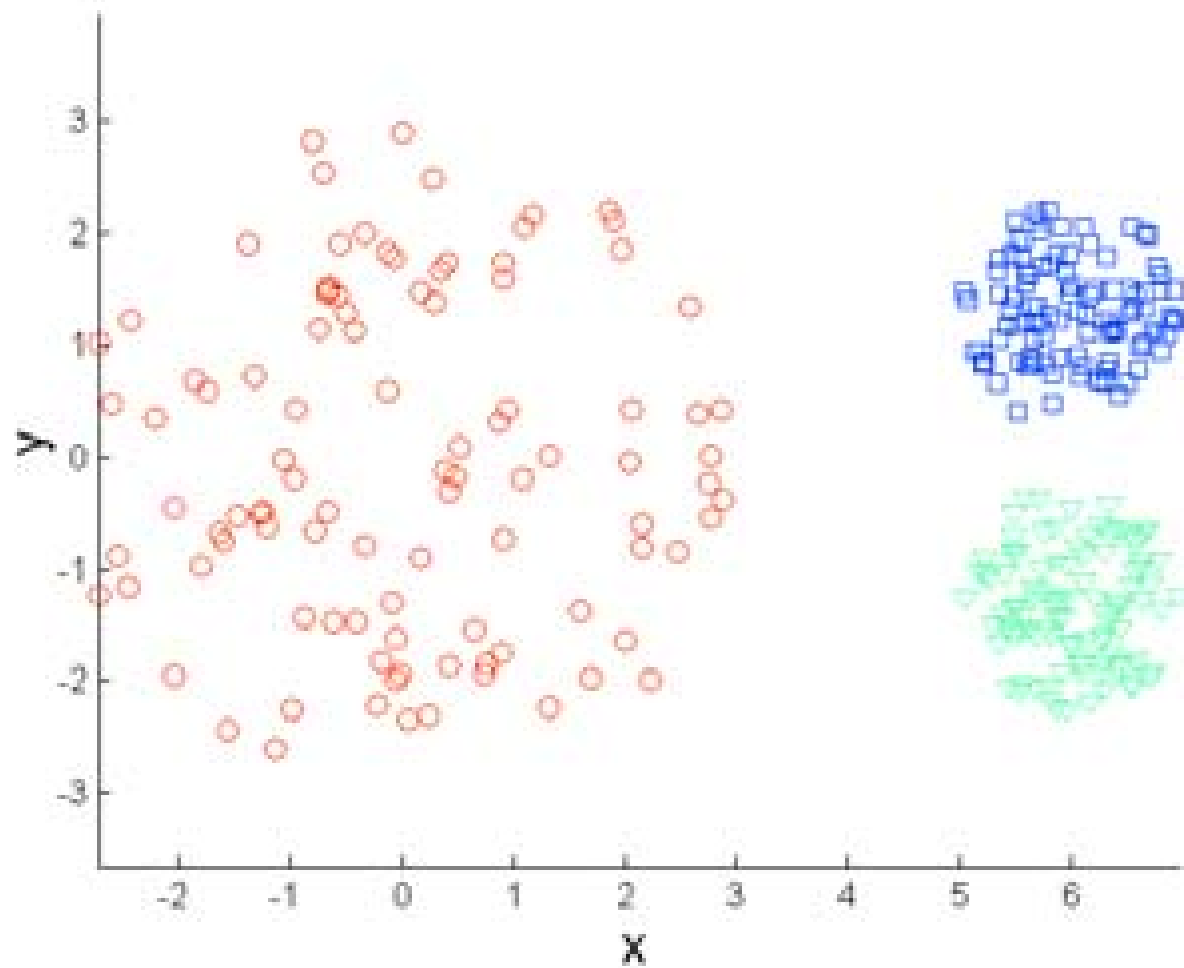


**Original Points**

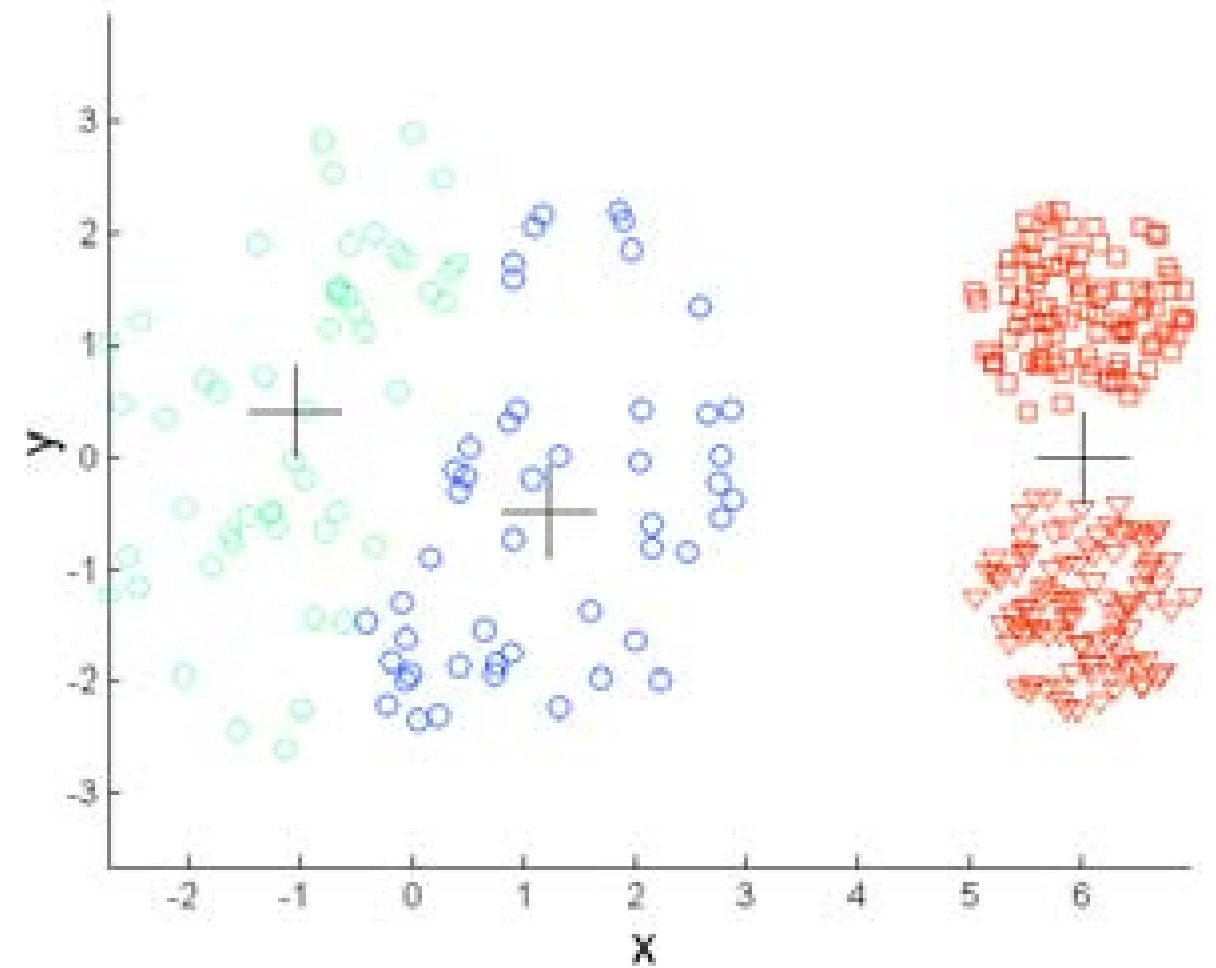


**K-means (3 Clusters)**

# limitations of k-means: different density



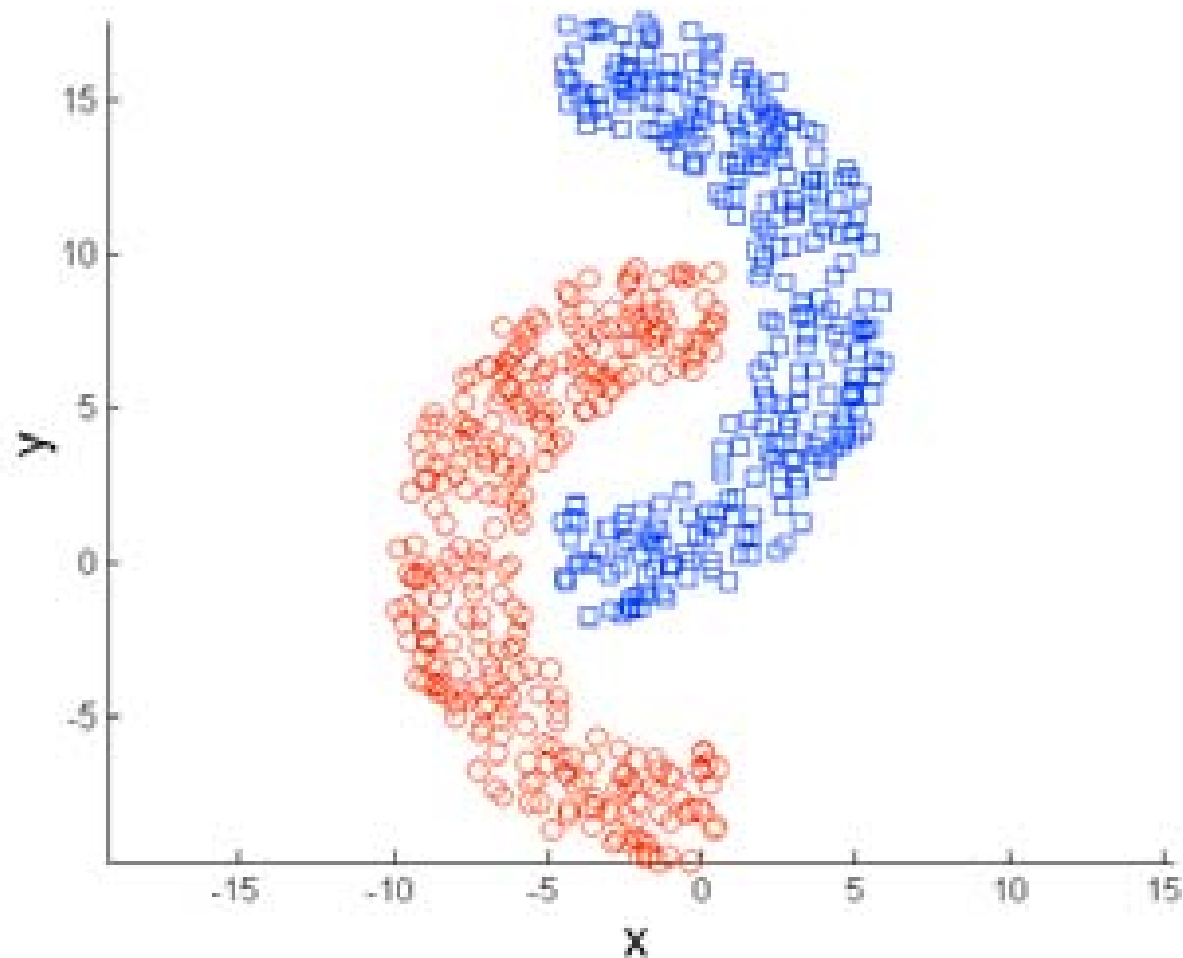
**Original Points**



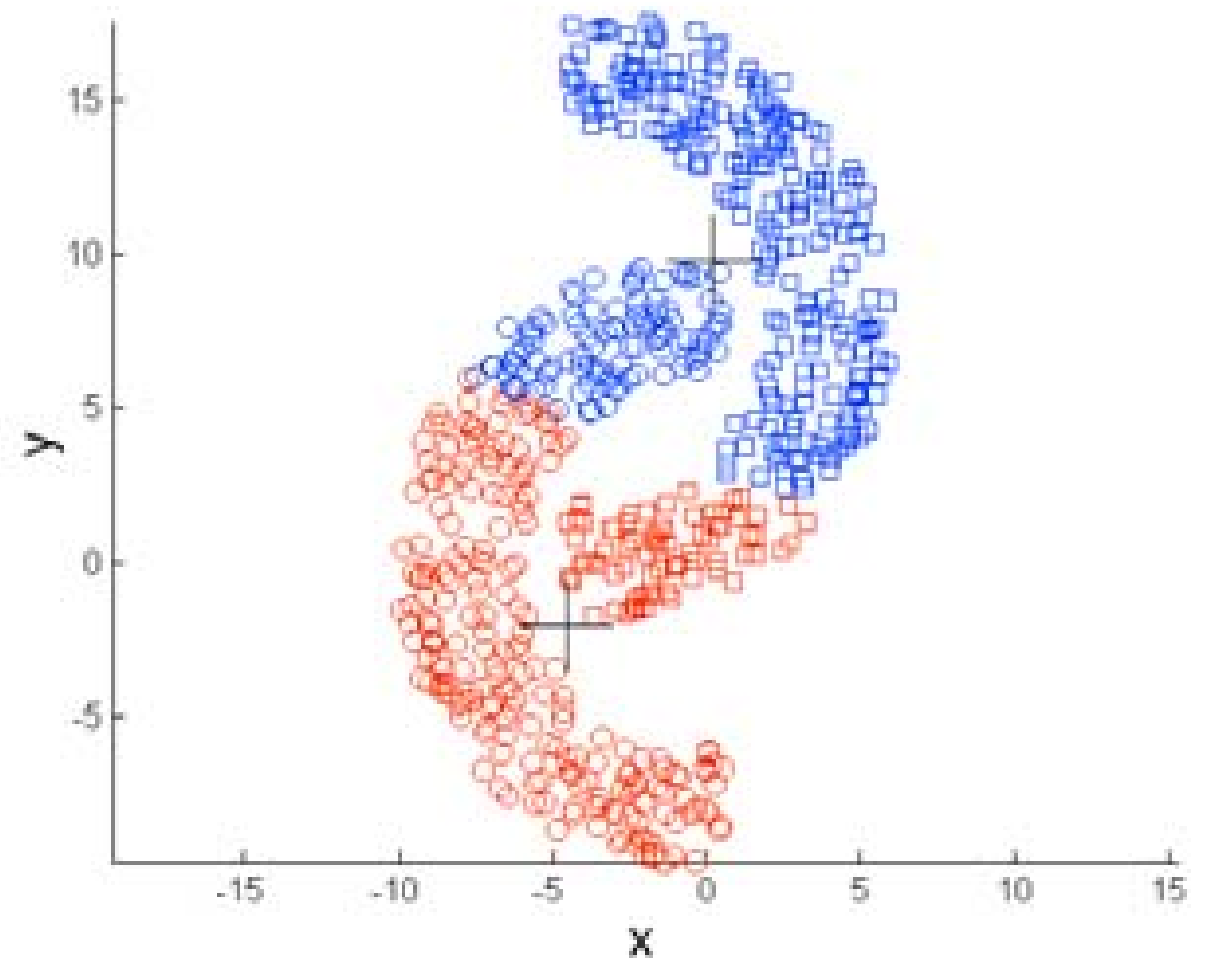
**K-means (3 Clusters)**



# limitations of k-means: non-spherical shapes



**Original Points**



**K-means (2 Clusters)**

# discussion on the k-means algorithm

finds a local optimum

often converges quickly

but not always

the choice of initial points can have large influence in the result

tends to find spherical clusters

outliers can cause a problem

different densities may cause a problem

# initialization

random initialization

repeat many times and take the best solution

helps, but solution can still be bad

pick points that are distant to each other

k-means++

provable guarantees

# generalizations, variants

can we generalize to non Euclidean points ?

yes, as long as we can compute means of clusters...

other problem formulations obtained by modifying the objective function

# variant : the k-median problem

consider set  $X = \{x_1, \dots, x_n\}$  of  $n$  points in  $\mathbb{R}^d$

assume that the number  $k$  is given

problem:

find  $k$  points  $c_1, \dots, c_k$  (named medians)

and partition  $X$  into  $\{X_1, \dots, X_k\}$  by assigning each point  $x_i$  in  $X$  to its nearest cluster median,

so that the cost

$$\sum_{i=1}^n \min_j ||x_i - c_j||_2 = \sum_{j=1}^k \sum_{x \in X_j} ||x - c_j||_2$$

is minimized

# the k-median problem

what about the **1-median** problem for Euclidean points?

also known as **Fermat's problem**

solution to the **1-median** problem (Torricelli point) can be approximated to a given precision by an iterative algorithm

the **general k-median problem** is **NP-hard**

there exist polynomial time **approximation algorithms**,  
assuming that the underlying distance is a metric

# the k-medoids algorithm

or **PAM** (partitioning around medoids)

1. **randomly** (or with another method) choose **k** medoids  $\{c_1, \dots, c_k\}$  from the original dataset **X**
2. assign the remaining **n-k** points in **X** to their **closest medoid**  $c_j$
3. for each cluster, replace each medoid by a point in the cluster that **improves the cost**
4. repeat (go to step 2) until convergence

# discussion on the k-medoids algorithm

k-medoids is a **practical algorithm** (**heuristic**) for solving the k-median problem

**no** approximation guarantee

very similar to the k-means algorithm

same advantages and disadvantages

how about efficiency?

it depends on how efficiently we can solve the 1-median problem for each cluster



# yet another variant : the k-center problem

consider set  $X = \{x_1, \dots, x_n\}$  of  $n$  points in  $\mathbb{R}^d$

assume that the number  $k$  is given

problem:

find  $k$  points  $c_1, \dots, c_k$  (named center)

and partition  $X$  into  $\{X_1, \dots, X_k\}$  by assigning each point  $x_i$  in  $X$  to its nearest cluster center,

so that the cost

$$\max_{i=1}^n \min_{j=1}^k \|x_i - c_j\|_2$$

is minimized

# properties of the k-center problem

NP-hard for dimension  $d \geq 2$

for  $d=1$  the problem is solvable in polynomial time (how?)

a simple combinatorial algorithm works well

# parenthesis...

## approximation algorithms

problem  $P$ :

given input  $I$  find solution  $S^*$  such that  $P(I, S^*)$  is optimized  
(say, minimized)

assume finding  $S^*$  is NP-hard

approximation algorithm  $A$ :

given input  $I$  of size  $n$ , finds  $A(I)$  in polynomial time, s.t.

$$P(I, A(I)) \leq f(n) P(I, S^*)$$

for all inputs  $I$

# approximation algorithms

given input  $I$  of size  $n$ , find  $A(I)$  s.t.

$$P(I, A(I)) \leq f(n) P(I, S^*)$$

$P(I, S^*)$ : value of the objective function for solution  $S$  on input  $I$

$f(n)$  : approximation factor (or approximation guarantee)

approximation scheme (arbitrarily close to  $I$ )  $I + \epsilon$

constant (independent of  $n$ )  $1.5, 2, 3, \dots$

logarithmic,  $f(n) = \log n, \log^2 n, \dots$

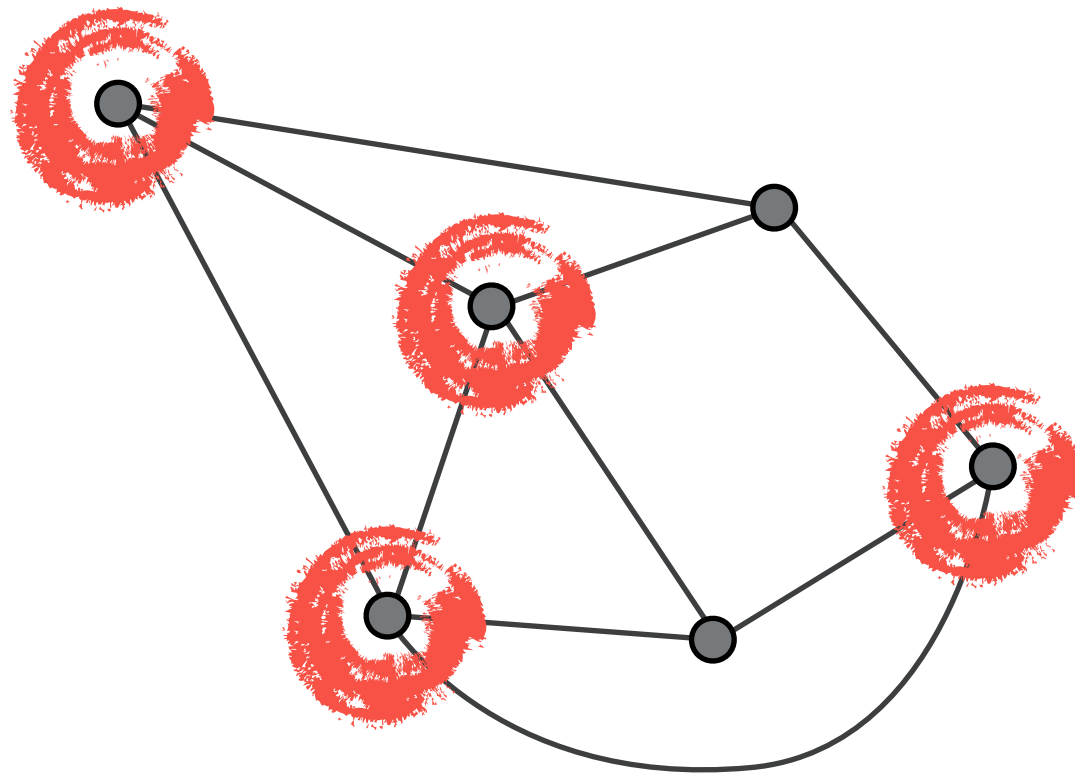
other,  $f(n) = \sqrt{n}, \dots$

# example

vertex cover problem

given a graph

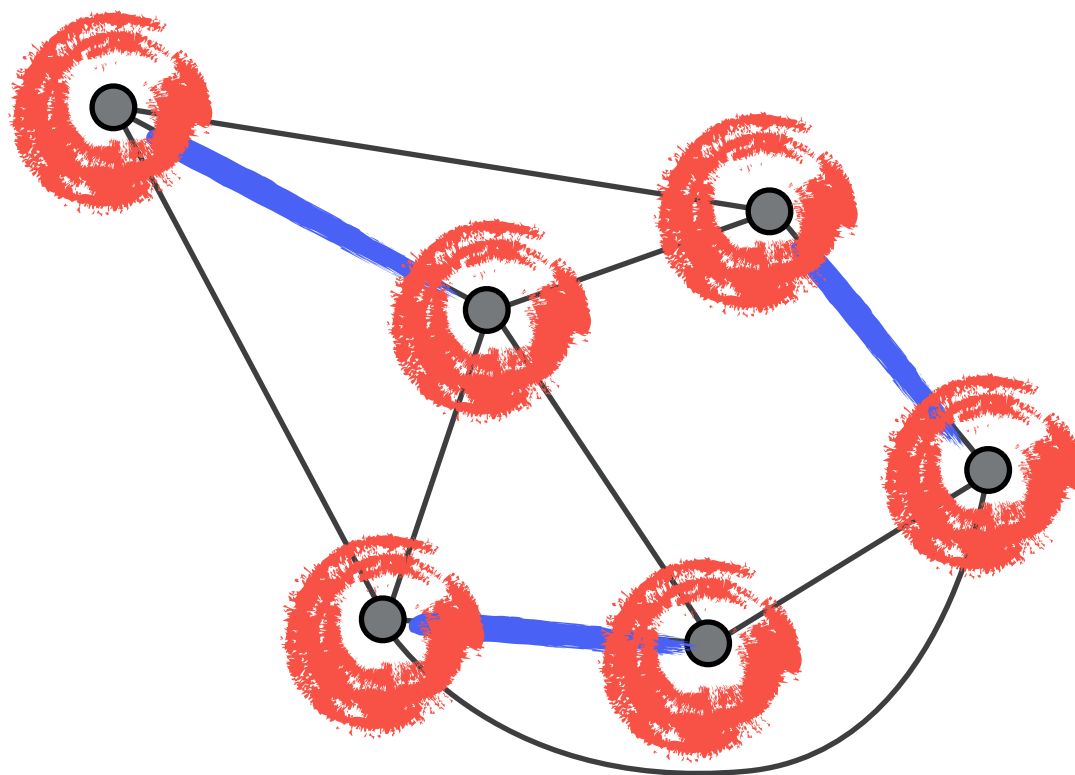
find the smallest set of vertices that cover all the edges



# vertex cover algorithm

find a maximal matching

take **all** the vertices of the matching



2-approximation !

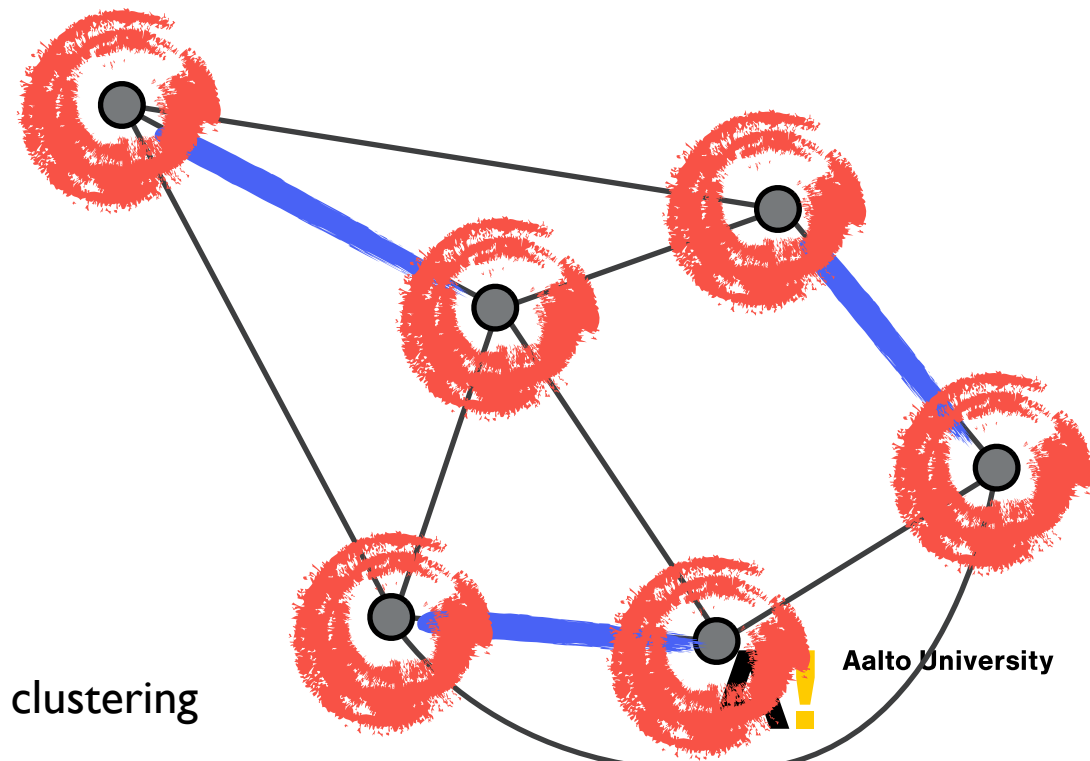
what about greedy?

# analysis

optimal has to cover the edges of the matching

$$(\text{optimal}) \geq (\text{matching size}) = (1/2)(\text{solution of algo})$$

$$(\text{solution of algo}) \leq 2 (\text{optimal})$$



...parenthesis



# recall : the k-center problem

consider set  $X = \{x_1, \dots, x_n\}$  of  $n$  points in  $\mathbb{R}^d$

assume that the number  $k$  is given

problem:

find  $k$  points  $c_1, \dots, c_k$  (named center)

and partition  $X$  into  $\{X_1, \dots, X_k\}$  by assigning each point  $x_i$  in  $X$  to its nearest cluster center,

so that the cost

$$\max_{i=1}^n \min_{j=1}^k \|x_i - c_j\|_2$$

is minimized

# furthest-first traversal algorithm

pick any data point and label it 1

for  $i=2,\dots,k$

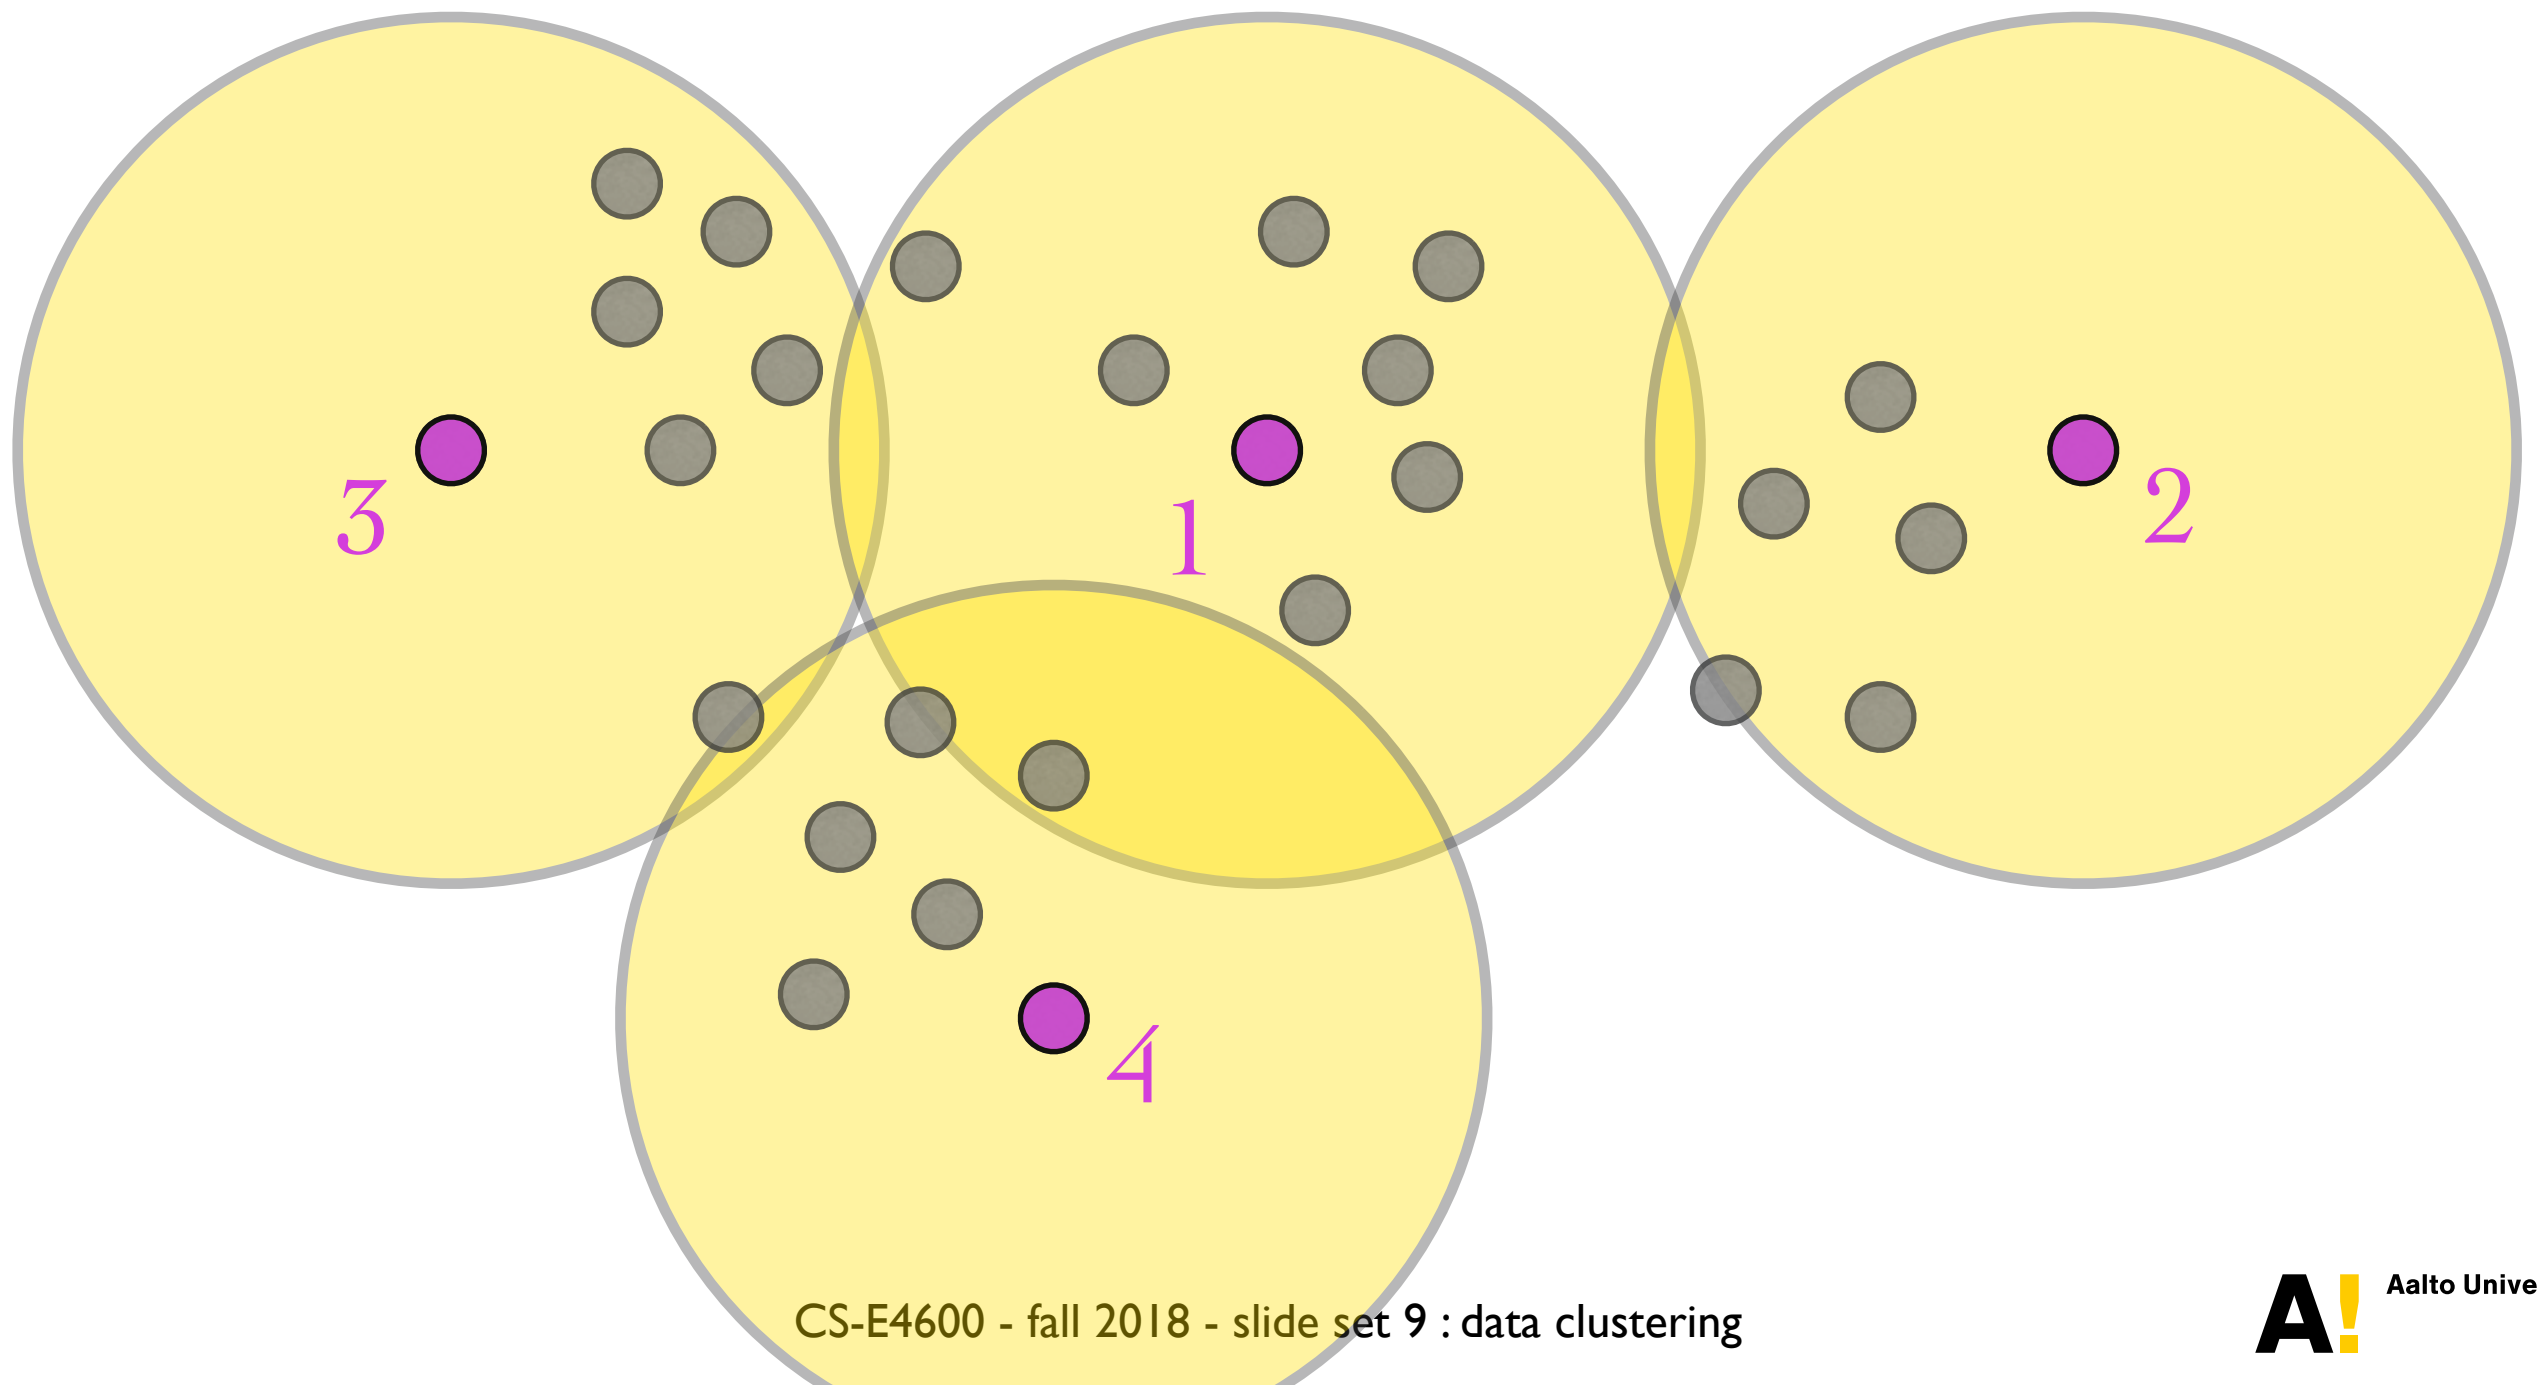
find the unlabeled point that is furthest from  $\{1,2,\dots,i-1\}$

// use  $d(x,S) = \min_{y \in S} d(x,y)$

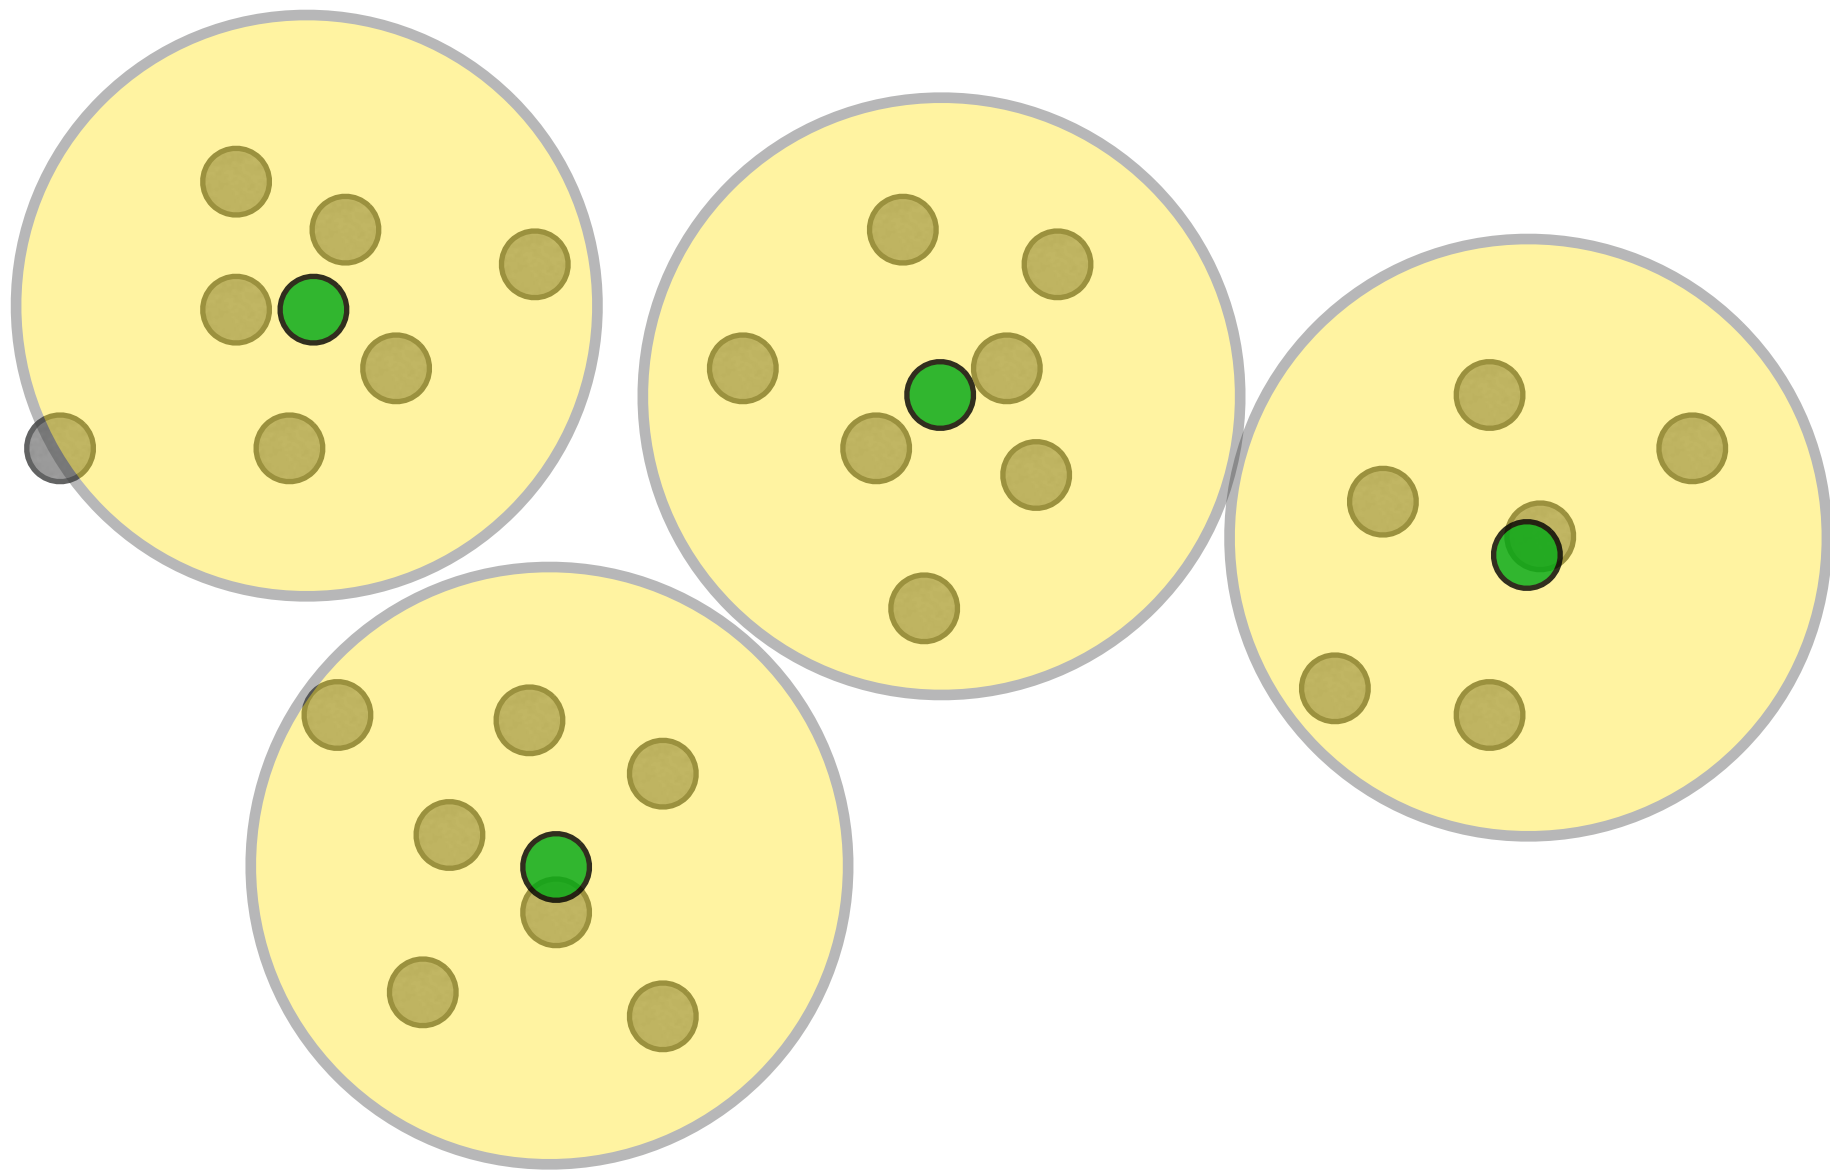
label that point  $i$

assign the remaining unlabeled data points to the closest labeled data point

# furthest-first traversal algorithm: example



# furthest-first traversal algorithm: example



# furthest-first traversal algorithm

furthest-first traversal algorithm gives a factor 2 approximation

# furthest-first traversal algorithm

pick any data point and label it 1

for  $i=2,\dots,k$

find the unlabeled point that is furthest from  $\{1,2,\dots,i-1\}$

// use  $d(x,S) = \min_{y \in S} d(x,y)$

label that point  $i$

$p(i) = \operatorname{argmin}_{j < i} d(i,j)$

$R_i = d(i,p(i))$

assign the remaining unlabeled data points to the closest labeled data point

# analysis

claim 1:  $R_2 \geq R_3 \geq \dots \geq R_k$

proof :

consider indices  $i$  and  $j$  with  $j > i$

$$R_j = d(j, p(j))$$

$$= d(j, \{1, 2, \dots, j-1\})$$

$$= d(j, \{1, 2, \dots, i-1, \dots, j-1\})$$

$$\leq d(j, \{1, 2, \dots, i-1\}) \quad // j > i$$

$$\leq d(i, \{1, 2, \dots, i-1\}) \quad // j \text{ was present when } i \text{ was selected}$$

$$= R_i$$

# analysis

claim 2 :

let  $C$  be the clustering produced by the FFT algorithm

let  $R(C)$  be the cost of that clustering

then  $R(C) = R_{k+1}$

proof :

for any  $i > k$  we have :

$$d(i, \{1, 2, \dots, k\}) \leq d(k+1, \{1, 2, \dots, k\}) = R_{k+1}$$



# analysis

## theorem

let  $C$  be the clustering produced by the FFT algorithm

let  $C^*$  be the optimal clustering

then  $R(C) \leq 2R(C^*)$

## proof:

let  $C_1^*, \dots, C_k^*$  be the clusters of the optimal  $k$ -clustering

if these clusters contain points  $\{1, \dots, k\}$  then

$$R(C) \leq 2R(C^*) \quad \star$$

otherwise suppose that one of these clusters contains two or more of the points in  $\{1, \dots, k\}$

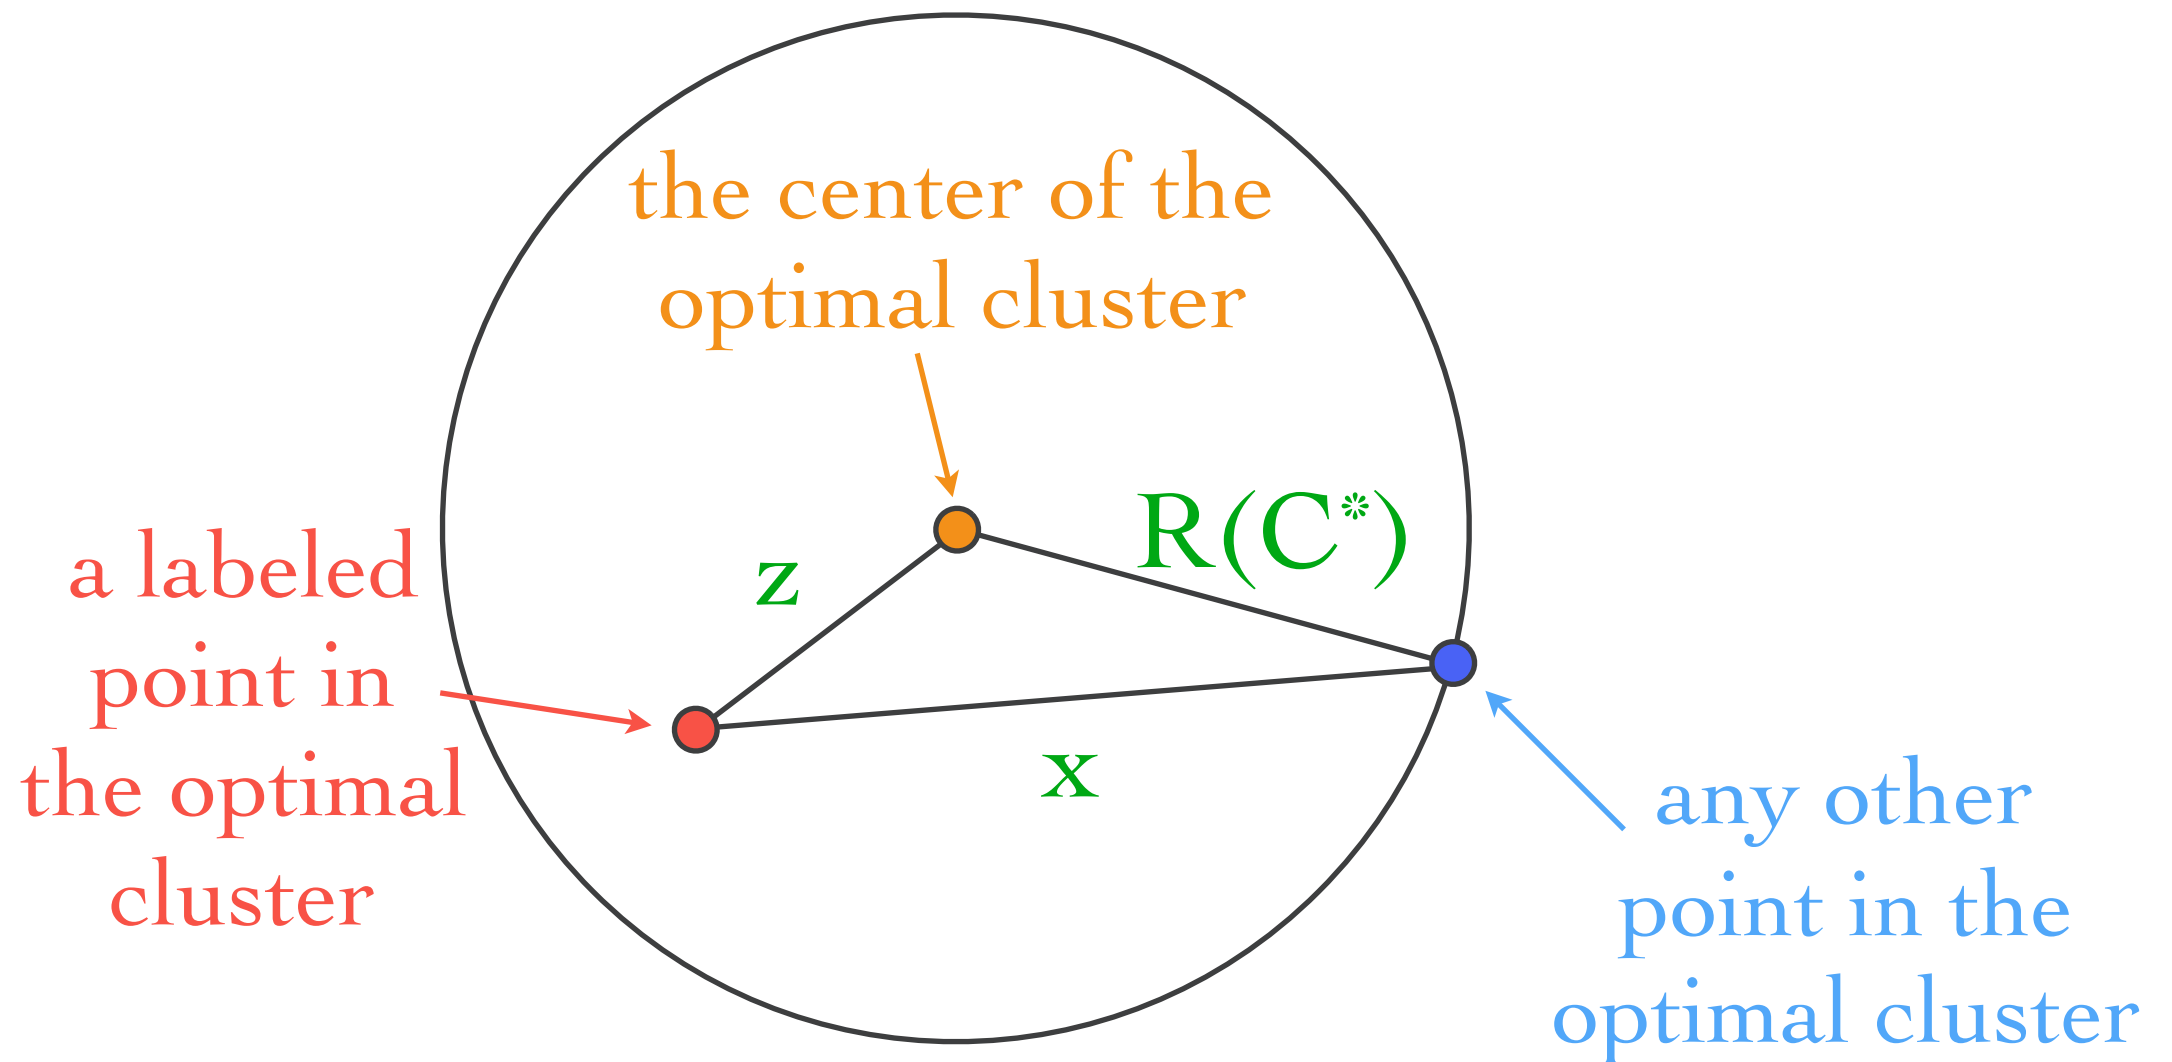
these points are at distance at least  $R_k$  from each other

this (optimal) cluster must have radius at least

$$\frac{1}{2} R_k \geq \frac{1}{2} R_{k+1} = \frac{1}{2} R(C)$$



$$R(C) \leq 2R(C^*)$$



$$R(C) = \max x \leq z + R(C^*) \leq 2R(C^*)$$

k-means++

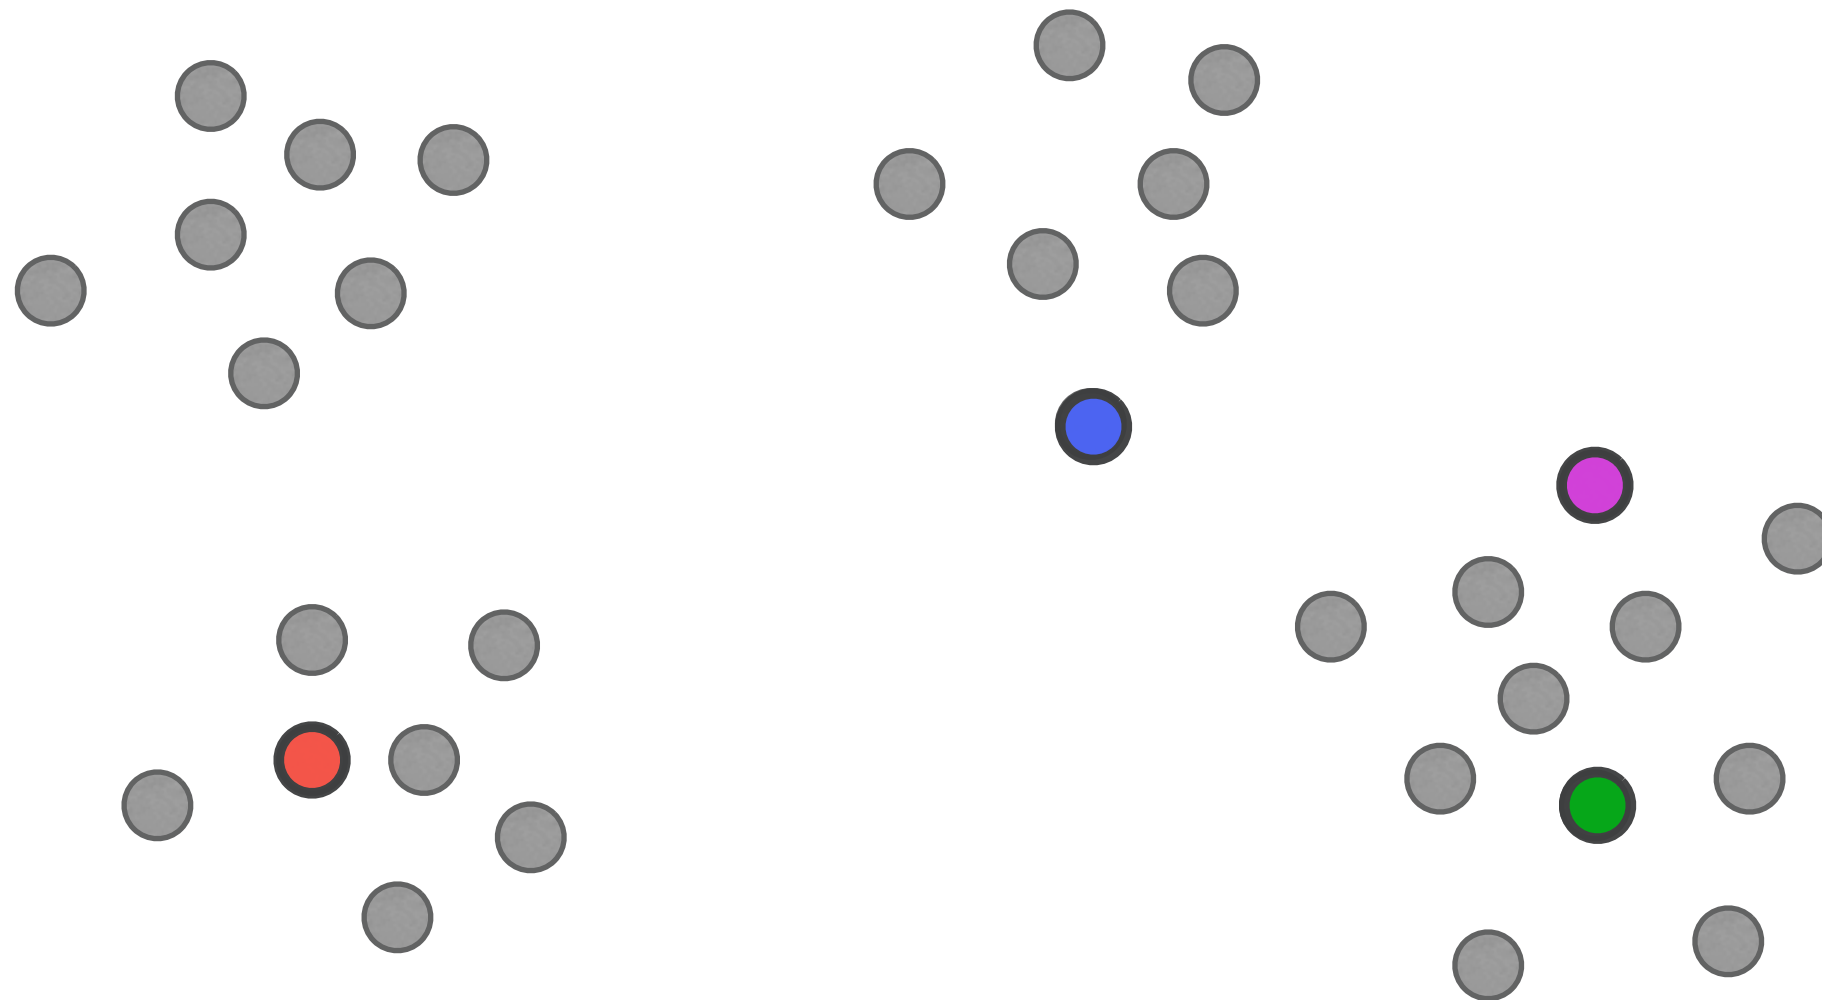
# optional reading assignment

David Arthur and Sergei Vassilvitskii

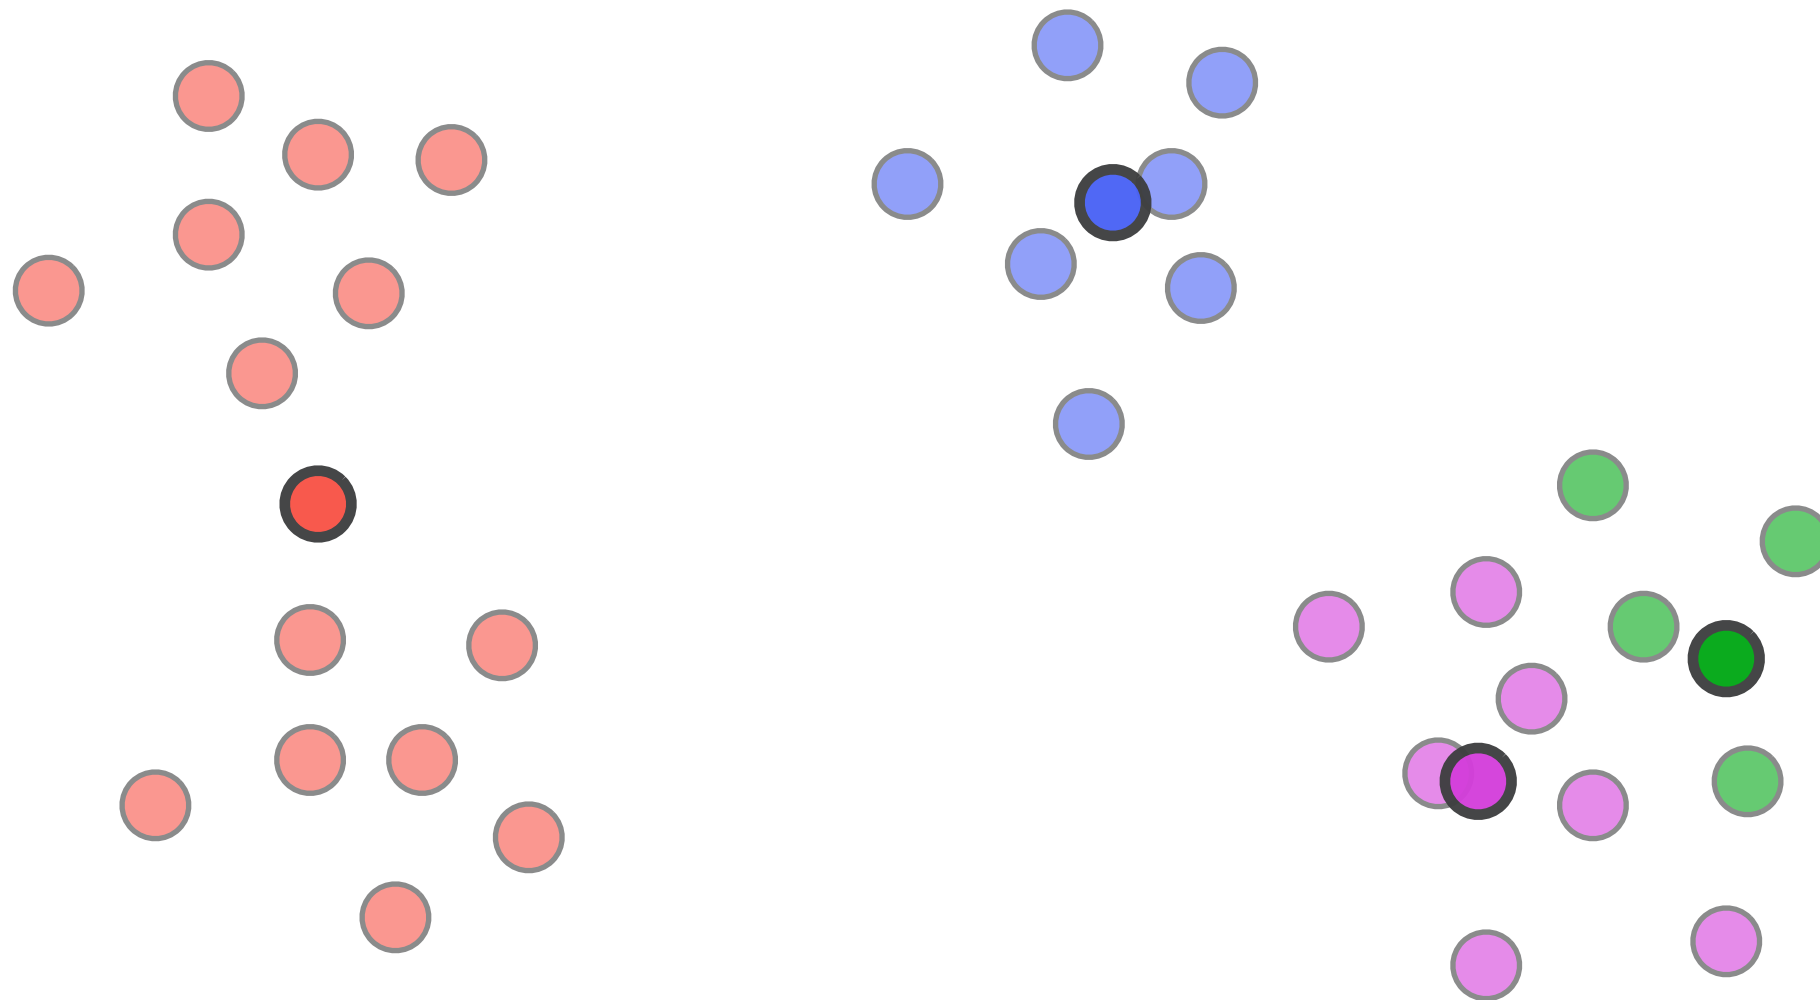
k-means++: The advantages of careful seeding

SODA 2007

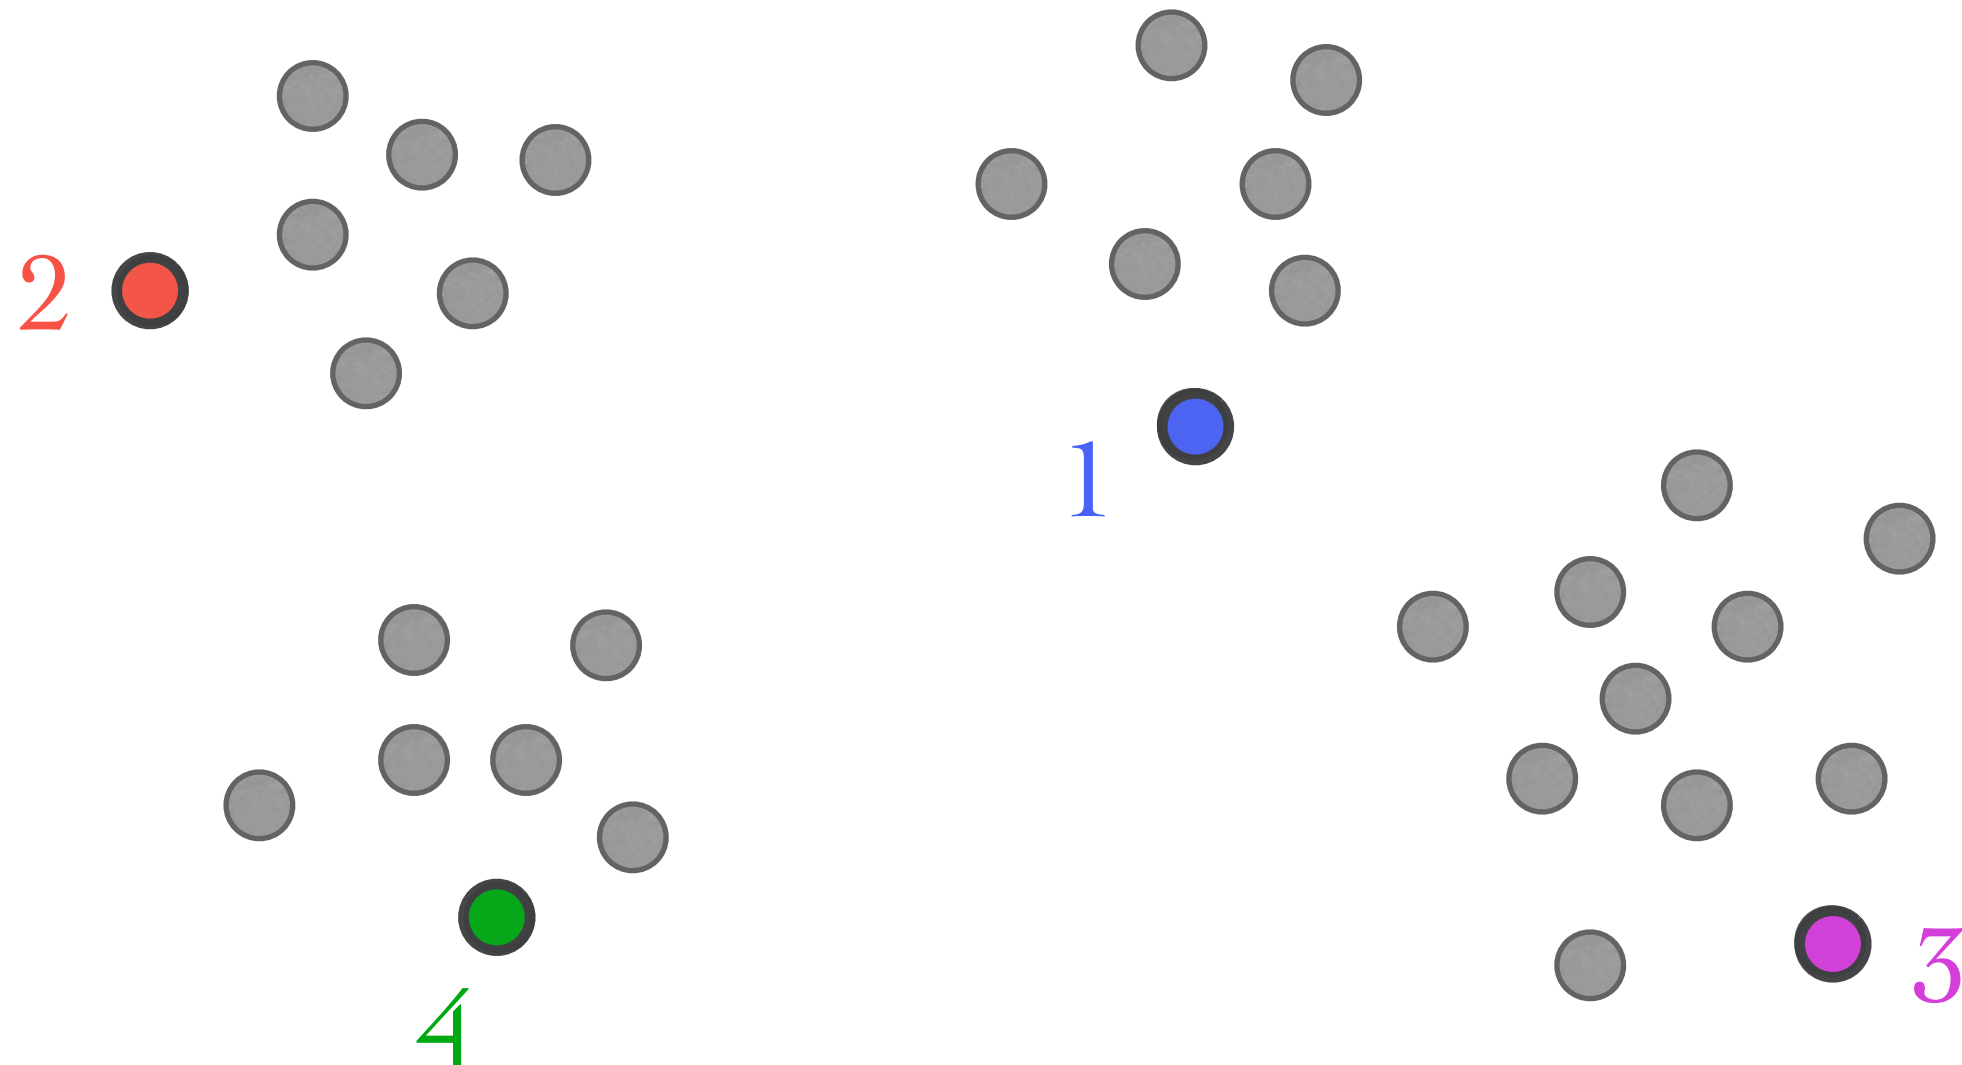
# k-means algorithm: random initialization



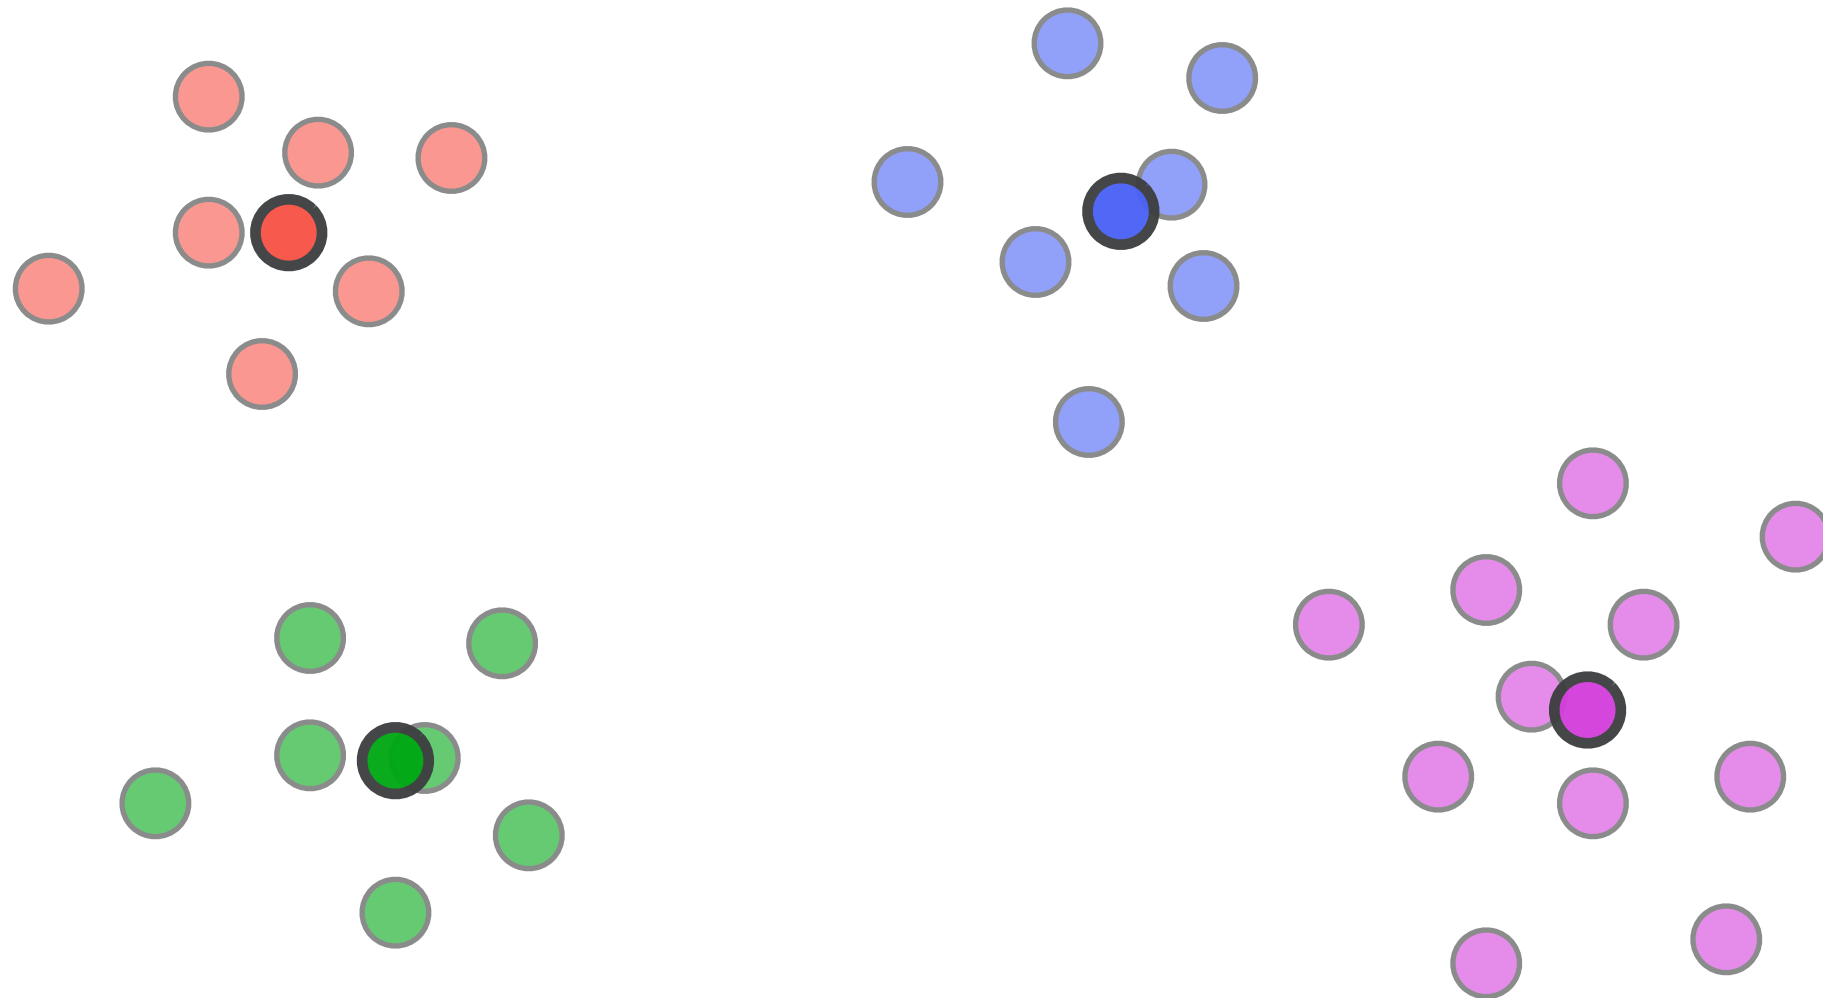
# k-means algorithm: random initialization



# k-means algorithm: initialization with further-first traversal

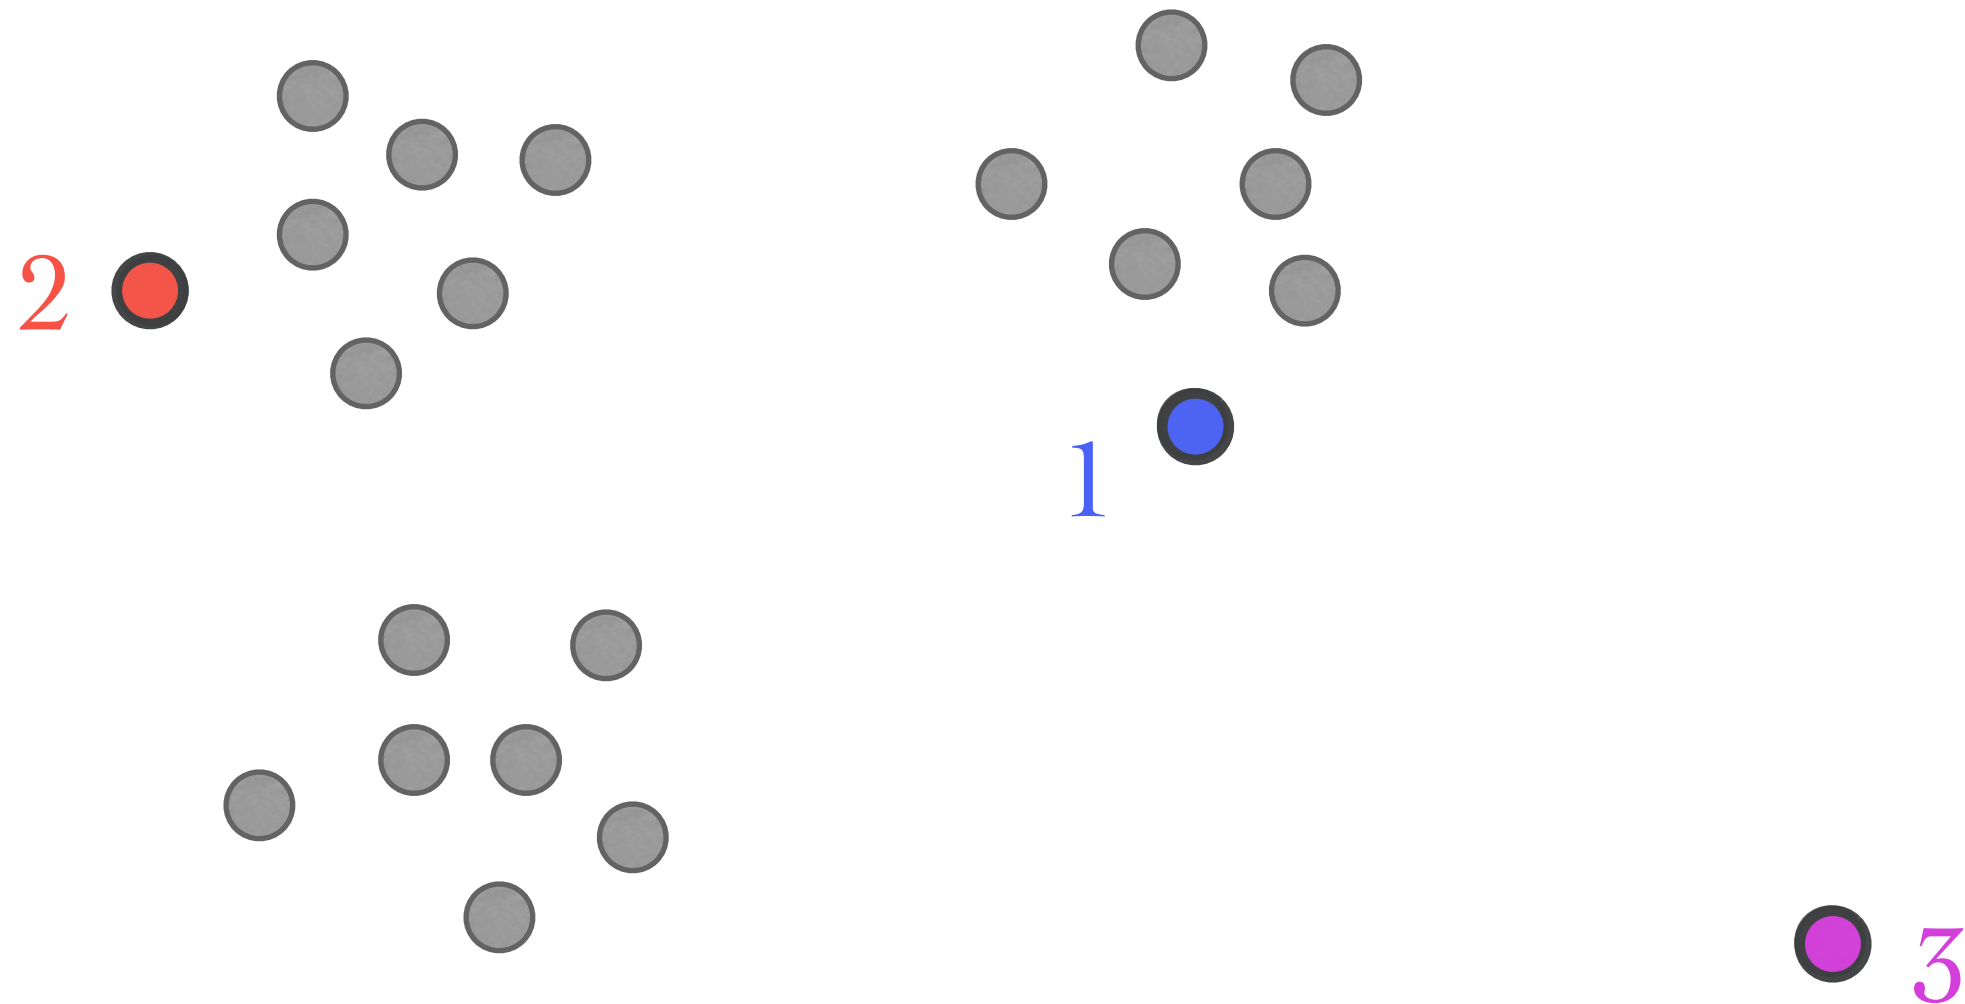


# k-means algorithm: initialization with further-first traversal

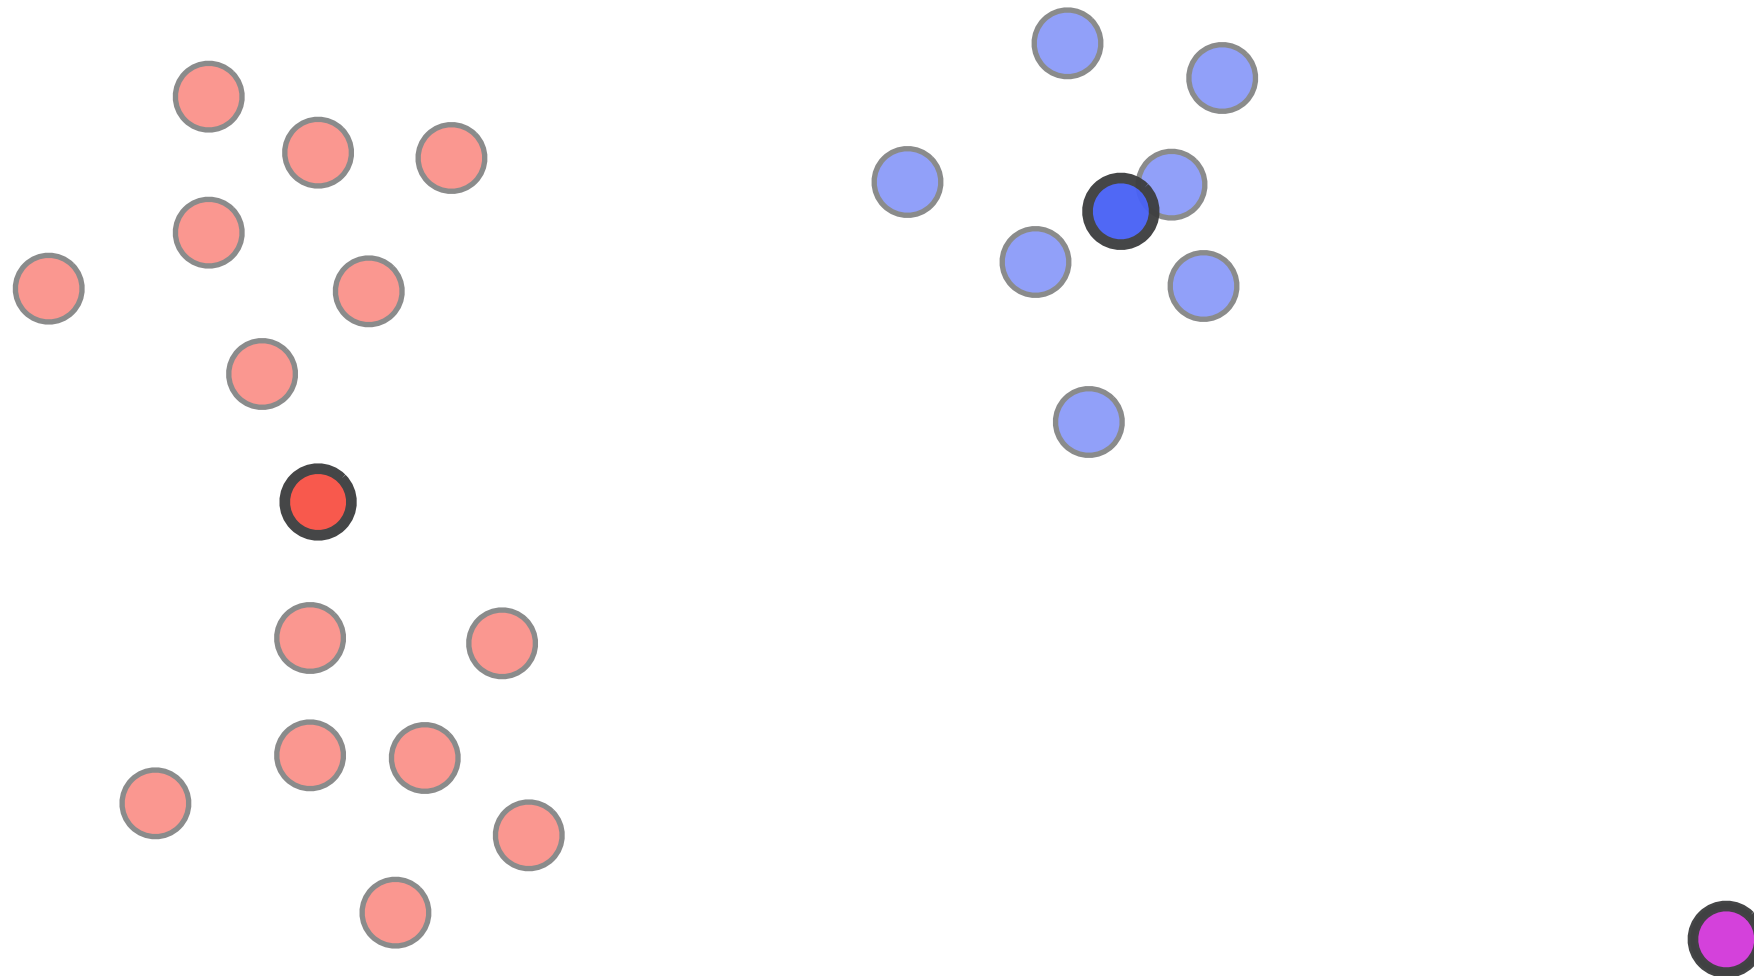




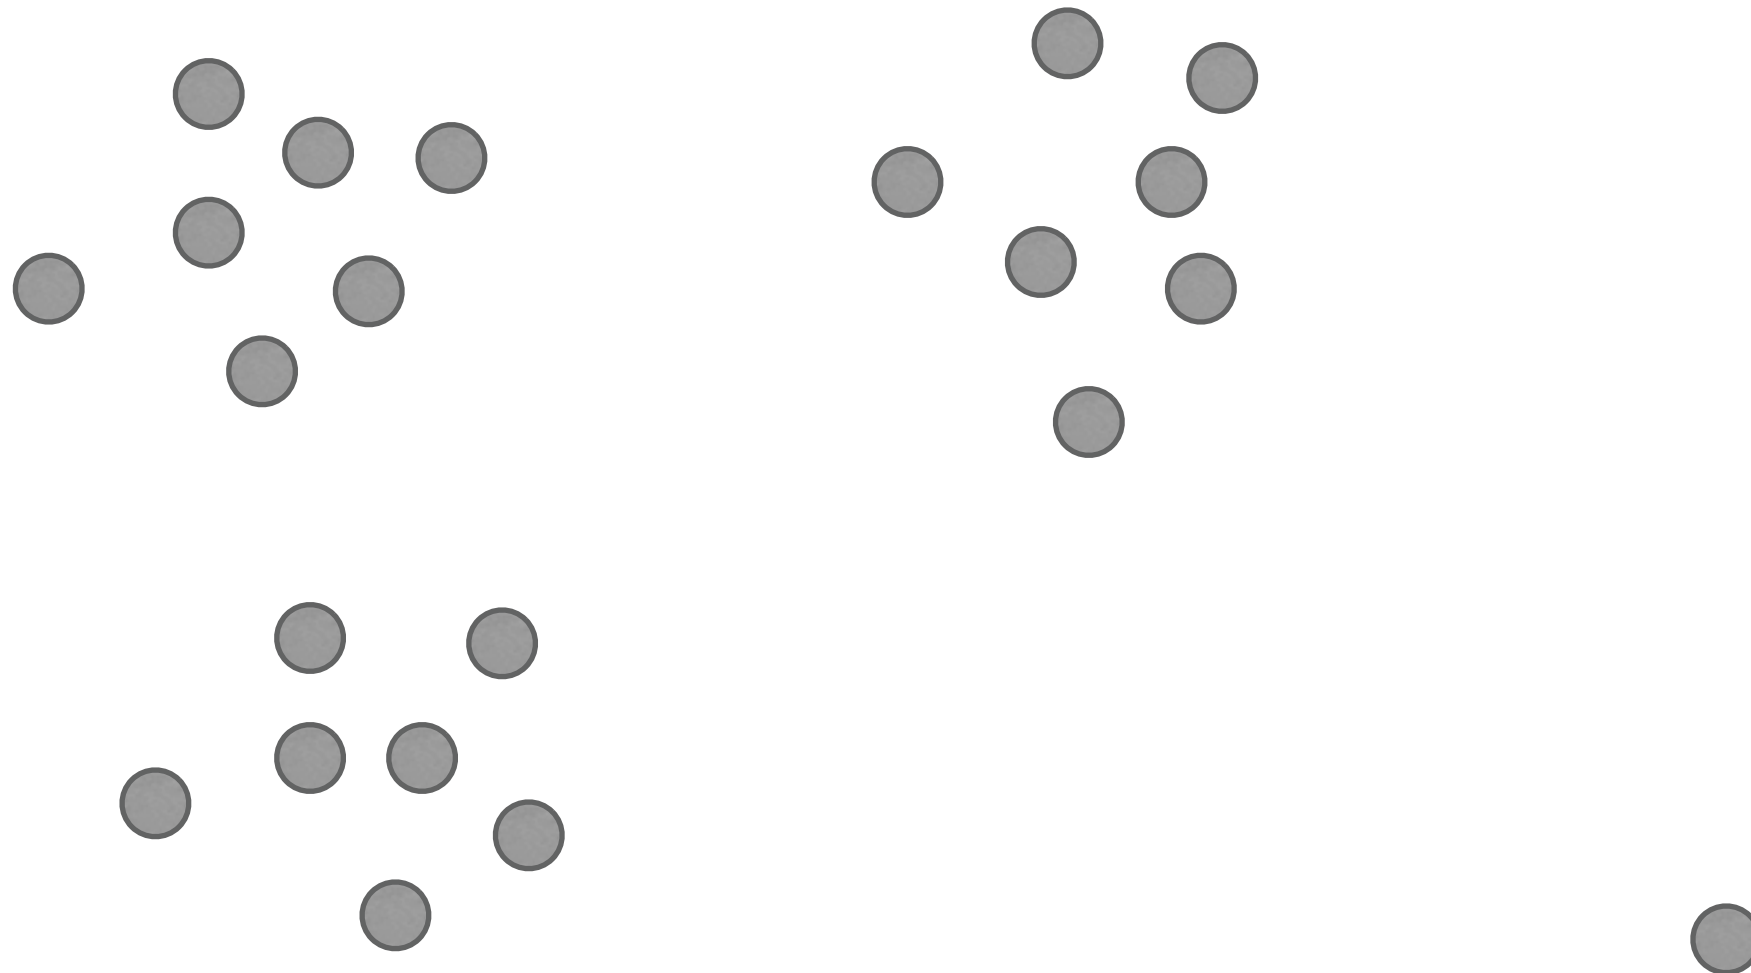
# but sensitive to outliers



# but sensitive to outliers



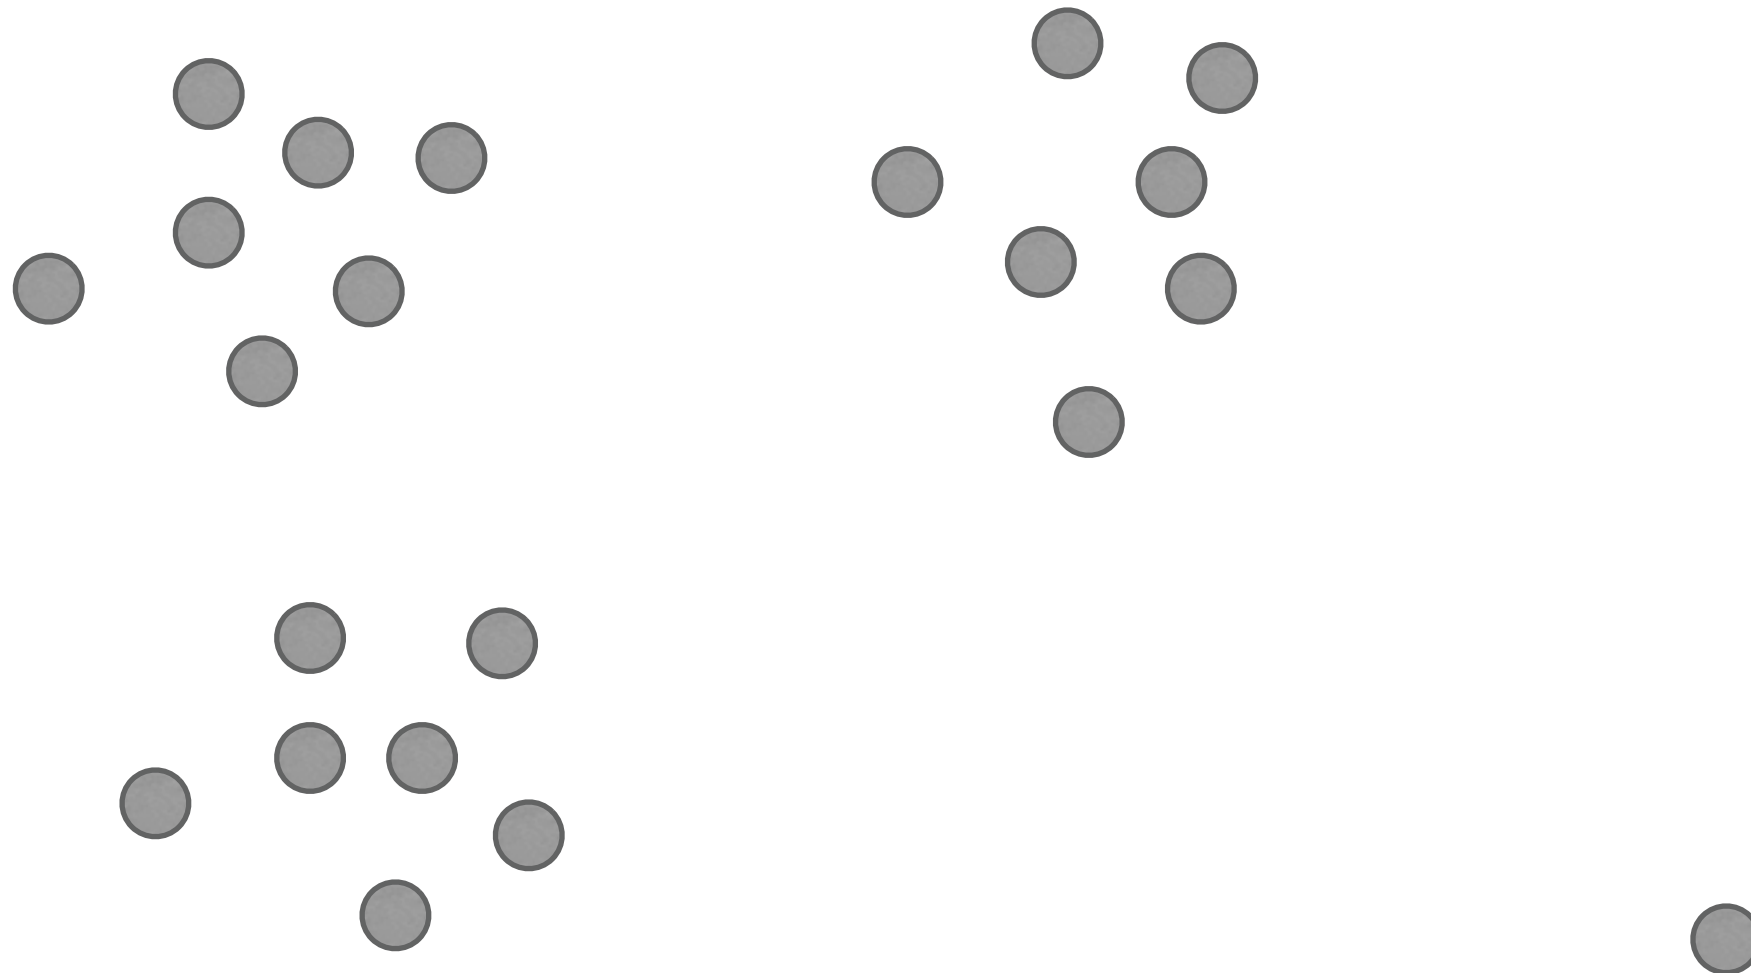
# here random may work well



we want to select seeds that that are :

1. **far** from existing seeds (explore a new area of the data)
2. have **many near-by points** (potentially discover a new cluster)

how do we accomplish both objectives?



# k-means++ algorithm

interpolate between the two methods (furthestst and random)

let  $D(x)$  be the distance between  $x$  and the nearest center selected so far

choose next center with probability proportional to

$$(D(x))^a = D^a(x)$$

$a = 0$       random initialization

$a = \infty$     furthest-first traversal

$a = 2$       k-means++

# k-means++ algorithm

initialization phase :

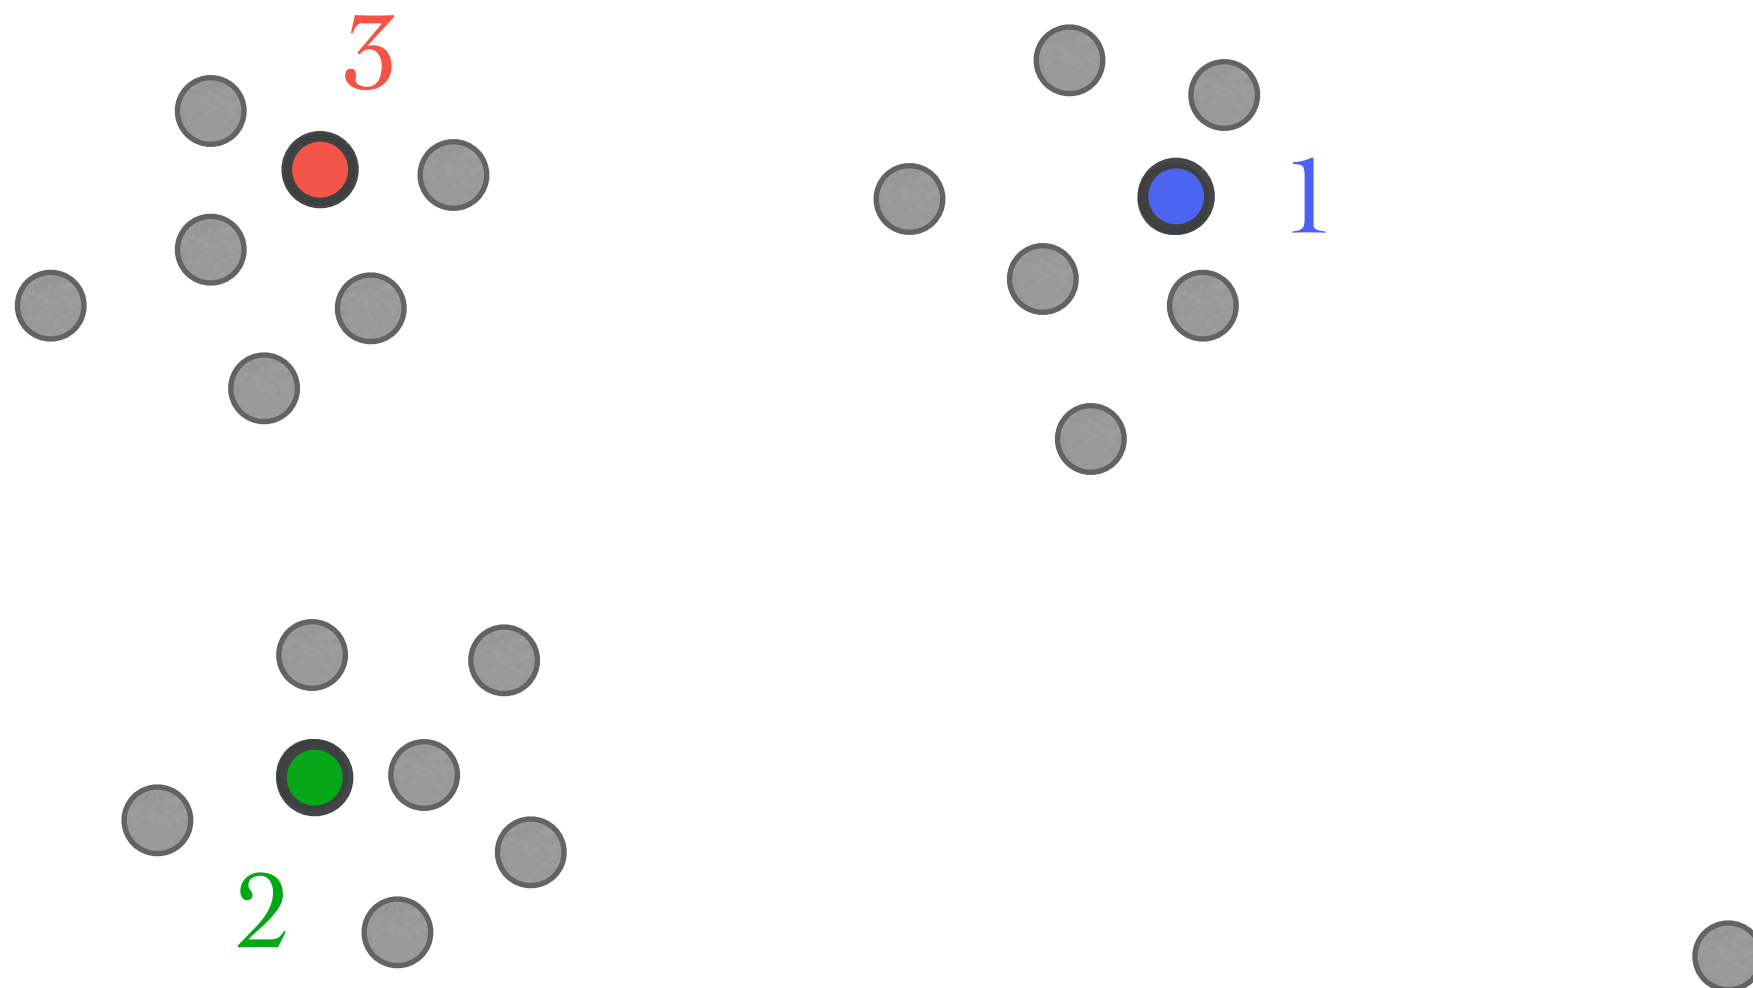
choose the first center uniformly at random

choose next center with probability proportional to  $D^2(x)$

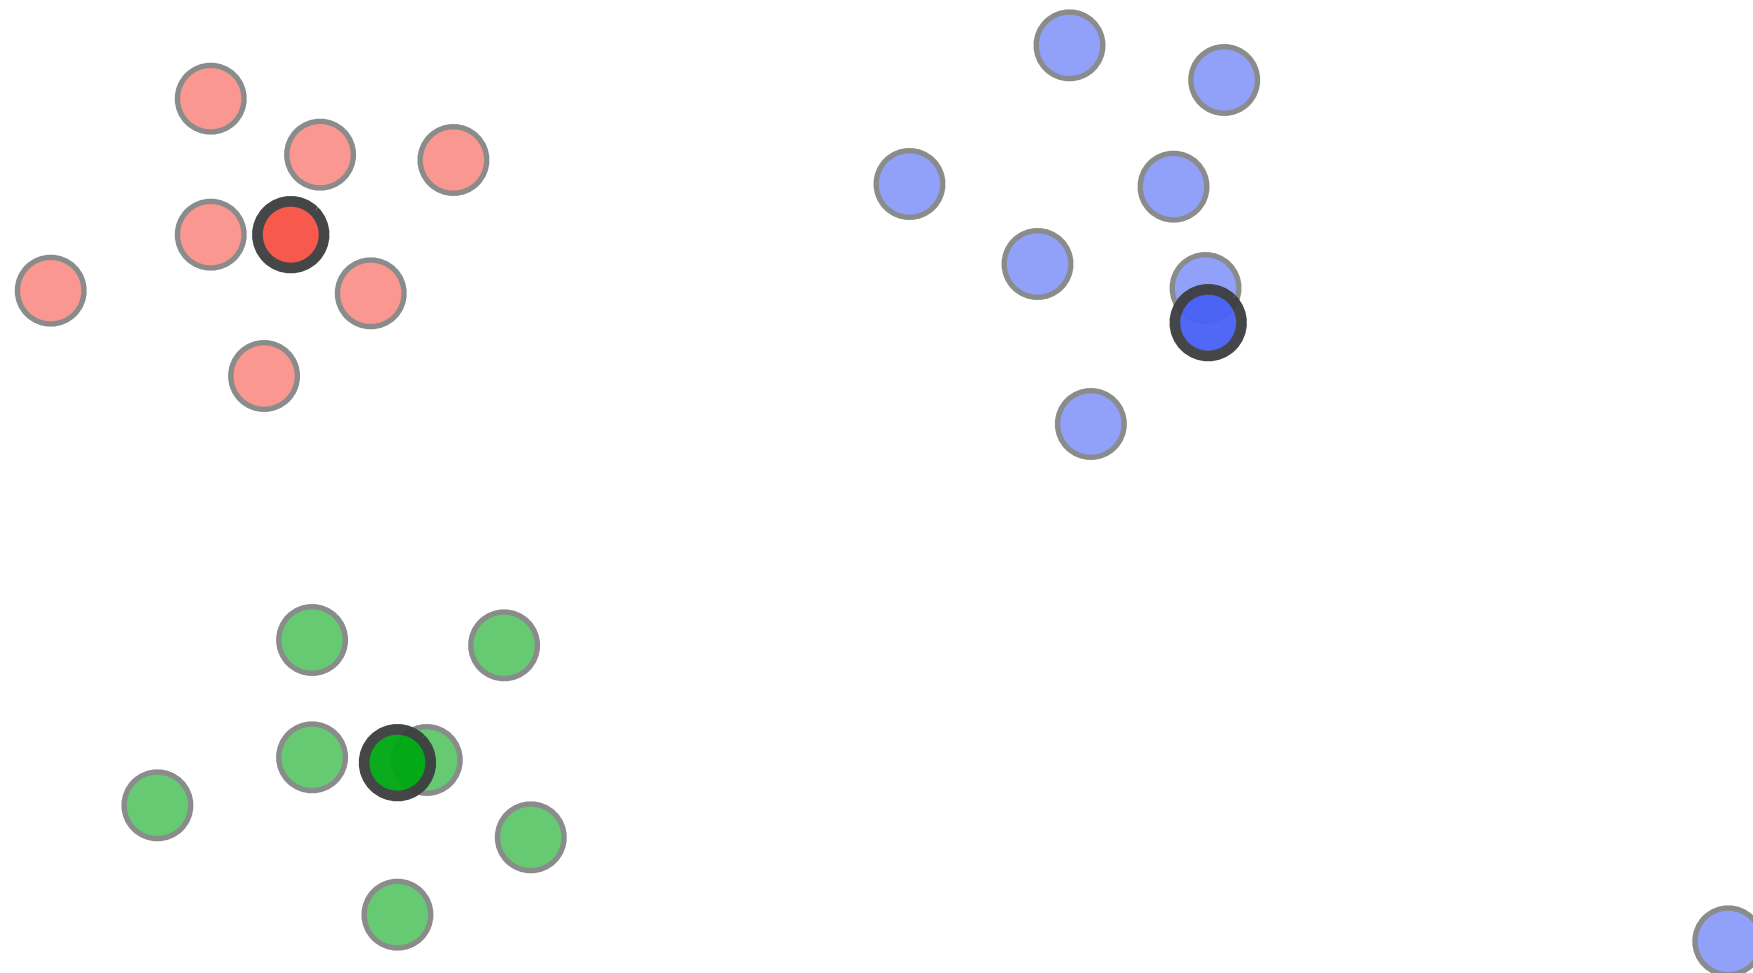
iteration phase :

iterate as in the k-means algorithm until convergence

# k-means++ initialization



# k-means++ initialization





# k-means++ provable guarantee

theorem:

k-means++ is  $O(\log k)$  approximate in expectation

# k-means++ provable guarantee

approximation guarantee comes just from the first iteration  
(initialization)

subsequent iterations can only improve cost

# k-means++ analysis

consider optimal clustering  $C^*$

assume k-means++ selects a center from a new optimal cluster  
then

k-means++ is 8-approximate in expectation

intuition: if no points from a cluster are picked, then it probably does not contribute much to the overall error

an inductive proof shows that the algorithm is  $O(\log k)$   
approximate

# k-means++ proof : first cluster

fix an optimal clustering  $C^*$

first center is selected uniformly at random

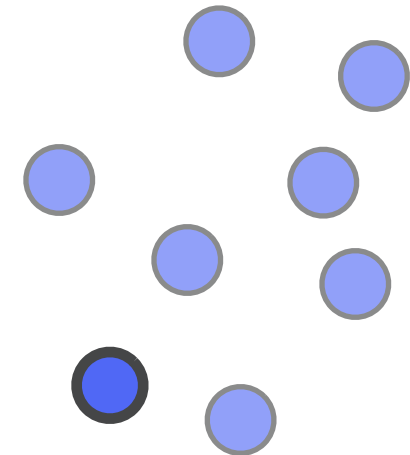
bound the total error of the points in the optimal cluster of the first center

# k-means++ proof : first cluster

let  $A$  be the first cluster

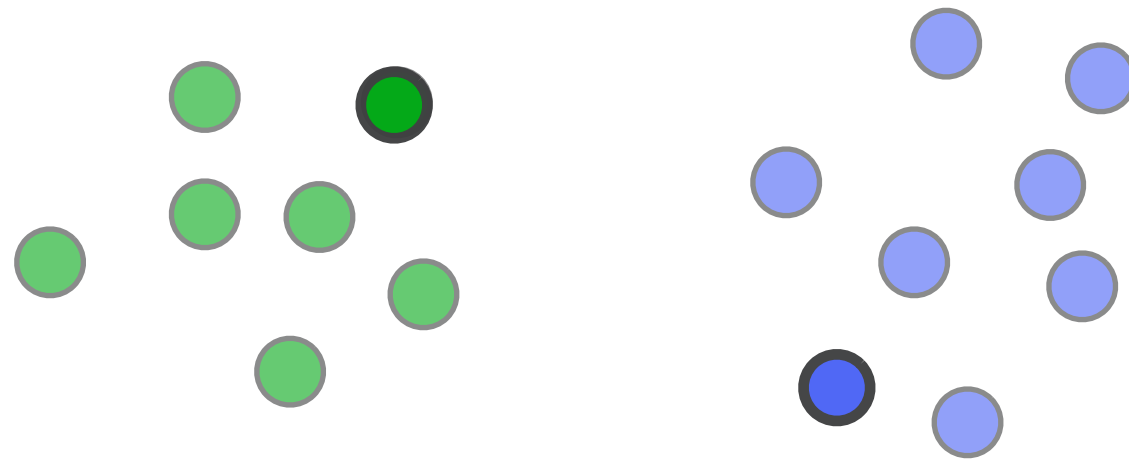
each point  $a_0 \in A$  is **equally likely** to be selected as center

**expected error:**



$$\begin{aligned} E[\phi(A)] &= \sum_{a_0 \in A} \frac{1}{|A|} \sum_{a \in A} \|a - a_0\|^2 \\ &= 2 \sum_{a \in A} \|a - \bar{A}\|^2 = 2\phi^*(A) \end{aligned}$$

# k-means++ proof : other clusters



suppose next center is selected from a **new cluster** in the optimal clustering  $C^*$

**bound** the **total error** of **that cluster**

# k-means++ proof : other clusters

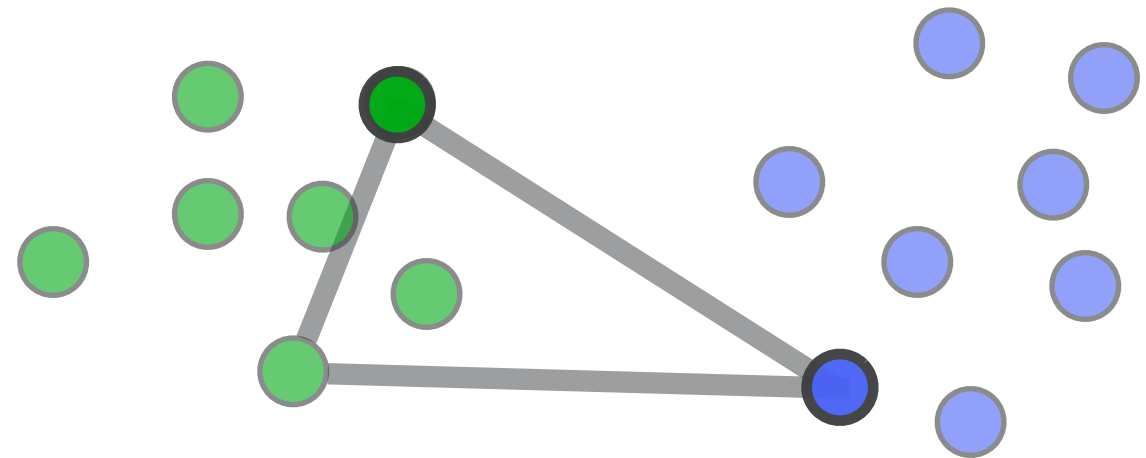
let **B** be the second cluster and **b<sub>0</sub>** the center selected

$$E[\phi(B)] = \sum_{b_0 \in B} \frac{D^2(b_0)}{\sum_{b \in B} D^2(b)} \sum_{b \in B} \min\{D(b), \|b - b_0\|^2\}$$

triangle inequality:

$$D(b_0) \leq D(b) + \|b - b_0\|$$

$$D^2(b_0) \leq 2D^2(b) + 2\|b - b_0\|^2$$



# k-means++ proof : other clusters

$$D^2(b_0) \leq 2D^2(b) + 2||b - b_0||^2$$

average over all points  $b$  in  $B$

$$D^2(b_0) \leq \frac{2}{|B|} \sum_{b \in B} D^2(b) + \frac{2}{|B|} \sum_{b \in B} ||b - b_0||^2$$

recall

$$\begin{aligned} E[\phi(B)] &= \sum_{b_0 \in B} \frac{D^2(b_0)}{\sum_{b \in B} D^2(b)} \sum_{b \in B} \min\{D(b), ||b - b_0||^2\} \\ &\leq 4 \sum_{b \in B} \frac{1}{|B|} \sum_{b_0 \in B} ||b - b_0||^2 = 4 \sum_{b \in B} 2||b - \bar{B}||^2 = 8\phi^*(B) \end{aligned}$$



# k-means++ analysis

if that k-means++ selects a center from a new optimal cluster  
then

k-means++ is 8-approximate in expectation

an inductive proof shows that the algorithm is  $O(\log k)$   
approximate

# lesson learned

no reason to use **k-means** and not **k-means++**

**k-means++** :

- easy to implement

- provable guarantee

- works well in practice