Aalto University
School of Science

CS-E4600 – Algorithmic methods of data mining

Slide set 11 : computing basic statistics

Aristides Gionis
Aalto University

fall 2018

algorithmic tools

efficiency considerations

- data in the web and social-media are typically of extremely large scale (easily reach to billions)
- how to compute simple graph statistics?
- even quadratic algorithms are not feasible in practice

hashing and sketching

- probabilistic / approximate methods
- sketching : create sketches that summarize the data and allow to estimate simple statistics with small space
- hashing : hash objects in such a way that similar objects have larger probability of mapped to the same value than non-similar objects

estimator theorem

- consider a set of items U
- a fraction ρ of them have a specific property
- estimate ρ by sampling

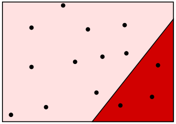

- how many samples N are needed?

estimator theorem

- consider a set of items U
- a fraction ρ of them have a specific property
- estimate ρ by sampling

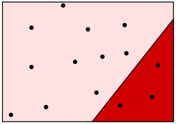

- how many samples N are needed?

$$N \geq \frac{4}{\epsilon^2 \rho} \log \frac{2}{\delta}.$$

for an ϵ-approximation with probability at least $1 - \delta$

estimator theorem

- consider a set of items U
- a fraction ρ of them have a specific property
- estimate ρ by sampling

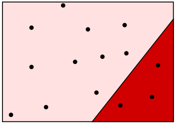

- how many samples N are needed?

$$N \geq \frac{4}{\epsilon^2 \rho} \log \frac{2}{\delta}.$$

  for an ϵ-approximation with probability at least $1 - \delta$
- notice: it does not depend on $|U|$ (!)
- but it depends on ρ

  useful when we have a lower bound on ρ

homework

use the Chernoff bound to derive the estimator theorem applications of the algorithmic tools to real scenarios

clustering coefficient and triangles

clustering coefficient

$$C = \frac{3 \times \text{number of triangles in the network}}{\text{number of connected triples of vertices}}$$

- how to compute it?
- how to compute the number of triangles in a graph?
- assume that the graph is very large, stored in disk

[Buriol et al., 2006]

- count triangles when graph is seen as a data stream
- two models:
  - edges are stored in any order
  - edges in order : all edges incident to one vertex are stored sequentially

counting triangles

- brute-force algorithm is checking every triple of vertices
- obtain an approximation by sampling triples

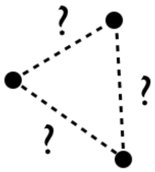

sampling algorithm for counting triangles

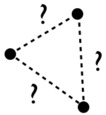

- how many samples are required?
- let T be the set of all triples and T_i the set of triples that have i edges, $i = 0, 1, 2, 3$
- by the estimator theorem, to get an ϵ-approximation, with probability $1 - \delta$, the number of samples should be

$$N \geq O\left(\frac{|T|}{|T_3|} \frac{1}{\epsilon^2} \log \frac{1}{\delta}\right)$$

- but $|T|$ can be very large compared to $|T_3|$

counting triangles

- incidence model : all edges incident to each vertex appear in order in the stream
- sample connected triples

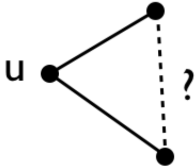

sampling algorithm for counting triangles

- incidence model
- consider sample space $\mathcal{S} = \{b\text{-}a\text{-}c \mid (a, b), (a, c) \in E\}$
- $|\mathcal{S}| = \sum_i d_i(d_i - 1)/2$

1: sample $X \subseteq \mathcal{S}$ (paths $b\text{-}a\text{-}c$)
2: estimate fraction of X for which edge (b, c) is present
3: scale by $|\mathcal{S}|$

- gives (ϵ, δ) approximation

counting triangles — incidence stream model

SAMPLETRIANGLE [Buriol et al., 2006]

1st pass

count the number of paths of length 2 in the stream

2nd pass

uniformly choose one path (a, b, c)

3rd pass

if $((b, c) \in E)$ $\beta = 1$ else $\beta = 0$

return β

counting triangles — incidence stream model

SAMPLETRIANGLE [Buriol et al., 2006]

1st pass

count the number of paths of length 2 in the stream

2nd pass

uniformly choose one path (a, b, c)

3rd pass

if $((b, c) \in E)$ $\beta = 1$ else $\beta = 0$

return β

we have $\mathrm{E}[\beta] = \frac{3|T_3|}{|T_2|+3|T_3|}$, with $|T_2| + 3|T_3| = \sum_u \frac{d_u(d_u-1)}{2}$, so

$$|T_3| = \mathrm{E}[\beta] \sum_u \frac{d_u(d_u - 1)}{6}$$

and space needed is $O((1 + \frac{|T_2|}{|T_3|})\frac{1}{\epsilon^2} \log \frac{1}{\delta})$

properties required to apply the estimator theorem

it should be possible to

- estimate the size of the sampling space
- sample an element uniformly at random

also

- quantity of interest should not be very small

  e.g., $\mathcal{O}(\frac{1}{n})$ is not very good

  while $\mathcal{O}(\frac{1}{\log n})$ or $\mathcal{O}(1)$ are good

counting graph minors

counting other minors

- count all minors in a very large graphs
  - connected subgraphs
  - size 3 and 4
  - directed or undirected graphs

- why?
- modeling networks, “signature” structures
  e.g., copying model
- anomaly detection, e.g., spam link farms
  [Alon, 2007, Bordino et al., 2008]

counting minors in large graphs

- characterize a graph by the distribution of its minors

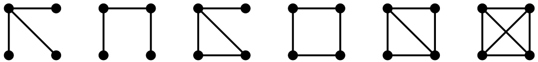

all undirected minors of size 4

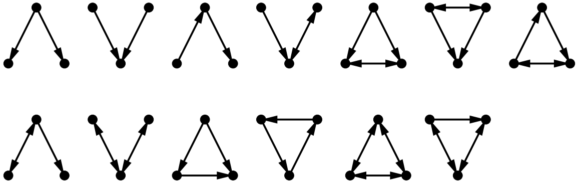

all directed minors of size 3

sampling algorithm for counting triangles

- incidence model
- consider sample space $\mathcal{S} = \{b\text{-}a\text{-}c \mid (a, b), (a, c) \in E\}$
- $|\mathcal{S}| = \sum_i d_i(d_i - 1)/2$

1: sample $X \subseteq \mathcal{S}$ (paths b-a-c)
2: estimate fraction of X for which edge (b, c) is present
3: scale by $|\mathcal{S}|$

- gives (ϵ, δ) approximation

adapting the algorithm

sampling spaces:

- 3-node directed

- 4-node undirected

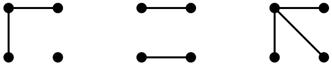

are the sampling space properties satisfied?

datasets

graph class	type	# instances
synthetic	un/directed	39
wikipedia	un/directed	7
webgraphs	un/directed	5
cellular	directed	43
citation	directed	3
food webs	directed	6
word adjacency	directed	4
author collaboration	undirected	5
autonomous systems	undirected	12
protein interaction	undirected	3
US road	undirected	12

clustering of undirected graphs

assigned to	0	1	2	3	4	5	6
AS graph	12	0	0	0	0	0	0
collaboration	0	0	3	2	0	0	0
protein	1	0	0	1	0	0	1
road-graph	0	12	0	0	0	0	0
wikipedia	0	0	0	0	2	5	0
synthetic	11	0	0	0	0	0	28
webgraph	2	0	0	1	0	0	0

clustering of directed graphs

feature class	accuracy compared to ground truth
standard topological properties (81)	0.74%
minors of size 3	0.78%
minors of size 4	0.84%
minors of size 3 and 4	0.91%

graph distance distributions

small-world phenomena

small worlds : graphs with short paths

- Stanley Milgram (1933-1984)
  “The man who shocked the world”
- obedience to authority (1963)
- small-world experiment (1967)

Milgram's experiment

- 300 people (starting population) are asked to dispatch a parcel to a single individual (target)
- the target was a Boston stockbroker
- the starting population is selected as follows:
  - 100 were random Boston inhabitants (group A)
  - 100 were random Nebraska strockbrokers (group B)
  - 100 were random Nebraska inhabitants (group C)

Milgram’s experiment

- rules of the game :
- parcels could be directly sent only to someone the sender knows personally
- 453 intermediaries happened to be involved in the experiments (besides the starting population and the target)

Milgram's experiment

questions Milgram wanted to answer:

1. how many parcels will reach the target?
2. what is the distribution of the number of hops required to reach the target?
3. is this distribution different for the three starting subpopulations?

Milgram's experiment

answers to the questions

1. how many parcels will reach the target?
   29%
2. what is the distribution of the number of hops required to reach the target?
   average was 5.2
3. is this distribution different for the three starting subpopulations?
   YES: average for groups A/B/C was 4.6/5.4/5.7

chain lengths

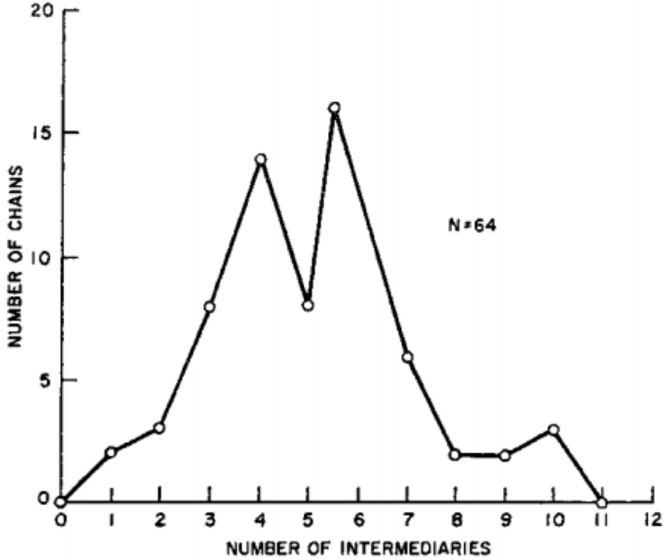

FIGURE 1

measuring what?

but what did Milgram’s experiment reveal, after all?

1. the the world is small
2. that people are able to exploit this smallness

graph distance distribution

- obtain information about a large graph, i.e., social network
- macroscopic level
- distance distribution
  - mean distance
  - median distance
  - diameter
  - effective diameter
  - ...

graph distance distribution

- given a graph, $d(x, y)$ is the length of the shortest path from x to y, defined as ∞ if one cannot go from x to y
- for undirected graphs, $d(x, y) = d(y, x)$
- for every t, count the number of pairs (x, y) such that $d(x, y) = t$
- the fraction of pairs at distance t is a distribution

exact computation

how can one compute the distance distribution?

exact computation

how can one compute the distance distribution?

- weighted graphs: Dijkstra (single-source: $O(m \log n)$),
- Floyd-Warshall (all-pairs: $O(n^3)$)
- in the unweighted case:
  - a single BFS solves the single-source version of the problem: $O(m)$
  - if we repeat it from every source: $O(nm)$

sampling pairs

- sample at random pairs of nodes (x, y)
- compute $d(x, y)$ with a BFS from x
- (possibly: reject the pair if $d(x, y)$ is infinite)

sampling pairs

- for every t, the fraction of sampled pairs that were found at distance t are an estimator of the value of the probability mass function
- takes a BFS for every pair — $O(m)$

sampling sources

- sample at random a source t
- compute a full BFS from t
- but it is an unbiased estimator only for undirected and connected graphs
- also BFS is not cache-friendly

idea : diffusion

[Palmer et al., 2002]

- let $B_t(x)$ be the ball of radius t around x
  (the set of nodes at distance $\leq t$ from x)
- clearly $B_0(x) = \{x\}$
- moreover $B_{t+1}(x) = \bigcup_{(x,y)} B_t(y) \bigcup \{x\}$
- so computing B_{t+1} from B_t just takes a single (sequential) scan of the graph

easy but costly

- every set requires $O(n)$ bits, hence $O(n^2)$ bits overall
- easy but costly
- too many!
- what about using approximated sets?
- we need probabilistic counters, with just two primitives: add and size
- very small!

estimating the number of distinct values (F_0)

[Flajolet and Martin, 1985]

- consider a bit vector **b** with $O(\log n)$ bits
- initialize **b** to $[0, \ldots, 0]$
- consider a hash function f that maps each item x to the j-th bit of the bit-vector **b** with probability $1/2^j$
- for each item x_i in the data stream

  set the bit $j = f(x_i)$ of **b** equal to 1

  (important: bits are set deterministically for each x_i)
- let R be the index of the largest bit set
- return $Y = 2^R$

ANF

- probabilistic counter for approximating the number of distinct values [Flajolet and Martin, 1985]
- ANF algorithm [Palmer et al., 2002]
  uses the original probabilist counters
- HyperANF algorithm [Boldi et al., 2011]
  uses HyperLogLog counters [Flajolet et al., 2007]

HyperANF

- HyperLogLog counter [Flajolet et al., 2007]
- with 40 bits you can count up to 4 billion with a standard deviation of 6%
- remember: one set per node

implementation tricks

[Boldi et al., 2011]

- use broad-word programming to compute union efficiently
- systolic computation for on-demand updates of counters
- exploit micro-parallelization of multicore architectures

performance

- HADI, a Hadoop-conscious implementation of ANF [Kang et al., 2011]
- takes 30 minutes on a 200K-node graph (on one of the 50 world largest supercomputers)
- HyperANF does the same in 2.25min on a workstation (20 min on a laptop).

experiments on facebook

[Backstrom et al., 2011]

considered only active users

- it : only italian users
- se : only swedish users
- it + se : only italian and swedish users
- us : only US users
- the whole facebook (750m nodes)

based on users current geo-IP location

distance distribution (it)

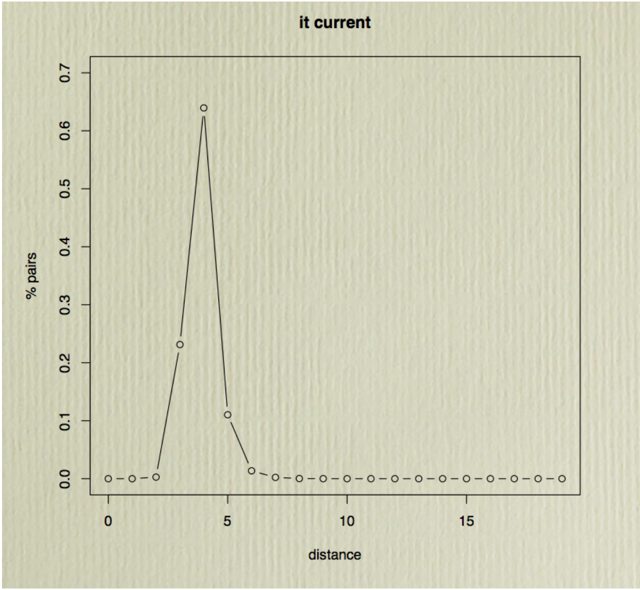

distance distribution (se)

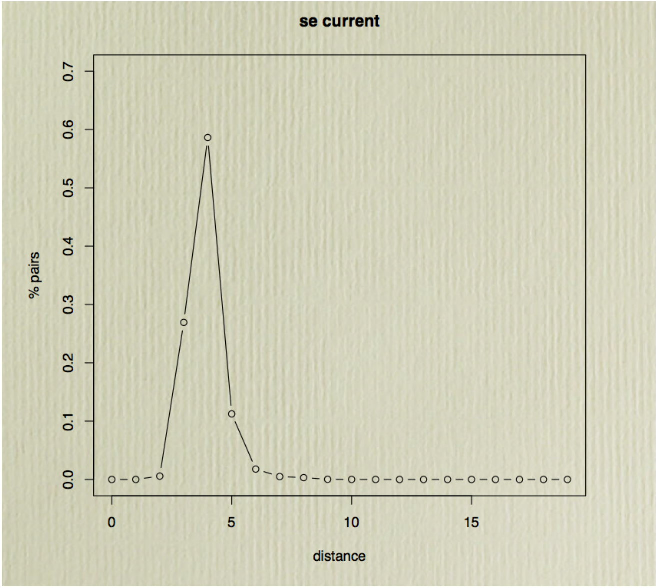

distance distribution (fb)

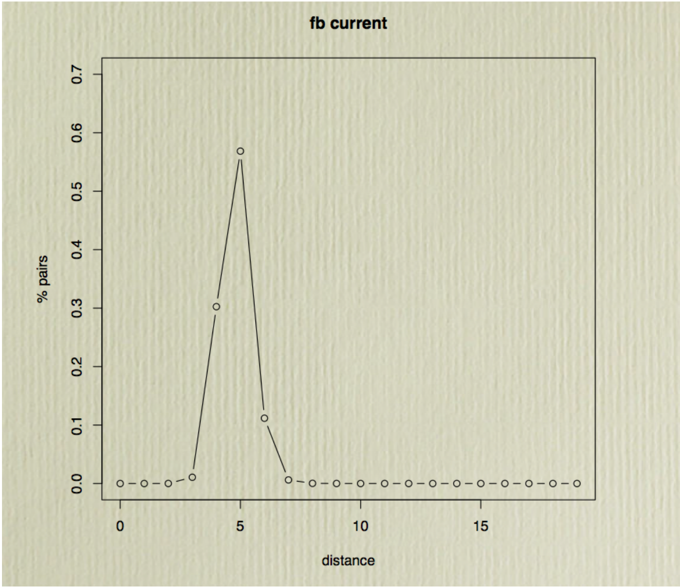

average distance

	2008	2012
it	6.58	3.90
se	4.33	3.89
it+se	4.90	4.16
us	4.74	4.32
fb	5.28	4.74

fb 2012 : 92% pairs are reachable!

effective diameter

	2008	2012
it	9.0	5.2
se	5.9	5.3
it+se	6.8	5.8
us	6.5	5.8
fb	7.0	6.2

actual diameter

	2008	2012
it	> 29	$= 25$
se	> 16	$= 25$
it+se	> 21	$= 27$
us	> 17	$= 30$
fb	> 17	> 58

breaking the news

HOME PAGE | TODAY'S PAPER | VIDEO | MOST POPULAR | TIMES TOPICS

Log In | Register Now | Help

The New York Times

Business Day

Technology

Search All NYTimes.com

Go

WORLD | U.S. | N.Y. / REGION | BUSINESS | TECHNOLOGY | SCIENCE | HEALTH | SPORTS | OPINION | ARTS | STYLE | TRAVEL | JOBS | REAL ESTATE | AUTOS

Separating You and Me? 4.74 Degrees

By JOHN MARKOFF and SOMINI SENGUPTA

Published: November 21, 2011

The world is even smaller than you thought.

Enlarge This Image

Cornell News Service

Jon Kleinberg of Cornell said weak ties could be important.

Adding a new chapter to the research that cemented the phrase “six degrees of separation” into the language, scientists at Facebook and the University of Milan reported on Monday that the average number of acquaintances separating any two people in the world was not six but 4.74.

The original “six degrees” finding, published in 1967 by the psychologist Stanley Milgram, was drawn from 296 volunteers who were asked to send a message by postcard, through friends and then friends of friends, to a specific person in a Boston suburb.

RECOMMEND

TWITTER

LINKEDIN

SIGN IN TO E-MAIL

PRINT

REPRINTS

SHARE

Log in to see what your friends are sharing on nytimes.com. Privacy Policy | What's This?

Log In With Facebook

What's Popular Now

Marvin Hamlisch, Composer, Dies at 68

France's 'les Riches' Vow to Leave if 75% Tax Rate Is Passed

acknowledgements

Paolo Boldi   Charalampos Tsourakakis

references

Alon, U. (2007).
Network motifs: theory and experimental approaches.
Nature Reviews Genetics.

Backstrom, L., Boldi, P., Rosa, M., Ugander, J., and Vigna, S. (2011).
Four degrees of separation.
CoRR, abs/1111.4570.

Boldi, P., Rosa, M., and Vigna, S. (2011).
HyperANF: approximating the neighborhood function of very large graphs on a budget.
In *WWW*.

Bordino, I., Donato, D., Gionis, A., and Leonardi, S. (2008).
Mining large networks with subgraph counting.
In *ICDM*.

references (cont.)

Buriol, L. S., Frahling, G., Leonardi, S., Marchetti-Spaccamela, A., and Sohler, C. (2006).
Counting triangles in data streams.
In *PODS '06: Proceedings of the twenty-fifth ACM SIGMOD-SIGACT-SIGART symposium on Principles of database systems*, pages 253–262, New York, NY, USA. ACM Press.

Flajolet, F., Fusy, E., Gandouet, O., and Meunier, F. (2007).
Hyperloglog: the analysis of a near-optimal cardinality estimation algorithm.
In *Proceedings of the 13th conference on analysis of algorithm (AofA)*.

Flajolet, P. and Martin, N. G. (1985).
Probabilistic counting algorithms for data base applications.
Journal of Computer and System Sciences, 31(2):182–209.

Kang, U., Tsourakakis, C. E., Appel, A. P., Faloutsos, C., and Leskovec, J. (2011).
HADI: Mining radii of large graphs.
ACM TKDD, 5.

references (cont.)

Palmer, C. R., Gibbons, P. B., and Faloutsos, C. (2002).

ANF: a fast and scalable tool for data mining in massive graphs.

In *Proceedings of the eighth ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 81–90, New York, NY, USA. ACM Press.